

CSCI596: Scientific Computing & Visualization

Aiichiro Nakano

Collaboratory for Advanced Computing & Simulations

Department of Computer Science

Department of Physics & Astronomy

Department of Chemical Engineering & Materials Science

Department of Biological Sciences

University of Southern California


Email: anakano@usc.edu

- **Computational science:** An area of scientific investigation, where computers play a central role—do science using computer (e.g., computational chemistry).
- **Scientific computing:** An area in computer science (CS) to support computational sciences by innovative use of computer systems—parallel computing, scientific visualization, etc.




New Computing Architecture


4.9 trillion-atom molecular dynamics
40 trillion-d.o.f. quantum mechanics



Global Grid of supercomputers




Toward exaflop/s computers



petaflop/s | exaflop/s = 10^{15} | 10^{18} floating-point operations per second




Many-core CPU computing



Intel 80-core chip
1Tflop/s@62W

10²-cluster Grid
10⁵-node cluster
10³-core node

Many-core CPU/GPU Computing



Godson-T Many-core Architecture

J. Parallel Distrib. Comput. 73 (2013) 1469–1482



Contents lists available at ScienceDirect

J. Parallel Distrib. Comput.

journal homepage: www.elsevier.com/locate/jpdc



Scalability study of molecular dynamics simulation on Godson-T many-core architecture



狗剩




Liu Peng^{a,*}, Guangming Tan^{b,*}, Rajiv K. Kalia^a, Aiichiro Nakano^a, Priya Vasishta^a, Dongrui Fan^b, Hao Zhang^b, Fenglong Song^b

^a Collaboratory for Advanced Computing and Simulations, University of Southern California, Los Angeles, CA, 90089, USA

^b Key Laboratory of Computer System and Architecture, Institute of Computing Technology, Chinese Academy of Sciences, Beijing, 100190, China




Multicore Processors



- Multiple simple processors (or cores) sharing common memory


World's Biggest Chip

Cerebras AI engine



Cerebras WSE

1.2 Trillion transistors
46,225 mm² silicon




Largest GPU


21.1 Billion transistors
815 mm² silicon

Computer Simulation


Run real applications on advanced computing architectures



Computer Experiment



Mathematical Model



TITLE PAGE OF THE FIRST EDITION OF THE PRINCIPIA
(See Appendix, Note 2, page 622)

AXIOMS, OR LAWS OF MOTION¹

LAW I

Every body continues in its state of rest, or of uniform motion in a right line, unless it is compelled to change that state by forces impressed upon it.

PROJECTILES continue in their motions, so far as they are not retarded by the resistance of the air, or impelled downwards by the force of gravity. A top, whose parts by their cohesion are continually drawn aside from rectilinear motions, does not cease its rotation, otherwise than as it is retarded by the air. The greater bodies of the planets and comets, meeting with less resistance in freer spaces, preserve their motions both progressive and circular for a much longer time.

LAW II²

The change of motion is proportional to the motive force impressed; and is made in the direction of the right line in which that force is impressed.

If any force generates a motion, a double force will generate double the motion, a triple force triple the motion, whether that force be impressed altogether and at once, or gradually and successively. And this motion (being always directed the same way with the generating force), if the body moved before, is added to or subtracted from the former motion, according as they directly conspire with or are directly contrary to each other; or obliquely joined, when they are oblique, so as to produce a new motion compounded from the determination of both.

LAW III

To every action there is always opposed an equal reaction: or, the mutual actions of two bodies upon each other are always equal, and directed to contrary parts.

Whatever draws or presses another is as much drawn or pressed by that other. If you press a stone with your finger, the finger is also pressed by the

[¹ Appendix, Note 14.] [² Appendix, Note 15.]

[13]

LAW I

Every body continues in its state of rest, or of uniform motion in a right line, unless it is compelled to change that state by forces impressed upon it.

LAW II²

The change of motion is proportional to the motive force impressed; and is made in the direction of the right line in which that force is impressed.

14


NEWTON'S MATHEMATICAL PRINCIPLES

stone. If a horse draws a stone tied to a rope, the horse (if I may so say) will be equally drawn back towards the stone; for the distended rope, by the same endeavor to relax or unbend itself, will draw the horse as much towards the stone as it does the stone towards the horse, and will obstruct the progress of the one as much as it advances that of the other. If a body impinge upon another, and by its force change the motion of the other, that body also (because of the equality of the mutual pressure) will undergo an equal change, in its own motion, towards the contrary part. The changes made by these actions are equal, not in the velocities but in the motions of bodies; that is to say, if the bodies are not hindered by any other impediments. For, because the motions are equally changed, the changes of the velocities made towards contrary parts are inversely proportional to the bodies. This law takes place also in attractions, as will be proved in the next Scholium.

COROLLARY I

A body, acted on by two forces simultaneously, will describe the diagonal of a parallelogram in the same time as it would describe the sides by those forces separately.

If a body in a given time, by the force M impressed apart in the place A, should with an uniform motion be carried from A to B, and by the force N impressed apart in the same place, should be carried from A to C, let the parallelogram ABCD be completed, and, by both forces acting together, it will in the same time be carried in the diagonal from A to D. For since the force N acts in the direction of the line AC, parallel to BD, this force (by the second Law) will not at all alter the velocity generated by the other



force M, by which the body is carried towards the line BD. The body therefore will arrive at the line BD in the same time, whether the force N be impressed or not; and therefore at the end of that time it will be found somewhere in the line BD. By the same argument, at the end of the same time it will be found somewhere in the line CD. Therefore it will be found in the point D, where both lines meet. But it will move in a right line from A to D, by Law I.

Towards 2020 Science

Some 400 years later, Newton, in his efforts to understand the natural laws of the rate of change in motion, used algebra to underpin another new branch of mathematics: calculus (a branch for which von Leibniz is simultaneously and independently credited). Calculus spurred scientists “to go off looking for other laws of nature that could explain natural phenomenon in terms of rates of change and found them by the bucketful - heat, sound, light, fluid dynamics, electricity and magnetism” [2].

A scientific revolution is just beginning.

However, what this report uncovers, for the first time, is a fundamentally important shift from *computers* supporting scientists to ‘do’ traditional science to computer science becoming embedded into the very fabric of science and how science is done, creating what I am prepared to go so far as to call ‘new kinds’ of science¹.

Stephen Emmott

<http://research.microsoft.com/towards2020science>
<http://www.nature.com/nature/journal/v440/n7083>


Molecular Dynamics Simulation

- Newton's equation of motion



$$m_i \frac{d^2 \mathbf{r}_i}{dt^2} = -\frac{\partial V(\mathbf{r}^N)}{\partial \mathbf{r}_i} \quad (i = 1, \dots, N)$$

- Many-body interatomic potential

$$V = \sum_{i < j} u_{ij}(|\mathbf{r}_{ij}|) + \sum_{i, j < k} v_{jik}(\mathbf{r}_{ij}, \mathbf{r}_{ik})$$




- Application: drug design, robotics, entertainment, etc.




Molecular Dynamics in Graphics

Particle modeling of dust & smoke



Jim Chen (George Mason)


Ron Fedkiw (Stanford)



MD Algorithm

Time discretization: differential → algebraic equation

$$\begin{cases} \vec{r}_i(t + \Delta) = \vec{r}_i(t) + \vec{v}_i(t)\Delta + \frac{1}{2}\vec{a}_i(t)\Delta^2 \\ \vec{v}_i(t + \Delta) = \vec{v}_i(t) + \frac{\vec{a}_i(t) + \vec{a}_i(t + \Delta)}{2}\Delta \end{cases} \quad \vec{a}_i = -\frac{1}{m}\frac{\partial V}{\partial \vec{r}_i}$$



Snapshots!



Time stepping: Velocity Verlet algorithm

Given $(\vec{r}_i(t), \vec{v}_i(t))$

1. Compute $\vec{a}_i(t)$ as a function of $\{\vec{r}_i(t)\}$
2. $\vec{v}_i\left(t + \frac{\Delta}{2}\right) \leftarrow \vec{v}_i(t) + \frac{\Delta}{2}\vec{a}_i(t)$
3. $\vec{r}_i(t + \Delta) \leftarrow \vec{r}_i(t) + \vec{v}_i\left(t + \frac{\Delta}{2}\right)\Delta$
4. Compute $\vec{a}_i(t + \Delta)$ as a function of $\{\vec{r}_i(t + \Delta)\}$
5. $\vec{v}_i(t + \Delta) \leftarrow \vec{v}_i\left(t + \frac{\Delta}{2}\right) + \frac{\Delta}{2}\vec{a}_i(t + \Delta)$

Physical phenomenon

Mathematical model

Discrete algebraic approximation

Numerical algorithms

Simulation program

Computer experiment

MD Program

```
for (n=0; n<nAtom; n++)
    for (k=0; k<3; k++)
        rv[n][k] = rv[n][k] + DeltaT/2*ra[n][k];
```

$$\vec{v}_i(t + \frac{\Delta}{2}) \leftarrow \vec{v}_i(t) + \frac{\Delta}{2} \vec{a}_i(t)$$


```
for (n=0; n<nAtom; n++)
    for (k=0; k<3; k++)
        r[n][k] = r[n][k] + DeltaT*rv[n][k];
```

$$\vec{r}_i(t + \Delta) \leftarrow \vec{r}_i(t) + \vec{v}_i(t + \frac{\Delta}{2})\Delta$$

```
ComputeAccel(); // r[][] → ra[][]
```





```
for (n=0; n<nAtom; n++)
    for (k=0; k<3; k++)
        rv[n][k] = rv[n][k] + DeltaT/2*ra[n][k];
```

$$\vec{v}_i(t + \frac{\Delta}{2}) \leftarrow \vec{v}_i(t) + \frac{\Delta}{2} \vec{a}_i(t)$$



Computer Experiment


200-500 million atom molecular dynamics simulation of hypervelocity (15 km/s) impact on AlN & Al₂O₃ plate



Type of Mathematical Models


	Particle model (ordinary differential equations)	Continuum model (partial differential equations)
Deterministic	molecular dynamics	computational fluid dynamics, continuum mechanics
Stochastic	Monte Carlo particle simulation	quantum Monte Carlo

Particle model of oxidation



Tim Campbell
USC-CACS

Continuum model of water flow



Ron Fedkiw (Stanford)
graphics.stanford.edu/~fedkiw

Continuum Model: Quantum Mechanics

Challenge: Complexity of quantum N -body problem

Density functional theory (DFT)

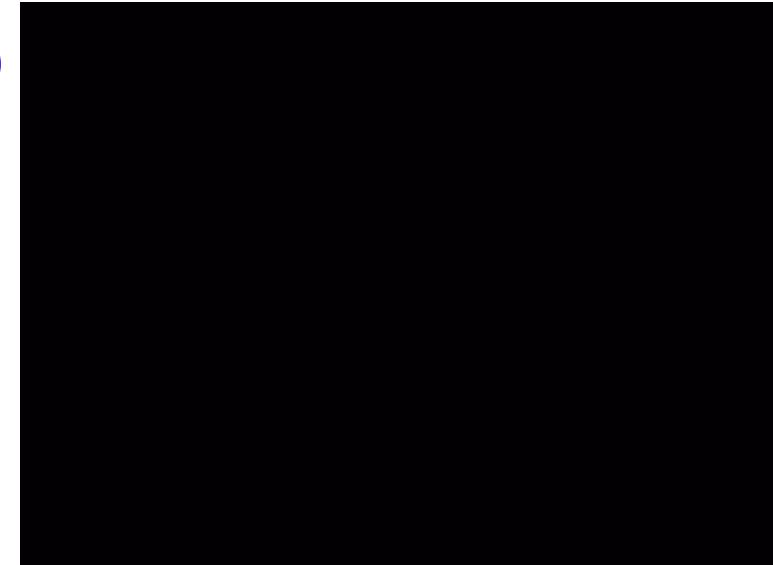
(Walter Kohn, Nobel Chemistry Prize, '98)



$$\psi(\mathbf{r}_1, \dots, \mathbf{r}_{N_{\text{el}}}) \quad O(C^N)$$



$$\{\psi_n(\mathbf{r}) | n = 1, \dots, N_{\text{el}}\} \quad O(N^3)$$



Constrained minimization problem:

Minimize:

$$E[\{\psi_n\}] = \sum_{n=1}^{N_{\text{el}}} \int d\mathbf{r} \psi_n^*(\mathbf{r}) \left(-\frac{\hbar^2}{2m_{\text{el}}} \frac{\partial^2}{\partial \mathbf{r}^2} + V_{\text{ion}}(\mathbf{r}) \right) \psi_n(\mathbf{r}) + \frac{e^2}{2} \iint d\mathbf{r} d\mathbf{r}' \frac{\rho(\mathbf{r})\rho(\mathbf{r}')}{|\mathbf{r}-\mathbf{r}'|} + E_{\text{xc}}[\rho(\mathbf{r})]$$

with orthonormal constraints: $\int d^3r \psi_m^*(\mathbf{r}) \psi_n(\mathbf{r}) = \delta_{mn}$

$$\text{Charge density: } \rho(\mathbf{r}) = \sum_{n=1}^{N_{\text{el}}} |\psi_n(\mathbf{r})|^2$$



DFT: It's Algorithm

Density functional theory (DFT)

Hohenberg & Kohn, *Phys. Rev.* **136**, B864 ('64); W. Kohn [Nobel chemistry prize, '98]

$$\begin{array}{ccc} O(C^N) & \xrightarrow{\hspace{1cm}} & O(N^3) \\ 1 \text{ } N\text{-electron problem} & & N \text{ } 1\text{-electron problems} \end{array}$$



Volume 140, Issue 18, 14 May 2014

SPECIAL TOPIC: ADVANCES IN DENSITY FUNCTIONAL THEORY

A divide-conquer-recombine algorithmic paradigm for large spatiotemporal quantum molecular dynamics simulations

Fuyuki Shimojo^{1,2}, Shinnosuke Hattori^{1,2}, Rajiv K. Kalia¹, Manaschai Kunaseth^{1,3}, Weiwei Mou¹, Aiichiro Nakano¹, Ken-ichi Nomura¹, Satoshi Ohmura^{1,2,4}, Pankaj Rajak¹, Kohei Shimamura^{1,2,5} and Priya Vashishta¹

+ VIEW AFFILIATIONS

J. Chem. Phys. **140**, 18A529 (2014); <http://dx.doi.org/10.1063/1.4869342>

$$\begin{array}{ccc} O(N^3) & \xrightarrow{\hspace{1cm}} & O(N) \\ \text{Mean-field theory} & & \text{Divide-conquer-recombine} \end{array}$$

Divide-Conquer-(Re)combine

- “The first was to never accept anything as true which I could not accept as obviously true. The second was to divide each of the problems in as many parts as I should to solve them. The third, beginning with the simplest and easiest to understand matters, little by little, to the most complex knowledge. And the last resolution was to make my enumerations so complete and my reviews so general that I could be assured that I had not omitted anything.” (René Descartes, *Discourse on Method*, 1637)

divide (conquer) recombine

- 「モデルの分割一再統合の方法の優れた点は、分割した要素的概念を、モデルの理解に役立つよう再構成することができ、そこに創造の入り込む余地があるという点にある。」/(福井謙一
学問の創造、1987)

room for creativity

Kenichi Fukui [Nobel Chemistry Prize, '98]



Stochastic Model of Stock Prices

Fluctuation in stock price

Market Summary > Zoom Video Communications Inc
NASDAQ: ZM

282.28 USD **-7.40 (2.55%)** +
Closed: Aug 24, 5:38 PM EDT · Disclaimer
After hours 282.35 **+0.070 (0.025%)**




Basis of Black-Scholes analysis of option prices



$$dS = \mu S dt + \sigma S \varepsilon \sqrt{dt}$$

(1997 Nobel Economy Prize to Myron Scholes)

Computational stock portfolio trading



Andrey Omeltchenko ([Quantlab](#))

First Gauss Prize

The International Mathematical Union (IMU) and
the Deutsche Mathematiker-Vereinigung (DMV)
jointly award the

Carl Friedrich Gauss Prize for Applications of Mathematics

to **Professor Dr. Kiyoshi Itô**



for laying the **foundations of the Theory of Stochastic Differential Equations and Stochastic Analysis**. Itô's work has emerged as one of the major mathematical innovations of the 20th century and has found a wide range of applications outside of mathematics. **Itô calculus** has become a key tool in areas such as **engineering** (e.g., filtering, stability, and control in the presence of noise), **physics** (e.g., turbulence and conformal field theory), and **biology** (e.g., population dynamics). It is at present of particular importance in **economics** and finance with **option pricing** as a prime example.

Madrid, August 22, 2006

Sir John Ball
President of IMU



$$dS = \mu S dt + \sigma S \varepsilon \sqrt{dt}$$

Günter M. Ziegler
President of DMV



Martin
Grötschel

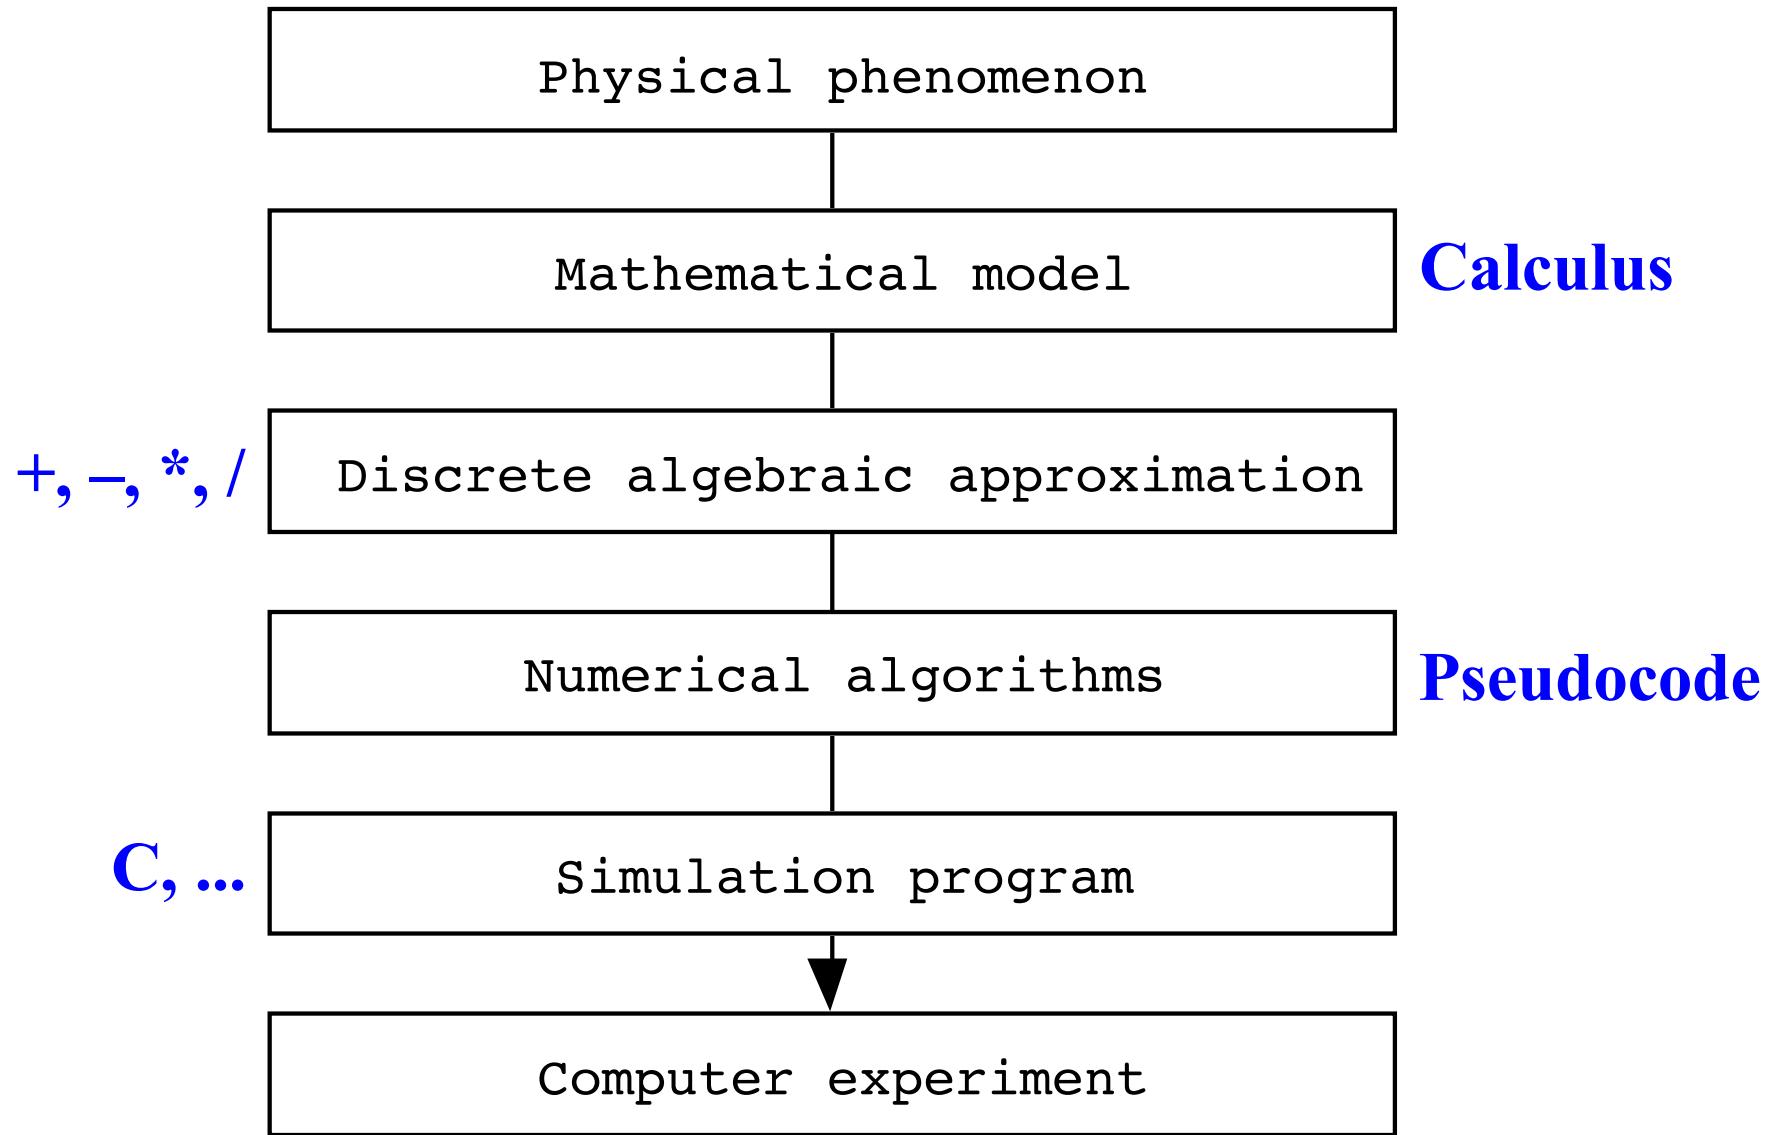
Monte Carlo Simulation



cf. Molecular dynamics

- Random trial → acceptance by a cost criterion
- Combinatorial optimization by simulated annealing

Recap: Scientific Computing



- Particle simulation as an archetypal example
- Use your own application in the final project


History of Particle Simulations

- '44 John von Neumann memo on a stored-program computer: "*Our present analytical methods seem unsuitable for the solution of the important problems arising in connection with nonlinear partial differential equations. The really efficient high-speed computing devices may provide us with those heuristic hints which are needed in all parts of mathematics for genuine progress*"
- '53 First Monte Carlo simulation of liquid by Metropolis, Rosenbluth, Rosenbluth, Teller, and Teller on MANIAC at Los Alamos Nat'l Lab
- '55 Enrico Fermi, John Pasta, and Stanislaw Ulam studied the dynamics of an one-dimensional array of particles coupled by anharmonic springs on MANIAC
- '56 Dynamics of hard spheres (billiards) studied by Alder and Wainwright at the Lawrence Livermore Nat'l Lab.
- '60 Radiation damage in crystalline Cu studied with short-range repulsion and uniform attraction toward the center by George Vineyard's group at Brookhaven Nat'l Lab
- '64 First MD simulation of liquid (864 argon atoms) using interatomic potentials by Aneesur Rahman at the Argonne Nat'l Lab on a CDC 3600

Moore's Law in Scientific Computing

Number of particles in MD simulations has doubled:

- Every 19 months in the past 50 years for classical MD
- Every 22 months in the past 30 years for DFT-MD



10¹²-atom MD & 10⁸-electron DFT on a petaflop/s computer with advances in algorithmic & parallel-computing techniques

ACM Best Dissertation Award



Association for
Computing Machinery

Advancing Computing as a Science & Profession

you are here: home → awards → doctoral dissertation award

ACM

myACM

Home

Membership

Publications

Digital Library

Special Interest Groups
(SIGs)

Online Books & Courses

Chapters

Conferences

Calendar of Events

Awards

Educational Activities

Public Policy

Online Communities

Computing News

Buy Proceedings & Videos

It's algorithm!

[site map](#) [accessibility](#) [contact](#)

[join](#) | [renew](#) | [subscribe](#)



Awards

[Home](#) | [ACM Awards](#) | [Nominations Process](#) | [Advanced Grades of Membership](#) | [Guide to Establishing an Award](#) | [Awards Committees](#) | [SIG Awards](#) |

1987 – Leslie Greengard

Series Winner (1987)

Citation

For his dissertation "The Rapid Evaluation of Potential Fields in Particle Systems."

$O(N^2) \rightarrow O(N)$



Top 10 Algorithms in History

- Metropolis Algorithm for Monte Carlo
- Simplex Method for Linear Programming
- Krylov Subspace Iteration Methods
- The Decompositional Approach to Matrix Computations
- The Fortran Optimizing Compiler
- QR Algorithm for Computing Eigenvalues
- Quicksort Algorithm for Sorting
- Fast Fourier Transform
- Integer Relation Detection
- Fast Multipole Method

IEEE CiSE, Jan/Feb ('00)

<http://awards.acm.org/doctoral%5Fdissertation>



History of Supercomputers

Early '40s: ENIAC by Presper Eckert & John Mauchly at Univ. of Pennsylvania—first general-purpose electronic computer

'76: Cray 1 by Seymour Cray—beginning of vector supercomputer era

Late 80's: massively parallel computers such as the Thinking Machines CM-2

('71: Intel 4004—invention of microprocessor)



ACM Best Dissertation Award




Association for
Computing Machinery

Advancing Computing as a Science & Profession

you are here: home → awards → doctoral dissertation award

[site map](#) [accessibility](#) [contact](#)



[join](#) | [renew](#) | [subscribe](#)

ACM

myACM



- [Home](#)
- [Membership](#)
- [Publications](#)
- [Digital Library](#)
- [Special Interest Groups \(SIGs\)](#)
- [Online Books & Courses](#)
- [Chapters](#)
- [Conferences](#)
- [Calendar of Events](#)
- [Awards](#)
- [Educational Activities](#)
- [Public Policy](#)
- [Online Communities](#)
- [Computing News](#)
- [Buy Proceedings & Videos](#)

Awards

[Home](#) | [ACM Awards](#) | [Nominations Process](#) | [Advanced Grades of Membership](#) | [Guide to Establishing an Award](#) | [Awards Committees](#) | [SIG Awards](#) |

1985 – Daniel Hillis

Series Winner (1985)

Citation

For his dissertation "The Connection Machine."

New Cancer Research Center to be Based at USC

By James Grant on October 26, 2009 7:55 AM

USC has been selected to establish a \$16 million cancer research center as part of a new strategy against the disease by the U.S. National Institutes of Health and its National Cancer Institute.

The five-year award will create a National Cancer Institute Physical Science-Oncology Center based at USC and involving a consortium of universities. Partnering in the USC grant will be Arizona State University, the California Institute of Technology, Cold Spring Harbor Laboratory, New York University, Stanford University, the University of Arizona and the University of Texas at Austin.

The Physical Science-Oncology Center initiative differs from past cancer research programs. While cancer biologists often work with scientists in other fields, this marks the first large-scale recruitment of outside scientists in the battle against the disease.



USC Viterbi School of Engineering professor and principal investigator W. Daniel Hillis

E-MAIL PRINT SHARE

Digress: Birth of FORTRAN Programming



N.S. Scott et al.,
Comput. Phys. Commun.
252, 107269 (2020)





Fig. 13. The range of programming languages used across the Program Library's 3224 published programs (1969–2016).



Math
vs.
Computing



Merge of PC & Supercomputers

Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,299,072	415,530.0	513,854.7	28,335
2	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148,600.0	200,794.9	10,096
3	Sierra - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	125,712.0	7,438
4	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371
5	Tianhe-2A - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000, NUDT National Super Computer Center in Guangzhou China	4,981,760	61,444.5	100,678.7	18,482

Theoretical performance
Measured performance
(in Tflop/s)

Flop/s =
floating-point
operations/second

M (mega) = 10^6
G (giga) = 10^9
T (Tera) = 10^{12}
P (Peta) = 10^{15}
X (Exa) = 10^{18}

<http://www.top500.org> (June '20)

- USC-HPC: 13,440 cores, 0.62 Pflop/s
- CACS: 4,096 cores
- CACS-INCITE: 200M core-hours/year on 280,320-core 11.7 Pflop/s Theta at Argonne Nat'l Lab



Supercomputing Tomorrow



Nation's first exaflop/s computer, Aurora A21

<http://aurora.alcf.anl.gov>

Delivery

CY 2021

Compute Node

2 Intel Xeon scalable "Sapphire Rapids" processors; 6 X^e arch-based GPUs; Unified Memory Architecture; 8 fabric endpoints

GPU Architecture

X^e arch-based "Ponte Vecchio" GPU Tile-based, chiplets, HBM stack, Foveros 3D integration, 7nm

On-Node Interconnect

CPU-GPU: PCIe
GPU-GPU: X^e Link

Aggregate System Memory

> 10 PB

System Interconnect

Cray Slingshot; Dragonfly topology with adaptive routing

Sustained Performance

≥ 1EF DP

High-End Computing at CACS

- Won two DOE supercomputing awards to develop & deploy metascalable (“design once, scale on future platforms”) simulation algorithms (2017-2020)



Innovative & Novel Computational Impact on Theory & Experiment

Title: “Petascale Simulations for Layered Materials Genome”

Principal Investigator:

Co-Investigator:



Aiichiro Nakano, University of Southern California

Priya Vashishta, University of Southern California

- NAQMD & RMD simulations on full 800K cores



786,432-core IBM Blue Gene/Q



Early Science Projects for Aurora


Supercomputer Announced

Metascalable layered materials genome

Investigator: Aiichiro Nakano, University of Southern California

- One of 10 initial simulation users of the next-generation DOE supercomputer

Exaflop/s = 10^{18} floating-point operations per second



Nation's first exaflop/s computer, Intel A21 (2021)

Your Platform: USC-CARC

Center for Advanced Research Computing (CARC)

- The CARC computing resource consists of 2 shared head nodes & a total of 30,000 CPU cores in 1,600 compute nodes. The typical compute node has dual 8 to 16 core processors accelerated by dual GPUs, residing on a 56 gigabit FDR InfiniBand backbone.
- **Discovery cluster (new):**
discovery.usc.edu
discovery2.usc.edu






Simulation Engines: NAQMD & RMD

Molecular Dynamics (MD)

Reactive MD (RMD)

Nonadiabatic quantum MD (NAQMD)





First principles-based reactive force-fields

- **Reactive bond order $\{BO_{ij}\}$**
→ Bond breakage & formation
- **Charge equilibration (QE_q) $\{q_i\}$**
→ Charge transfer

Tersoff, Brenner, Sinnott *et al.*; Streitz & Mintmire *et al.*;
van Duin & Goddard (ReaxFF)

Scalable Simulation Algorithm Suite



QMD (quantum molecular dynamics): DC-DFT


RMD (reactive molecular dynamics): F-ReaxFF

MD (molecular dynamics): MRMD

- 4.9 trillion-atom space-time multiresolution MD (MRMD) of SiO_2
 - 8.5 billion-atom fast reactive force-field (F-ReaxFF) RMD of RDX
 - 39.8 trillion grid points (50.3 million-atom) DC-DFT QMD of SiC
- parallel efficiency over 0.98 on 786,432 Blue Gene/Q cores

Nanobubble Collapse Near Silica Surface

- Billion-atom MD simulation of shock-induced nanobubble collapse in water near silica surface (67 million core-hours on 163,840 Blue Gene/P cores)




- Water nanojet formation and its collision with silica surface

A. Shekhar *et al.*, *Phys. Rev. Lett.* **111**, 184503 ('13)


Novel Nano-carbon Synthesis

112 million-atom reactive molecular dynamics (RMD) simulation of high-temperature oxidation of SiC nanoparticle on 786,432 IBM Blue Gene/Q cores



H₂ Production from Water Using LiAl Particles

16,661-atom quantum molecular dynamics (QMD) simulation
of Li₄₄₁Al₄₄₁ in water on 786,432 IBM Blue Gene/Q cores




Enabling Computer Science Technologies

What can computer science do to enable extreme-scale computational science?

A lot! That's what you will learn in CSCI 596:

- **Parallel computing**
- **Visualization**
- **Grid/cloud computing**




Parallel Computing

Parallel MD algorithm: Easy!

1. $\vec{v}_i(t + \frac{\Delta}{2}) \leftarrow \vec{v}_i(t) + \frac{\Delta}{2} \vec{a}_i(t)$
2. $\vec{r}_i(t + \Delta) \leftarrow \vec{r}_i(t) + \vec{v}_i(t + \frac{\Delta}{2}) \Delta$
3. atom_migrate()
4. atom_cache()
5. Compute $\vec{a}_i(t + \Delta)$ as a function of $\{\vec{r}_i(t + \Delta)\}$
6. $\vec{v}_i(t + \Delta) \leftarrow \vec{v}_i(t + \frac{\Delta}{2}) + \frac{\Delta}{2} \vec{a}_i(t + \Delta)$


Parallel Computing Hardware




- **Processor:** Executes arithmetic & logic operations
- **Memory:** Stores program & data (**stored program computer**)
- **Communication interface:** Performs signal conversion & synchronization between communication link & a computer
- **Communication link:** A wire capable of carrying a sequence of bits as electrical (or optical) signals

Parallel Molecular Dynamics


Spatial decomposition (short ranged): $O(N/P)$ computation



Atom caching: $O((N/P)^{2/3})$




Atom migration



MPI Programming


```
#include "mpi.h"
#include <stdio.h>
main(int argc, char *argv[ ]) {
    MPI_Status status;
    int myid;
    int n;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    if (myid == 0) {
        n = 777;
        MPI_Send(&n, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
    }
    else {
        MPI_Recv(&n, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
        printf("n = %d\n", n);
    }
    MPI_Finalize();
}
```



OpenMP Programming

```
#include <stdio.h>
#include <omp.h>
void main () {
    int nthreads,tid;
    nthreads = omp_get_num_threads();
    printf("Sequential section: # of threads = %d\n",nthreads);
    /* Fork multi-threads with own copies of variable */
    #pragma omp parallel private(tid)
    {
        /* Obtain & print thread id */
        tid = omp_get_thread_num();
        printf("Parallel section: Hello world from thread %d\n",tid);
        /* Only master thread does this */
        if (tid == 0) {
            nthreads = omp_get_num_threads();
            printf("Parallel section: # of threads = %d\n",nthreads);}
        } /* All created threads terminate */
    }
```

parallel section



- Obtain the number of threads & my thread ID
- By default, all variables are shared unless selectively changing storage attributes using private clauses

GPU Programming: CUDA

- Compute Unified Device Architecture
- Integrated host (CPU) + device (GPU) application programming interface based on C language developed at NVIDIA
- CUDA homepage

http://www.nvidia.com/object/cuda_home_new.html


- Compilation

\$ nvcc pi.cu


- Execution

\$ a.out

PI = 3.141593



Immersive & Interactive Visualization




Basics:

- OpenGL programming/CAVE library

Billion-atom walkthrough:

- Octree-based view frustum culling
- Probabilistic occlusion culling
- Multiresolution rendering
- Parallel & distributed processing



OpenGL Programming




```
...
#include <OpenGL/gl.h>
#include <OpenGL/glu.h>
#include <GLUT/glut.h>
...
/* Set a glut callback function */
glutDisplayFunc(display);

/* generate an OpenGL display list for single sphere */
sphereid = glGenLists(1);
makeFastNiceSphere(sphereid,atom_radius);

/* generate an OpenGL display list for the atoms' geometry */
atomsid = glGenLists(1);
/* make the geometry of the current frame's atoms */
makeCurframeGeom();

/* Start main display loop */
glutMainLoop();
```

3D Is Back in Hollywood (and Home)




<https://github.com/USCCACS/GEARs>




Oculus Rift & HTC Vive head mounted displays

Scientific Mixed Reality



VR



MR




Microsoft Hololens academic seeding program at USC

cf. CSCI 538: Augmented, Virtual and Mixed Reality

Grid Computing

- **World Wide Web:** Universal interface to digital library on the Internet
- **Information Grid:** Pervasive (from any place in the world at any time) access to everything (computing, mass storage, experimental equipments, distributed sensors, etc., on high-speed networks)




Grid Computing

- **Globus middleware**
 - > **Resource monitoring, discovery, & management**
 - > **Security**
 - > ...
- http://www.globus.org
- **Globus-enabled MPI**
http://www.globus.org/grid_software/computation/mpich-g2.php


Global Collaborative Simulation

Multiscale MD/QM simulation on a Grid of distributed PC clusters in the US & Japan

- Task decomposition (MPI Communicator) + spatial decomposition
- MPICH-G2/Globus



Japan: Yamaguchi—65 P4 2.0GHz
Hiroshima, Okayama, Niigata—3×24 P4 1.8GHz
US: Louisiana—17 Athlon XP 1900+



MD — 91,256 atoms
QM (DFT) — $76n$ atoms on n clusters

- Scaled speedup, $P = 1$ (for MD) + $8n$ (for QM)
- Efficiency = 94.0% on 25 processors over 3 PC clusters

Kikuchi et al.
IEEE/ACM SC02

Multiscale Modeling

The Nobel Prize in Chemistry 2013




© Nobel Media AB
Martin Karplus



Photo: Keilana via
Wikimedia Commons
Michael Levitt



Photo: Wikimedia
Commons
Arieh Warshel



QM/MM:
quantum-
mechanical/molecular-
mechanical modeling



The Nobel Prize in Chemistry 2013 was awarded jointly to Martin Karplus, Michael Levitt and Arieh Warshel *"for the development of multiscale models for complex chemical systems"*.

- A. Warshel & M. Karplus, *J. Am. Chem. Soc.* **94**, 5612 ('72)
A. Warshel & M. Levitt, *J. Mol. Biol.* **103**, 227 ('76)

Sustainable Grid Supercomputing

- Sustained ($>$ months) supercomputing ($> 10^3$ CPUs) on a Grid of geographically distributed supercomputers
- Hybrid Grid remote procedure call (GridRPC) + message passing (MPI) programming
- Dynamic allocation of computing resources on demand & automated migration due to reservation schedule & faults

Ninf-G GridRPC: ninf.apgrid.org; MPICH: www.mcs.anl.gov/mpi




Takemiya et al., IEEE/ACM SC06
Song et al., IJCS ('09)

Multiscale QM/MD simulation of high-energy beam oxidation of Si



Above the Clouds: A Berkeley View of Cloud Computing

*Michael Armbrust
Armando Fox
Rean Griffith
Anthony D. Joseph
Randy H. Katz
Andrew Konwinski
Gunho Lee
David A. Patterson
Ariel Rabkin
Ion Stoica
Matei Zaharia*




Electrical Engineering and Computer Sciences
University of California at Berkeley

Technical Report No. UCB/EECS-2009-28
<http://www.eecs.berkeley.edu/Pubs/TechRpts/2009/EECS-2009-28.html>

February 10, 2009

MapReduce

- Parallel programming model for data-intensive applications on large clusters
 - > User implements **Map()** and **Reduce()**
- Parallel computing framework
 - > Libraries take care of everything else
 - Parallelization
 - Fault tolerance
 - Data distribution
 - Load balancing
- Developed at Google




Cybermanufacturing

NSF 2036359 FMRG: Artificial Intelligence Driven Cybermanufacturing of Quantum Material Architectures

9/1/2020 – 8/31/2025

R. Nagpal (*Harvard*); R. Kalia, A. Nakano, H. Wang (*USC*); D. Rawat (*Howard*)



This project develops a transformative future manufacturing platform for quantum material architectures using a cybermanufacturing approach, which combines artificial intelligence, robotics, multiscale modeling, and predictive simulation for the automated & parallel assembly of multiple two-dimensional materials into complex three-dimensional structures.

National AI & Quantum Initiative

THE WHITE HOUSE
Office of Science and Technology Policy

FOR IMMEDIATE RELEASE

August 26, 2020

NEWS | Trump Administration Announces \$1 Billion in AI and Quantum Research Institutes Nationwide

Today, the White House Office of Science and Technology Policy, the National Science Foundation (NSF), and the U.S. Department of Energy (DOE) announced over \$1 billion for the establishment of 12 new artificial intelligence (AI) and quantum information science (QIS) research institutes nationwide.

The \$1 billion will go towards seven NSF-led AI Research Institutes and five DOE QIS Research Centers over five years, establishing a total of 12 multi-disciplinary and multi-institutional national hubs for research and workforce development in these critical emerging technologies. Together, the institutes will spur cutting edge innovation, support regional economic growth, and advance American leadership in these critical industries of the future.

“Thanks to the leadership of President Trump, the United States is accomplishing yet another milestone in our efforts to strengthen research in AI and quantum. We are proud to announce that over \$1 billion in funding will be geared towards that research, a defining achievement as we continue to shape and prepare this great Nation for excellence in the industries of the future,” said Advisor to the President Ivanka Trump.

Related Courses

- CSCI 503—Parallel Programming
 - > Parallel programming paradigms
 - > Parallel computing architectures
 - > Hands-on parallel programming assignments

- EE 451—Parallel & Distributed Computation
 - > Parallel programming techniques
 - > Performance optimization
 - > Application mapping to multicore, accelerator & cloud

See previous syllabus by Prof. Viktor Prasanna:

<http://cacs.usc.edu/education/cs596/EE451-Prasanna.pdf>

Computational Research Survival Guide

Talk to user
Regular expression (text processing)

<https://regexone.com>

CSCI 596: Talk to (C) compiler via libraries & directives
Parallel processing: MPI, OpenMP, CUDA, ...
Visualization: OpenGL

Talk to operating system
Shell programming

A. Scopatz and K. D. Huff, *Effective Computation in Physics* (O'Reilly, '15)
<http://www.amazon.com/Effective-Computation-Physics-Anthony-Scopatz/dp/1491901535>