

Message Passing Interface (MPI) Programming

Aiichiro Nakano

*Department of Computer Science
Department of Physics & Astronomy
Department of Quantitative & Computational Biology
University of Southern California*

Email: anakano@usc.edu



MPI: Standard parallel programming language

— Default parallel-programming language



Preparation

Minimal knowledge required for the hands-on projects in this course:

- Able to log in & use the Discovery computing cluster at USC Center for Advanced Research Computing (CARC) at the level of its “getting started” tutorial—logging in; transferring files; installing software; running jobs,...:

<https://www.carc.usc.edu/user-guides/hpc-systems/discovery/getting-started-discovery>

- Use shell commands to interact with the operating system at the level of “Chapter 1—Introduction to the Command Line” of *Effective Computation in Physics* by Scopatz and Huff; USC students have free access to the book through Safari Online:

<https://libraries.usc.edu/databases/safari-books>

Chapter 1. Introduction to the Command Line

The command line, or *shell*, provides a powerful, transparent interface between the user and the internals of a computer. At least on a Linux or Unix computer, the command line provides total access to the files and processes defining the state of the computer—including the files and processes of the operating system.

- Need to log in from USC secure network or USC VPN if off campus

<https://itservices.usc.edu/vpn>

How to Use USC CARC Cluster

System: Intel/AMD-based computing cluster

<https://carc.usc.edu>

Log in

```
> ssh anakano@discovery.usc.edu
```

Alternatively, you can use `discovery2.usc.edu`

To use MPI library:

If using Bash shell, add these in .bashrc

```
module purge
module load usc
```

Use text editor like vim, nano, emacs
To set up standard software environment

Compile an MPI program

```
> mpicc -o mpi_simple mpi_simple.c
```

Execute an MPI program


```
> mpirun -n 2 mpi_simple
```

```
[anakano@discovery ~]$ which mpicc
/spack/2206/apps/linux-centos7-x86_64_v3/gcc-11.3.0/openmpi-4.1.4-4w23jca/bin/mpicc
[anakano@discovery ~]$ more /proc/cpuinfo
```

To find absolute path to mpicc command

To find processor information

Email carc-support@usc.edu for assistance



Shell is a language you speak with the operating system

Type `echo $0` to find which shell you are using

VPN Issue

- **It is now required to use VPN (virtual private network) to access Discovery from off-campus:**

<https://itservices.usc.edu/vpn>

- **Cisco AnyConnect software for VPN on Mac may have a DNS (domain name system) problem, which could be bypassed using IP addresses instead of login server names (note discovery.usc.edu is a generic name for the two login servers, discovery1 and discovery2)**

discovery1.usc.edu: 10.72.0.13

discovery2.usc.edu: 10.72.0.14

Submit a Slurm Batch Job

Prepare a script file, mpi_simple.sl

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=2
#SBATCH --time=00:00:10
#SBATCH --output=mpi_simple.out
#SBATCH -A anakano_429
mpirun -n $SLURM_NTASKS ./mpi_simple
```

Slurm (Simple Linux Utility for Resource Management): Open-source job scheduler that allocates compute resources on clusters for queued jobs

Class project account; type `myaccount` to check all accounts

Total number of processes = `ntasks-per-node` × `nodes`

Submit a Slurm job

```
discovery: sbatch mpi_simple.sl
```

```
Submitted batch job 63695
```

Check the status of a Slurm job

```
discovery: squeue -u anakano
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
63695	main	mpi_simple	anakano	PD	0:00	1	(Resources)

Cancel a Slurm job

```
discovery: scancel 63695
```


Check the output

```
discovery: more mpi_simple.out
```

```
n = 777
```

For detailed explanation, see the lecture note

<https://aiichironakano.github.io/cs596/02MPI.pdf>



Interactive Job at CARC

When debugging your MPI program, you may want to access computing nodes interactively, so that you can edit, compile & run MPI program in real time unlike the batch job

Reserve 2 processors for 20 minutes

```
[anakano@discovery cs596]$ salloc -n 2 -t 20
salloc: Granted job allocation 63754
salloc: Waiting for resource configuration
salloc: Nodes d05-05 are ready for job
[anakano@d05-05 cs596]$ mpirun -n 2 ./mpi_simple
n = 777
[anakano@d05-05 cs596]$ exit
exit
salloc: Relinquishing job allocation 63754
[anakano@discovery cs596]$
```

Note you are now using a computing node named d05-05

Back to the login node

Type `less /proc/cpuinfo` to find what kind of node you got

Symbolic Link to Work Directory

- Your home directory has small (but enough for assignments) quota (type `myquota` to confirm), so use the scratch file system (`/scratch1/anakano` for user `anakano`) if needed
- It is convenient to make a symbolic link to a directory you use often, rather than typing its long absolute path every time

```
[anakano@discovery ~]$ ln -s /scratch1/anakano/cs596 cs596
[anakano@discovery ~]$ ls -lt
total 81985
lrwxrwx--- 1 anakano anakano 22 Aug 23 12:14 cs596 -> /scratch1/anakano/cs596
drwxrwx--- 3 anakano anakano  1 Aug 20 10:07 FFTW
lrwxrwx--- 1 anakano anakano 16 Aug 14 15:48 scr -> /scratch1/anakano
...
[anakano@discovery ~]$ cd cs596
[anakano@discovery cs596]$ pwd -P
/scratch1/anakano/cs596
```

Instead of typing
`cd /scratch1/anakano/cs596`

Print physical working directory


File Transfer

- Use secure file transfer protocol to transfer files between your laptop and Discovery

```
macbook-pro $ sftp anakano@discovery.usc.edu
Connected to discovery.usc.edu.
sftp> cd cs596
sftp> put md.*  Transfer files from local computer (your laptop) to remote computer (Discovery)
sftp> ls  Check whether the files have been transferred
md.c      md.h      md.in
sftp> exit
macbook-pro $
```


- To transfer files from remote computer to local computer, use **get** instead

Parallel Computing Hardware



- **Processor:** Executes arithmetic & logic operations.
- **Memory:** Stores program & data.
- **Communication interface:** Performs signal conversion & synchronization between communication link and a computer.
- **Communication link:** A wire capable of carrying a sequence of bits as electrical (or optical) signals.

Motherboard




Key Features

1. Dual Intel® Xeon™ EM64T Support up to 3.60 GHz
2. Intel® E7525 (Tumwater) Chipset
3. Up to 16GB DDRII-400 SDRAM
4. Intel® 82546GB Dual-port Gigabit Ethernet Controller
5. Adaptec AIC-7902 Dual Channel Ultra320 SCSI
6. 2x SATA Ports via ICH5R SATA Controller
7. 1 (x16) & 1 (x4) PCI-Express, 1 x 64-bit 133MHz PCI-X, 2 x 64-bit 100MHz PCI-X, 1 x 32-bit 33MHz PCI Slots
8. Zero Channel RAID Support
9. AC'97 Audio, 6-Channel Sound

Supermicro X6DA8-G2


Communication Network

**Mesh
(torus)**




NEC Earth Simulator (640x640 crossbar)

IBM Blue Gene/Q (5D torus)



**Crossbar
switch**



See [Grama](#),
Secs. 2.4.2-2.4.4

Message Passing Interface

MPI (Message Passing Interface)

A standard message passing system that enables us to write & run applications on parallel computers

Download for Unix & Windows:

<http://www.mcs.anl.gov/mpi/mpich>

Compile

```
> mpicc -o mpi_simple mpi_simple.c
```

Run

```
> mpirun -np 2 mpi_simple
```

MPI Programming

mpi_simple.c: Point-to-point message send & receive

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {
    MPI_Status status;
    int myid;
    int n;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    if (myid == 0) {
        n = 777;
        MPI_Send(&n, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
    }
    else {
        MPI_Recv(&n, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
        printf("n = %d\n", n);
    }
    MPI_Finalize();
    return 0;
}
```

MPI rank

Matching message tags


Data triplet

To/from whom

```
graph TD
    P0((P0)) -- "send to 1" --> MPI_daemon([MPI daemon])
    P1((P1)) -- "recv from 0" --> MPI_daemon
    MPI_daemon -- "requests" --> P1
```


cf. p. 3 of <https://aiichironakano.github.io/cs596/02MPI.pdf>

Single Program Multiple Data (SPMD)



Single Program Multiple Data (SPMD)

What really happens?



MPI Minimal Essentials

We only need `MPI_Send()` & `MPI_Recv()`
within `MPI_COMM_WORLD`

```
MPI_Send(&n, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
```

```
MPI_Recv(&n, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
```

└──────────┘

Data triplet

└──────────────────┘

To/from whom

└──┘


Information

Global Operation

All-to-all reduction: Each process contributes a partial value to obtain the global summation. In the end, all the processes will receive the calculated global sum.


```
MPI_Allreduce(&local_value, &global_sum, 1, MPI_INT, MPI_SUM,  
MPI_COMM_WORLD)
```

```
int l_v, g_s; // local variable & global sum  
l_v = myid;   // myid is my MPI rank  
MPI_Allreduce(&l_v, &g_s, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
```



Hypercube Algorithm

Hypercube algorithm: Communication of a reduction operation is structured as a series of pairwise exchanges, one with each neighbor in a hypercube (**butterfly**) structure. Allows a computation requiring all-to-all communication among p processes to be performed in $\log_2 p$ steps.




Butterfly network

$$\begin{aligned}
 & a000 + a001 + a010 + a011 + a100 + a101 + a110 + a111 \\
 = & ((a000 + a001) + (a010 + a011)) \\
 + & ((a100 + a101) + (a110 + a111)) \\
 \textcircled{2} & \qquad \qquad \textcircled{1} \qquad \qquad \textcircled{0}
 \end{aligned}$$

Barrier

```
<A>;  
barrier();  
<B>;
```



MPI_Barrier(MPI_Comm communicator)

Useful for debugging (but would slow down the program)

MPI Communication

MPI communication functions:

1. Point-to-point

`MPI_Send()`

`MPI_Recv()`

2. Global

`MPI_Allreduce()`

`MPI_Barrier()`


`MPI_Bcast()`

Hypercube Template

```

procedure hypercube(myid, input, log2P, output)
begin
    myid ∈ [0, P - 1]
    mydone := input;
    for l := 0 to log2P-1 do
    begin
        partner := myid XOR 2l;
        send mydone to partner;
        receive hisdone from partner;
        mydone = mydone OP hisdone
    end
    output := mydone
end

```



level	2 ^l	bitvalue
0	001	
1	010	
2	100	

Exclusive OR

a	b	a XOR b
0	0	0
0	1	1
1	0	1
1	1	0

Associative operator (e.g., sum, max)

$$(a \text{ OP } b) \text{ OP } c = a \text{ OP } (b \text{ OP } c)$$

$$abcdefg \text{ XOR } 0000100 = abcd\bar{e}fg$$

In C, ^ (caret operator) is bitwise XOR applied to int

Driver for Hypercube Test

```
#include "mpi.h"
#include <stdio.h>
int nprocs; // Number of processes
int myid;   // My rank

double global_sum(double partial) {
    /* Implement your own global summation here */
}

int main(int argc, char *argv[]) {
    double partial, sum, avg;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    partial = (double) myid;

    printf("Rank %d has %le\n", myid, partial);
    sum = global_sum(partial);
    if (myid == 0) {
        avg = sum/nprocs;
        printf("Global average = %le\n", avg);
    }
    MPI_Finalize();
    return 0;
}
```

Global variables are visible in both *global_sum()* & *main()*

Who am I?

How big is the world? (see p. 5 in [MPI lecture note](#))

C Implementation of global_sum()


```
double mydone = partial;
for (int bitvalue=1; bitvalue<nprocs; bitvalue*=2)
{
    int partner = myid ^ bitvalue;
    send mydone to partner;
    receive hisdone from partner;
    mydone = mydone + hisdone;
}
return mydone;
```

Multiplied by 2

level	2^l	bitvalue
0	001	1
1	010	2
2	100	4


Implement with MPI_Send() & MPI_Recv()

Use *bitvalue* as counter & bitmask



It is recommended to use distinct labels (tags) for different messages, e.g., bitmask (= stride) as a tag

MPI Send & Receive Revisited



MPI_Datatype	C data type
MPI_CHAR	char
MPI_INT	int
MPI_FLOAT	float
MPI_DOUBLE	double
...	...

MPI_Status status;
 Filled with information about the received message
 status.MPI_SOURCE Source process rank
 status.MPI_TAG Tag of the received message

- Only tag-matching message passing between matching source/destination pair of ranks take place
- It is recommended to use distinct tags for different messages to avoid accidental receipt of unintended messages

Sample Slurm Script

Run two MPI runs in a single Slurm job

```
#!/bin/bash
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=4
#SBATCH --time=00:00:59
#SBATCH --output=global.out
#SBATCH -A anakano_429      mpicc -o global_avg global_avg.c

mpirun -n $SLURM_NTASKS ./global_avg
mpirun -n 4 ./global_avg
```

Total number of processors
= ntasks-per-node (4) × nodes (2) = 8

- **Type** `sbatch global_avg.sl` **in the directory where the executable** `global_avg` **resides, or** `cd` **(change directory) to where it is**

Output of global.c

- **4-processor job**

```
Rank 0 has 0.000000e+00
Rank 1 has 1.000000e+00
Rank 2 has 2.000000e+00
Rank 3 has 3.000000e+00
Global average = 1.500000e+00
```

- **8-processor job**


```
Rank 0 has 0.000000e+00
Rank 1 has 1.000000e+00
Rank 2 has 2.000000e+00
Rank 3 has 3.000000e+00
Rank 5 has 5.000000e+00
Rank 6 has 6.000000e+00
Rank 4 has 4.000000e+00
Rank 7 has 7.000000e+00
Global average = 3.500000e+00
```

**Actual output is
random order in ranks
— Why?**

References on Hypercube Algorithms

1. [https://en.wikipedia.org/wiki/Hypercube_\(communication_pattern\)](https://en.wikipedia.org/wiki/Hypercube_(communication_pattern))
2. I. Foster, *Designing and Building Parallel Programs* (Addison-Wesley, 1995) Chap. 11 — Hypercube algorithms:
<https://www.mcs.anl.gov/~itf/dbpp/text/node123.html>

Distributed-Memory Parallel Computing



Communicator

mpi_comm.c: Communicator = process group + context

```
#include "mpi.h"
#include <stdio.h>
#define N 64
int main(int argc, char *argv[]) {
    MPI_Comm world, workers;
    MPI_Group world_group, worker_group;
    int myid, nprocs;
    int server, n = -1, ranks[1];
    MPI_Init(&argc, &argv);
    world = MPI_COMM_WORLD;
    MPI_Comm_rank(world, &myid); // My rank in the world (see next slide)
    MPI_Comm_size(world, &nprocs); // How big is the world?
    server = nprocs-1; // The last guy becomes the server
    MPI_Comm_group(world, &world_group);
    ranks[0] = server; // Note the rest will become workers
    MPI_Group_excl(world_group, 1, ranks, &worker_group);
    MPI_Comm_create(world, worker_group, &workers);
    MPI_Group_free(&worker_group); // Release resources no longer needed
    if (myid != server). // All, except for the server, are workers
        MPI_Allreduce(&myid, &n, 1, MPI_INT, MPI_SUM, workers);
    printf("process %2d: n = %6d\n", myid, n);
    MPI_Comm_free(&workers);
    MPI_Finalize();
    return 0;
}
```

Usage

- Avoid accidental match of unintended Send-Receive pairs
- Global operations in a subgroup of processes

Code at <https://aiichironakano.github.io/cs596/src/mpi/>
For detail, see p. 4 in <https://aiichironakano.github.io/cs596/02MPI.pdf>

Example: Ranks in Different Groups

World Rank	Institution*	Country /Region	National Rank	Total Score	Score on Alumni ▾
1	Harvard University		1	100	100
2	Stanford University		2	72.1	41.8
3	Massachusetts Institute of Technology (MIT)		3	70.5	68.4
4	University of California-Berkeley		4	70.1	66.8
5	University of Cambridge		1	69.2	79.1
51	University of Southern California		33	31	31.7

`MPI_Comm_rank(world, &usc_world);`
`MPI_Comm_rank(us, &usc_national);`

Rank is relative in each communicator!


Output from mpi_comm.c

Slurm script

```
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=2
...
mpirun -n $SLURM_NTASKS ./mpi_comm
```


```
process 3: n = -1
process 0: n = 3
process 1: n = 3
process 2: n = 3
```

**What Has
Happened?**



Grid Computing & Communicators

H. Kikuchi *et al.*, "Collaborative simulation Grid: multiscale quantum-mechanical/classical atomistic simulations on distributed PC clusters in the US & Japan, *IEEE/ACM SC02*




Communicator = a nice migration path to distributed computing

- Single MPI program run with the Grid-enabled MPI implementation, MPICH-G2
- Processes are grouped into MD & QM groups by defining multiple MPI communicators as subsets of MPI_COMM_WORLD; a machine file assigns globally distributed processors to the MPI processes

Global Grid QM/MD

- *One of the largest (153,600 cpu-hrs) sustained Grid supercomputing* at 6 sites in the US (USC, Pittsburgh, Illinois) & Japan (AIST, U Tokyo, Tokyo IT)





Takemiya *et al.*, "Sustainable adaptive Grid supercomputing: multiscale simulation of semiconductor processing across the Pacific," *IEEE/ACM SC06*

Sustainable Grid Supercomputing

- Sustained (> months) supercomputing (> 10^3 CPUs) on a Grid of geographically distributed supercomputers
- Hybrid Grid remote procedure call (GridRPC) + message passing (MPI) programming
- Dynamic allocation of computing resources on demand & automated migration due to reservation schedule & faults



Ninf-G GridRPC: ninf.apgrid.org; MPICH: www.mcs.anl.gov/mpi



Multiscale QM/MD simulation of high-energy beam oxidation of Si

Computation-Communication Overlap

H. Kikuchi *et al.*, "Collaborative simulation Grid: multiscale quantum-mechanical/classical atomistic simulations on distributed PC clusters in the US & Japan, *IEEE/ACM SC02*



$$\frac{\text{Earth's circumference}}{\text{Light speed}} = \frac{40,000 \text{ [km]} = 4 \times 10^7 \text{ [m]}}{3 \times 10^8 \text{ [m/s]}} = 0.1 \text{ s} = 100 \text{ ms}$$

Try on Discovery:

`tracert www.u-tokyo.ac.jp`
vs. `ping hpc-transfer.usc.edu`

- **How to overcome 200 ms latency & 1 Mbps bandwidth?**
- **Computation-communication overlap:** To hide the latency, the communications between the MD & QM processors have been overlapped with the computations using **asynchronous messages**

Synchronous Message Passing

MPI_Send () : (blocking), synchronous


- Safe to modify original data immediately on return
- Depending on implementation, it may return whether or not a matching receive has been posted, or it may block (especially if no buffer space available)

MPI_Recv () : blocking, synchronous

- Blocks for message to arrive
- Safe to use data on return



Experienced a lot of blocking on iPSC/860 with 12 MB user & 4 MB system memory per node



Asynchronous Message Passing

Allows computation-communication overlap

MPI_Isend() : non-blocking, asynchronous

- Returns immediately whether or not a matching receive has been posted
- Not safe to modify original data immediately (use **MPI_Wait()** system call)

MPI_Irecv() : non-blocking, asynchronous

- Does not block for message to arrive
- Cannot use data before checking for completion with **MPI_Wait()**

```
A...;  
MPI_Isend();  
B...;  
MPI_Wait();  
C...; // Reuse the send buffer
```

MPI_Irecv() is just a “request” for data delivery, when a matching message arrives


```
A...;  
MPI_Irecv();  
B...;  
MPI_Wait();  
C...; // Use the received message
```

See p. 6 in <https://aiichironakano.github.io/cs596/02MPI.pdf>

Program irecv_mpi.c

```
#include "mpi.h"
#include <stdio.h>
#define N 1000
int main(int argc, char *argv[]) {
    MPI_Status status;
    MPI_Request request;
    int send_buf[N], recv_buf[N];
    int send_sum = 0, recv_sum = 0;
    long myid, left, Nnode, msg_id, i;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    MPI_Comm_size(MPI_COMM_WORLD, &Nnode);
    left = (myid + Nnode - 1) % Nnode;
    for (i=0; i<N; i++) send_buf[i] = myid*N + i; // Compose a big message
    MPI_Irecv(recv_buf, N, MPI_INT, MPI_ANY_SOURCE, 777, MPI_COMM_WORLD,
              &request); // Post a receive
    /* Perform tasks that don't use recv_buf */
    MPI_Send(send_buf, N, MPI_INT, left, 777, MPI_COMM_WORLD);
    for (i=0; i<N; i++) send_sum += send_buf[i];
    MPI_Wait(&request, &status); // Complete the receive
    /* Now it's safe to use recv_buf */
    for (i=0; i<N; i++) recv_sum += recv_buf[i];
    printf("Node %d: Send %d Recv %d\n", myid, send_sum, recv_sum);
    MPI_Finalize();
    return 0;
}
```


Wrap-around/torus
via modulo (%) operator
(cf. periodic boundary condition)



Code at <https://aiichironakano.github.io/cs596/src/mpi/>

Output from `irecv_mpi.c`

```
Node 1: Send 1499500 Recv 2499500
Node 3: Send 3499500 Recv 499500
Node 0: Send 499500 Recv 1499500
Node 2: Send 2499500 Recv 3499500
```




Multiple Asynchronous Messages

```
MPI_Request requests[N_message];
MPI_Status statuses[N_message];
MPI_Status status;
int index;

/* Wait for all messages to complete */
MPI_Waitall(N_message, requests, statuses);

/* Wait for any specified messages to complete */
MPI_Waitany(N_message, requests, &index, &status);
```



returns the index ($\in [0, N_message-1]$) of the message that completed

Polling MPI_Irecv


```
int flag;

/* Post an asynchronous receive */
MPI_Irecv(recv_buf, N, MPI_INT, MPI_ANY_SOURCE, 777,
          MPI_COMM_WORLD, &request);


/* Perform tasks that don't use recv_buf */
...

/* Polling */
MPI_Test(&request, &flag, &status); // Check completion
if (flag) { // True if message received
    /* Now it's safe to use recv_buf */
    ...
}
```


Where to Go from Here



Basic MPI



Advanced MPI, including MPI-3

- **Complete MPI reference at** <http://www.netlib.org/utk/papers/mpi-book/mpi-book.html>
- **MPI is evolving (MPI-2 to MPI-5) to include advanced features like remote memory access (`MPI_Put()` & `MPI_Get()`; cf. `sftp`), parallel I/O and dynamic process management**
- **Various versions of MPI standard are specified at** <https://www.mpi-forum.org/docs/>

MPI Basics: Recap

- **Parallel computing = Who does what**
- **Single program multiple data (SPMD) programming: Do it with MPI rank (who am I) & selection constructs (`if`, *etc.*)**
- **Only need `MPI_Send()` & `MPI_Recv()` within communicators to implement any distributed-memory parallel computing**
- **Asynchronous message passing (`MPI_Isend()` & `MPI_Irecv()`) to overlap computation & communication**
- **You can survive professionally only with a few global communication functions, *e.g.*, `MPI_Allreduce()`, `MPI_Barrier()` & `MPI_Bcast()`**

Start using MPI for your research & projects!

20 Years-Unleashing the Power of HPC

SC2001

2001 Chair
Charles Slocumb
Denver, CO



2001

Notable Systems first mentioned this year in the proceedings:

- SGI Origin 3000
- Sun Fire 6000
- ASCI White
- Blue Horizon
- ASCI Blue Mountain

Notable Processors:

- MIPS R 12000
- Intel Pentium 4
- Intel Itanium

Noteworthy Architecture Topics:

- Cache coherence through snooping
- Application speedups through custom on-the-fly FPGA function units
- Interactive program steering
- Grid-enabled parallel computing

Notable Programming Languages:

- HDL
- PThreads

Research Machines:


- CPlant




A WINE-2 system board



A discrete particle simulation of 1.5 billion atoms



Adaptive mesh simulation of advecting sinusoidal density contours




Adaptive mesh simulation of star formation



The MDM system


Adaptive mesh simulation of a spherical shock



Solution of a three body quantum mechanics problem

Then, only (basic) MPI

It's not how fancy programming constructs you use, but a compelling application instead



Best Paper	Aiichiro Nakano, Rajiv K. Kalia, Priya Vashishta, Timothy J. Campbell, Shuji Ogata, Fuyuki Shimojo, and Subhash Saini Scalable atomistic simulation algorithms for materials research
Best Student Paper	Shava Smullen, Henri Cazsanova and Francine Berman Applying Scheduling and Tuning to On-line Parallel Tomography
ACM Gordon Bell Prize	See list of ACM Gordon Bell Prize winners
Best Research Poster	Sumir Chandra, Johan Steensland, and Manish Parashar ??? If you know, please contact chair@SIGHPC.org


Multi-Exaflop/s Light-Matter Dynamics

2025 Gordon Bell prize finalist

Argonne's Aurora supercomputer drives groundbreaking simulations to explore how light shapes quantum materials

A USC-led team uses exascale computing power and AI to simulate light-matter interactions at an unprecedented scale

The team's work has been named [a finalist for the Association of Computing Machinery's Gordon Bell Prize](#), one of the top honors in high-performance computing.



- **60,000 GPUs & 20,000 CPUs (on 10,000 nodes) of Aurora: 1.87 Exaflop/s; 152× & 3,780× improvements of time-to-solution over state-of-the-art for 15.4 million-electron DC-MESH NAQMD & 1.2 trillion-atom excited XS-NNQMD; nearly perfect parallel efficiency**

DC-MESH: Divide-&-conquer Maxwell-Ehrenfest-surface hopping

NAQMD: Nonadiabatic quantum molecular dynamics

NNQMD: Neural-network quantum molecular dynamics

- **Open standard programming: Portable across supercomputers through MPI (message passing) + OpenMP (multithreading) + OpenMP target (GPU) + PyTorch (NNQMD)**



Now, (still basic) MPI + OpenMP

Will learn soon

T. M. Razakh *et al.*, *Proc. SC* ('25) doi: 10.1145/3712285.3771785