Neutral Territory Decomposition for Parallel MD

Aiichiro Nakano

Collaboratory for Advanced Computing & Simulations
Department of Computer Science
Department of Physics & Astronomy
Department of Quantitative & Computational Biology
University of Southern California

Email: anakano@usc.edu

D. E. Shaw, A fast, scalable method for the parallel evaluation of distance-limited pairwise particle interactions, *J. Comput. Chem*. **26**, 1318 ('05)

King of Quant
https://www.youtube.com/watch?v=U6rNTQtDhA4&t=3s

Fine Granularity

Number of atoms per process (N/P) ~ 1

***cf.* Biomolecular simulations**

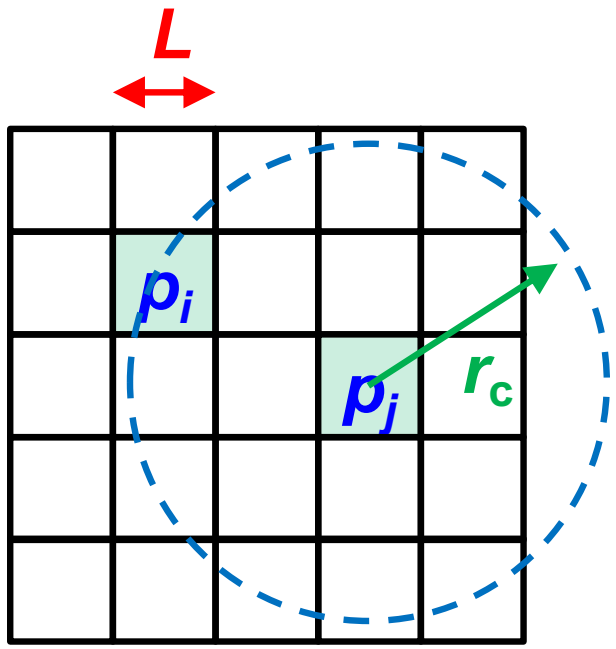

spatial subsystem length (L) ≪ interaction cutoff (r_c)

Spatial (Half-Shell) *vs.* NT Decompositions

Locus of interaction — who does what (2-dimensional example)

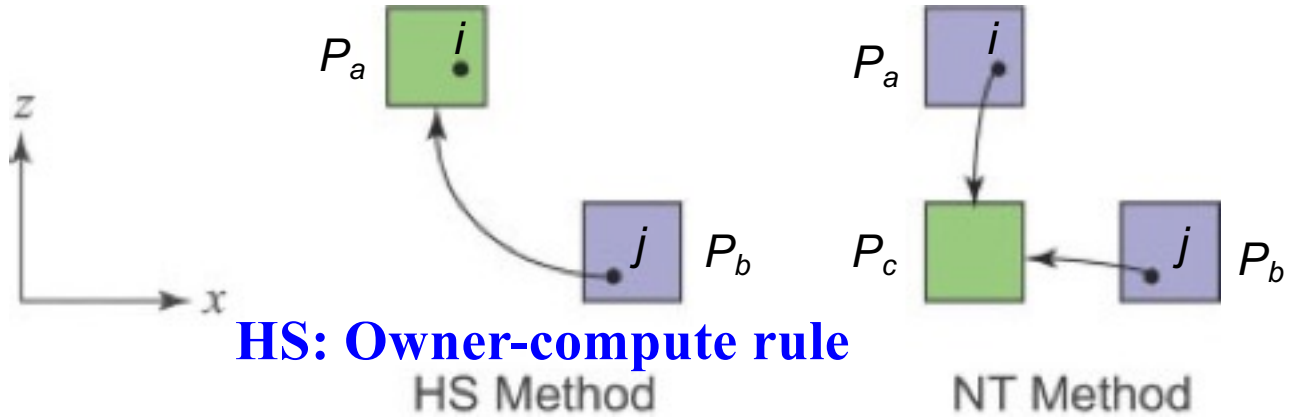

HS: Owner-compute rule

NT = hybrid spatial (data) & force (computation) decomposition with well-designed order/layout

Import regions or communication volume (2-dimensional example)

HS

$4bR+\pi R^2$

$\rightarrow$ const.

$(b \rightarrow 0)$

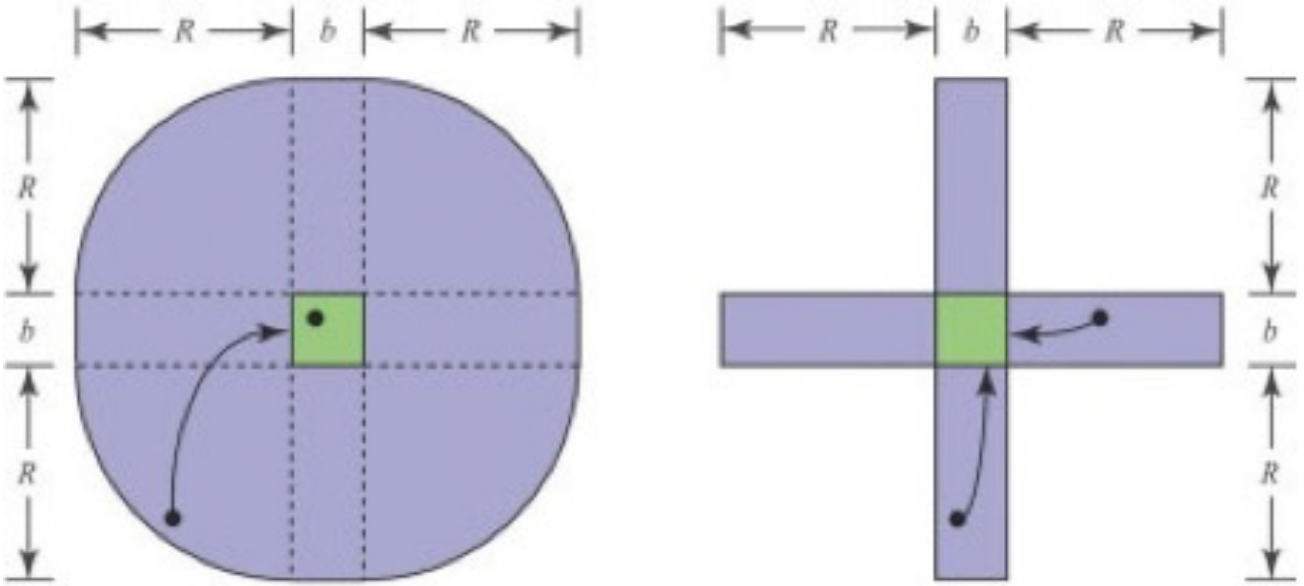

NT

$4bR$

$\rightarrow 0$

$(b \rightarrow 0)$

Import volume will be halved using Newton's 3rd law

3D Import Regions

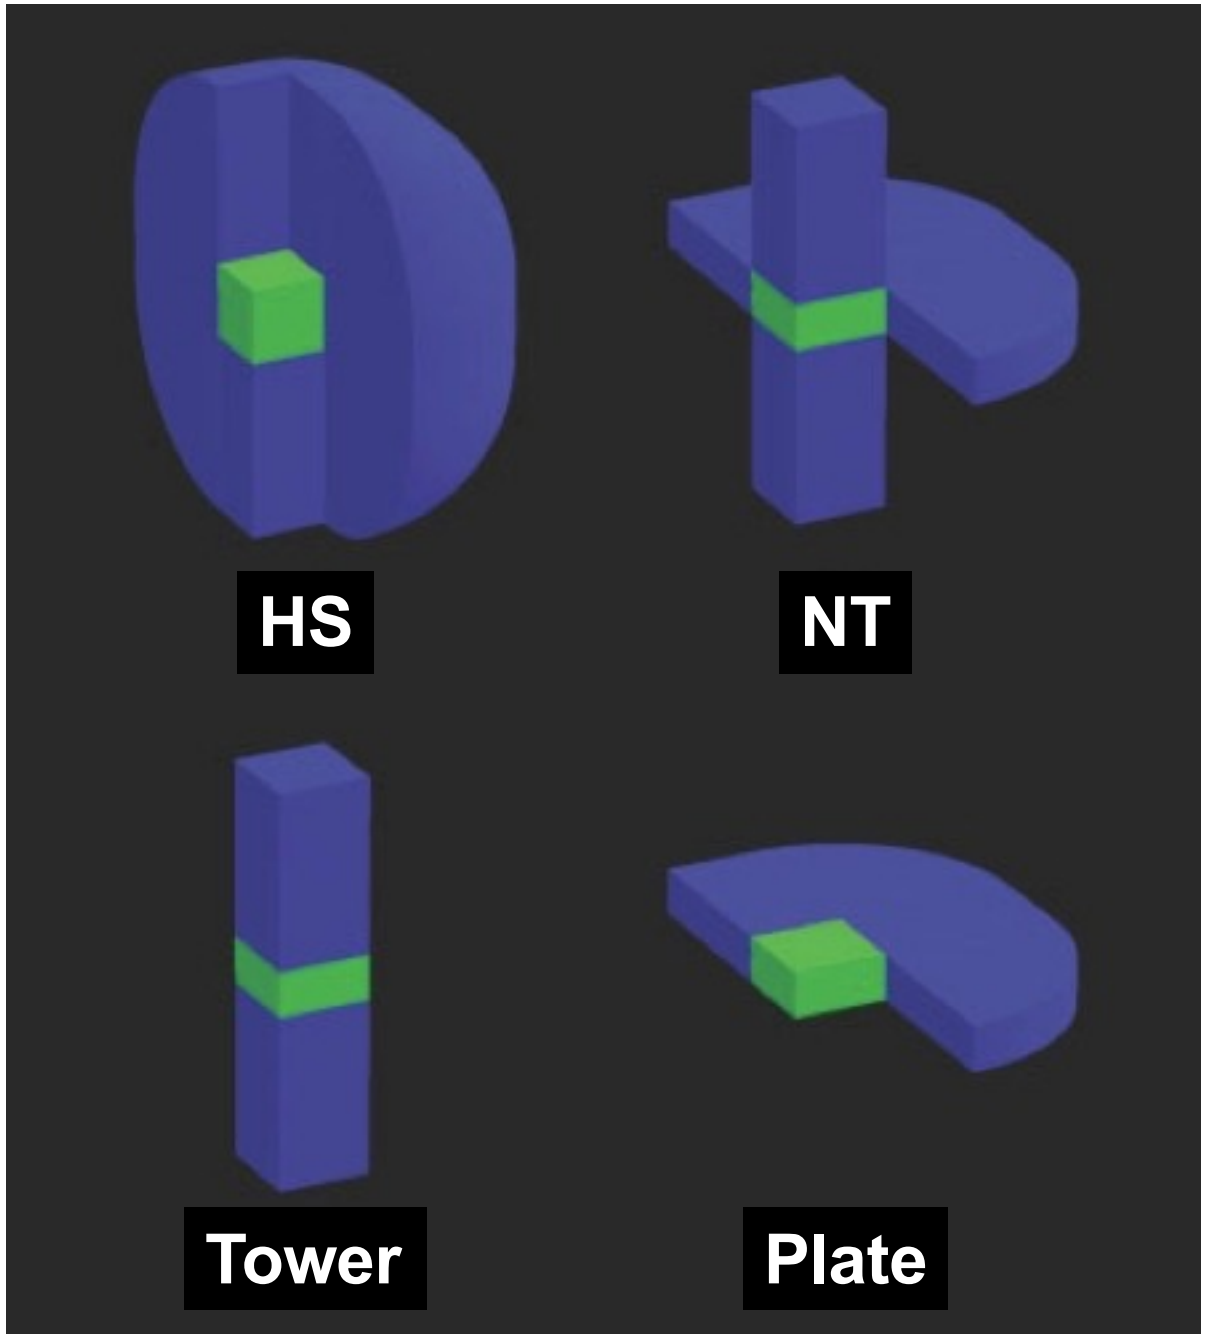

Scaling of Import Regions

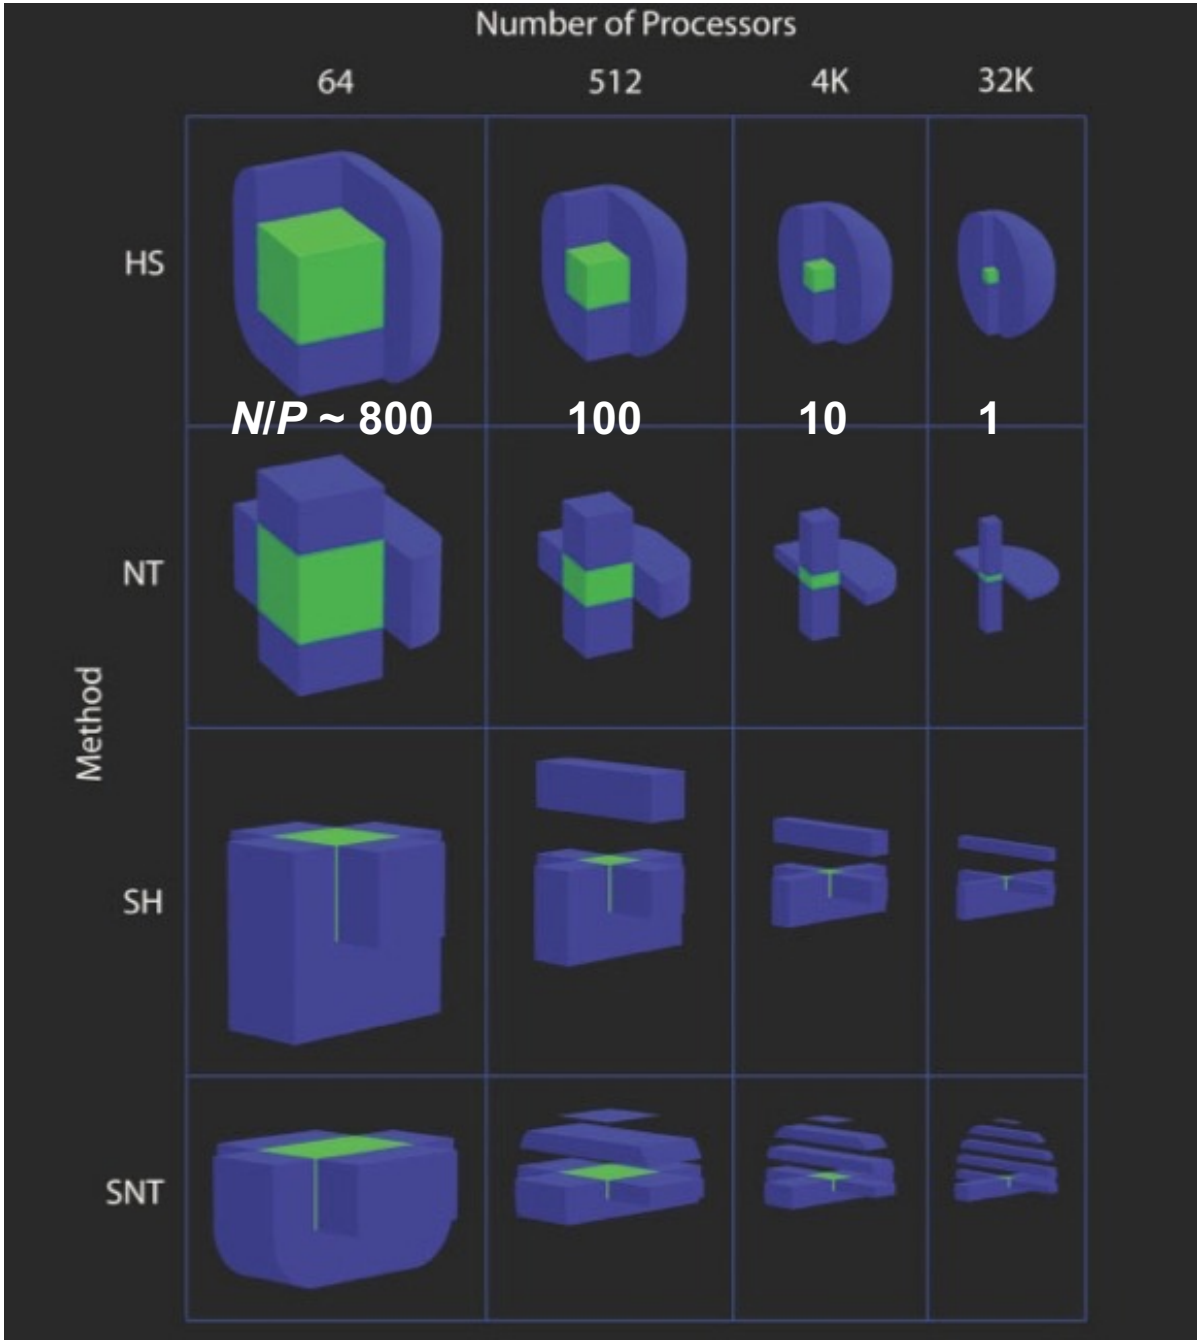

Marc Snir

Scaling of the Volume of Import Regions

HS decomposition

$$V_i = O\left(R^3\right)$$

NT decomposition

$$V_i = O\left(R^{3/2} p^{-1/2}\right)$$

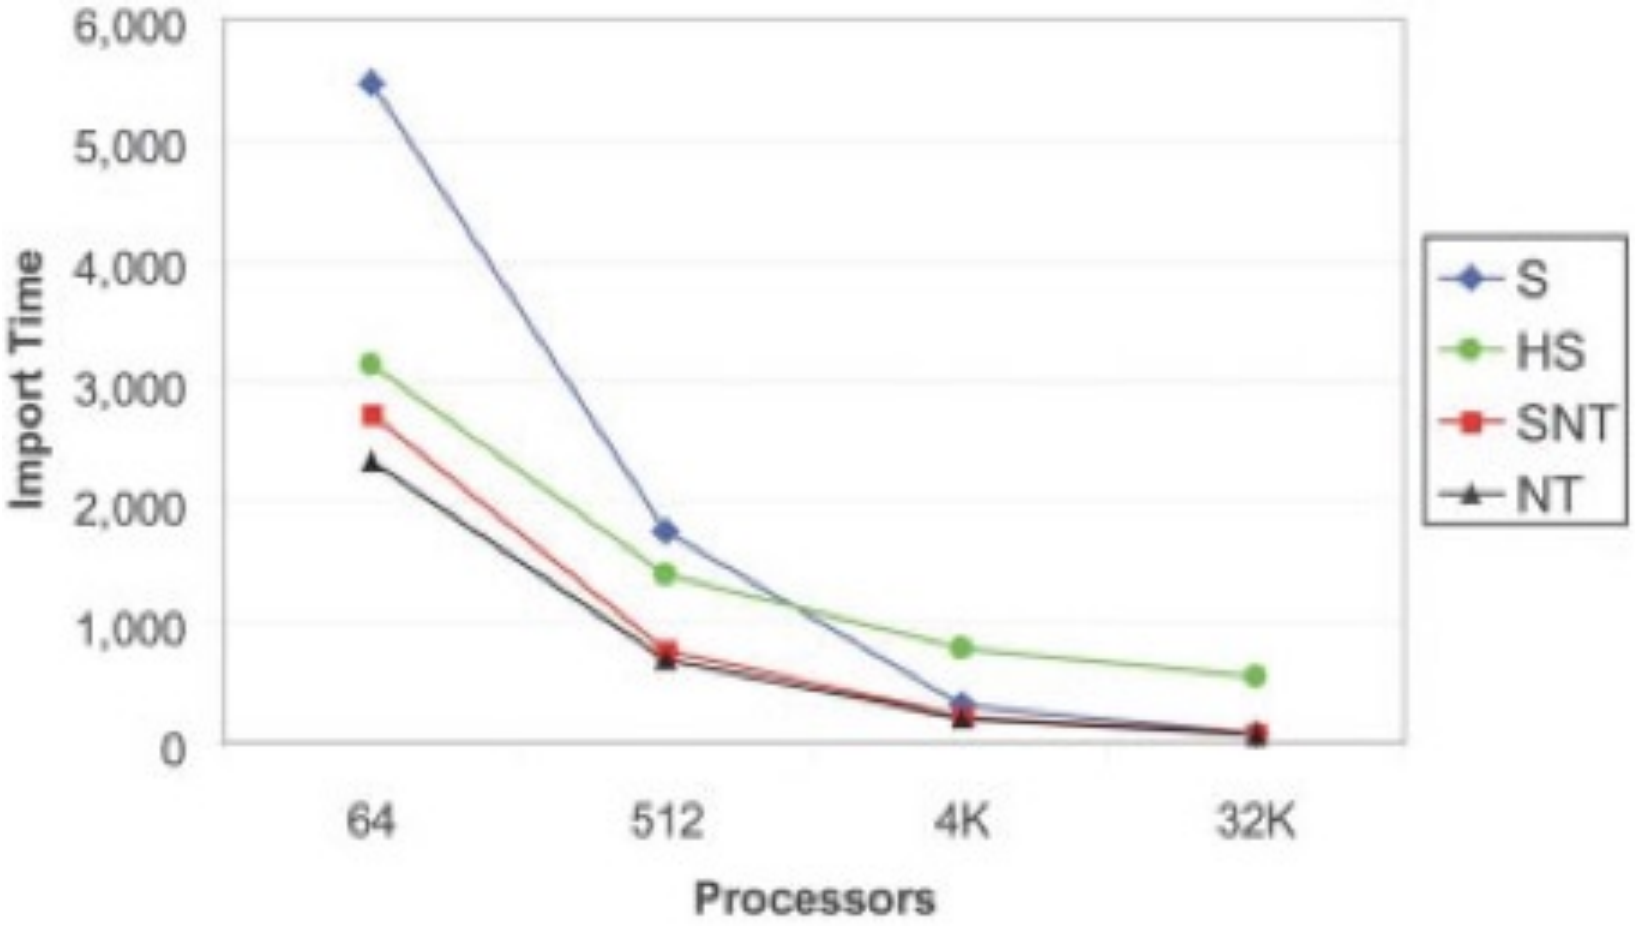

Communication time

$$T_{\text{comm}} = t_{\text{latency}} N_{\text{message}} + \frac{1}{b_{\text{bandwidth}}} V_{\text{message}}$$

# of messages — **volume (Bytes) of messages**

ns ~ many μs — **GB/s**

Combine NT with ...

Cache-oblivious recursive blocking?

Cache-Oblivious Algorithms

EXTENDED ABSTRACT SUBMITTED FOR PUBLICATION. FOCS99

Matteo Frigo Charles E. Leiserson Harald Prokop Sridhar Ramachandran
MIT Laboratory for Computer Science, 545 Technology Square, Cambridge, MA 02139
{athena,cel,prokop,sridhar}@supertech.lcs.mit.edu

SIAM REVIEW
Vol. 46, No. 1, pp. 3–45

© 2004 Society for Industrial and Applied Mathematics

Recursive Blocked Algorithms and Hybrid Data Structures for Dense Matrix Library Software*

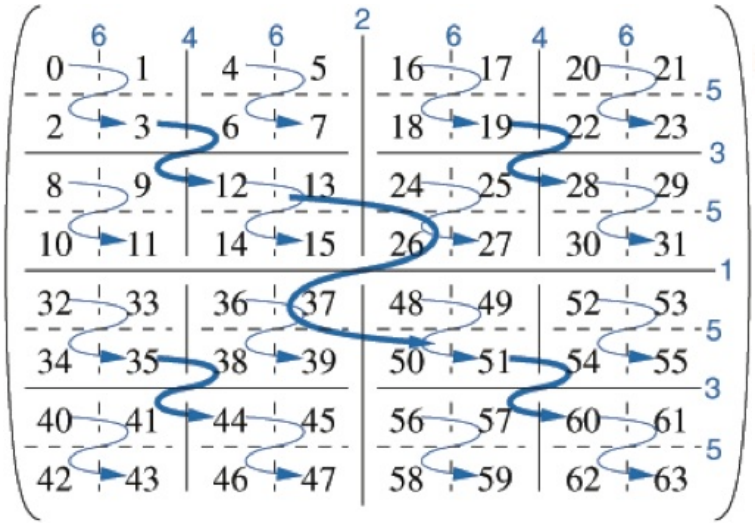

Erik Elmroth†
Fred Gustavson‡
Isak Jonsson†
Bo Kågström†

Combine NT with ...

Optimal data/computation layout (on Cell, GPU, multicore,...)?

Improving Memory Hierarchy Performance for Irregular Applications*

John Mellor-Crummey†, David Whalley‡, Ken Kennedy†

† Department of Computer Science, MS 132
Rice University
6100 Main
Houston, TX 77005
{johnmc,ken}@cs.rice.edu

‡ Computer Science Department
Florida State University
Tallahassee, FL 32306-4530
whalley@cs.fsu.edu
phone: (850) 644-3506

ISC99

IEEE TRANSACTIONS ON KNOWLEDGE AND DATA ENGINEERING, VOL. 13, NO. 1, JANUARY/FEBRUARY 2001

Analysis of the Clustering Properties of the Hilbert Space-Filling Curve

Bongki Moon, H.V. Jagadish, Christos Faloutsos, *Member, IEEE*, and Joel H. Saltz, *Member, IEEE*

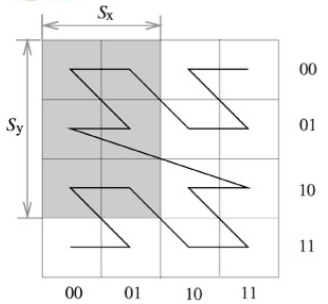

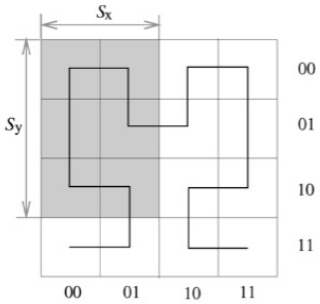

Metrics and Models for Reordering Transformations

Morton or Hilbert?

MSP04

Michelle Mills Strout
Mathematics and Computer Science Division
Argonne National Laboratory
Argonne, IL 60439 USA
mstrout@mcs.anl.gov

Paul D. Hovland
Mathematics and Computer Science Division
Argonne National Laboratory
Argonne, IL 60439 USA
hovland@mcs.anl.gov

Hypergraph

G.M. Morton, "A computer oriented geodetic data base & a new technique in file sequencing," *IBM Tech. Report* ('66)

Shift-Collapse (SC) Algorithm

- **Generalization of Shaw's eighth-cell method (non-owner-compute method on high-latency cluster) for pair computation to general dynamic range-limited *n*-tuples**

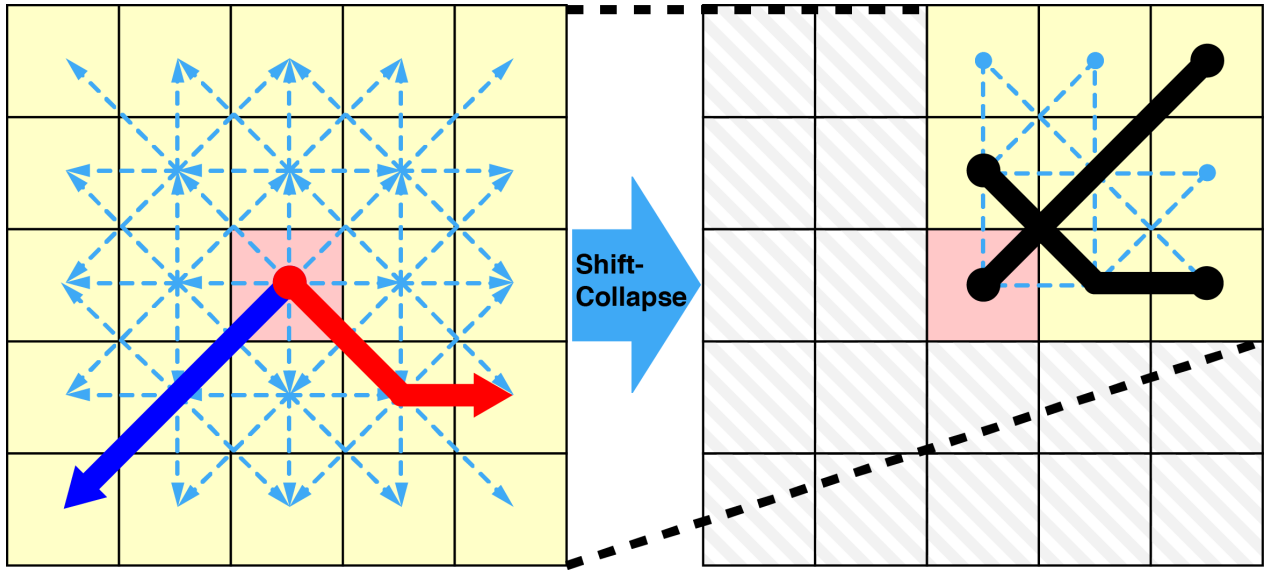

M. Kunaseth *et al.*, *IEEE/ACM Supercomputing* (*SC13*)

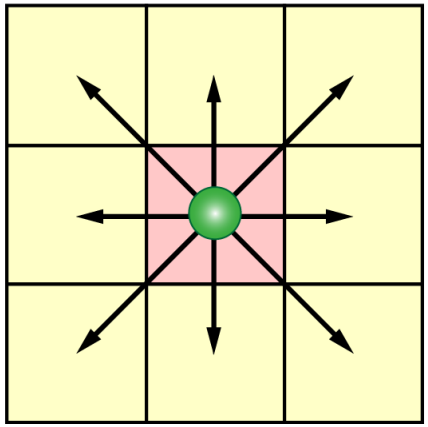

Full-shell (FS) method
e.g. Rappaport, '88

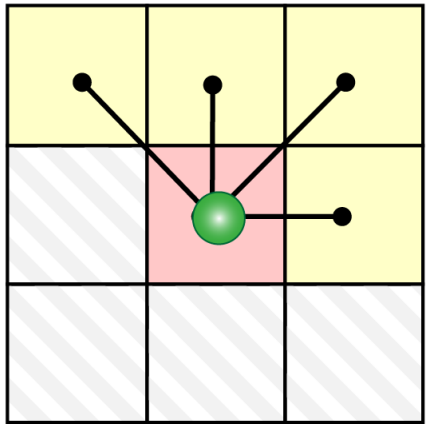

Half-shell (HS) method
e.g. Rappaport, '88

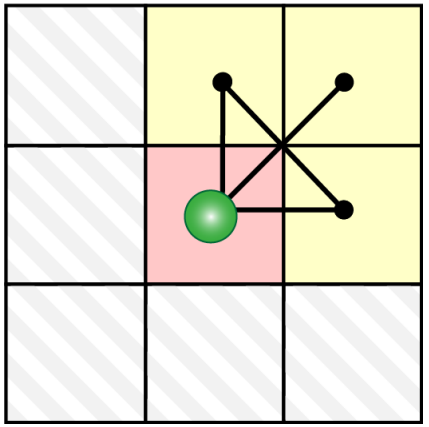

Eighth-shell (ES) method
Bower *et al.*, '06

Shift-Collapse (SC) Performance

Runtime comparison on 48 Intel-Xeon nodes and 64 Blue Gene/Q nodes

- **SC-MD is always faster than FS-MD**
- **At the smallest grain, SC-MD is 9.7- and 5.1-fold speedups over the state-of-the-art hybrid linked-cell & neighbor list code**
- **Crossover of optimal algorithm from SC-MD to hybrid MD at larger granularity (*i.e.*, $N/P > 2{,}095$ on Intel Xeon and $N/P > 425$)**

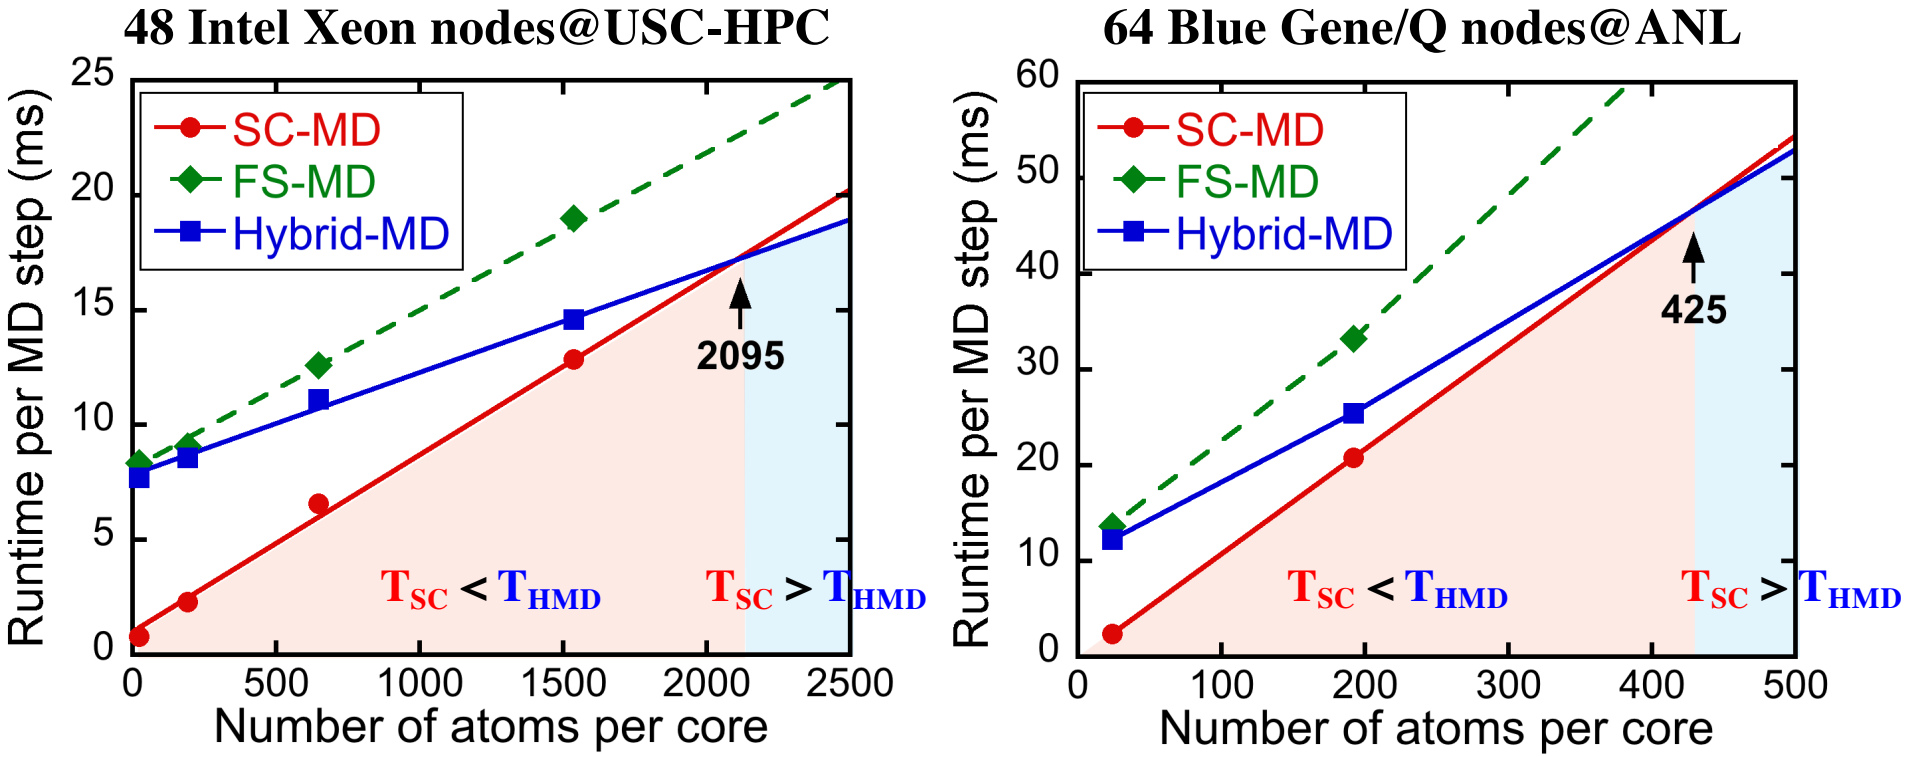

M. Kunaseth *et al.*, *IEEE/ACM Supercomputing* (*SC13*)

Shift-Collapse on Neighbor List (SC-NBL)

- **Apply shift-collapse operations to the hybrid linked-cell & neighbor list code (best of both)**

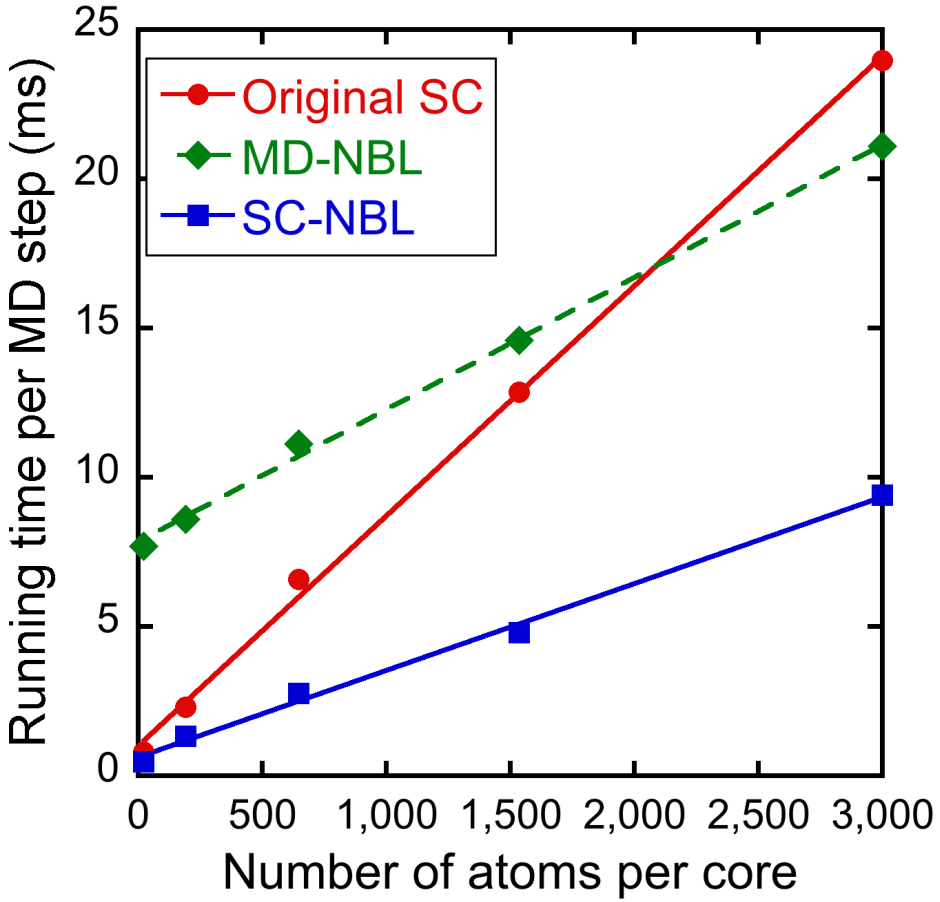

Shift/collapse on neighbor list (SC-NBL): fast evaluation of dynamic many-body potentials in molecular dynamics simulations, M. Kunaseth, S. Hannongbua, & A. Nakano, *Comput. Phys. Commun.* **235**, 88 ('19)

Challenge: Expose massive data parallelism for SC on graphics processing unit (GPU)