

Order-Invariant Real Number Summation: Circumventing Accuracy Loss for Multimillion Summands on Multiple Parallel Architectures

Patrick E. Small, Rajiv K. Kalia, Aiichiro Nakano, Priya Vasishta

Collaboratory for Advanced Computing & Simulations

Department of Computer Science

Department of Physics & Astronomy

Department of Biological Sciences

University of Southern California

Email: [\(patrices,rkalia,anakano,priyav\)@usc.edu](mailto:(patrices,rkalia,anakano,priyav)@usc.edu)

Proc. IEEE International Parallel & Distributed Processing Symposium, IPDPS, p. 152 ('16)

<https://aiichironakano.github.io/cs596/Small-OrderInvariantSum-IPDPS16.pdf>

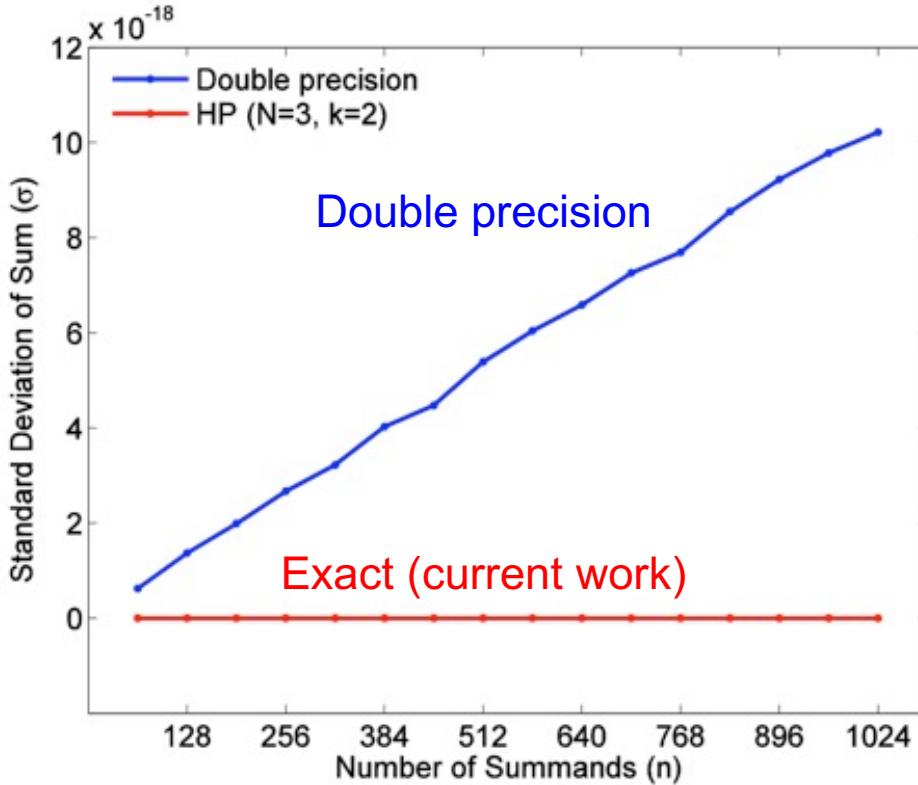


Optimize bits!

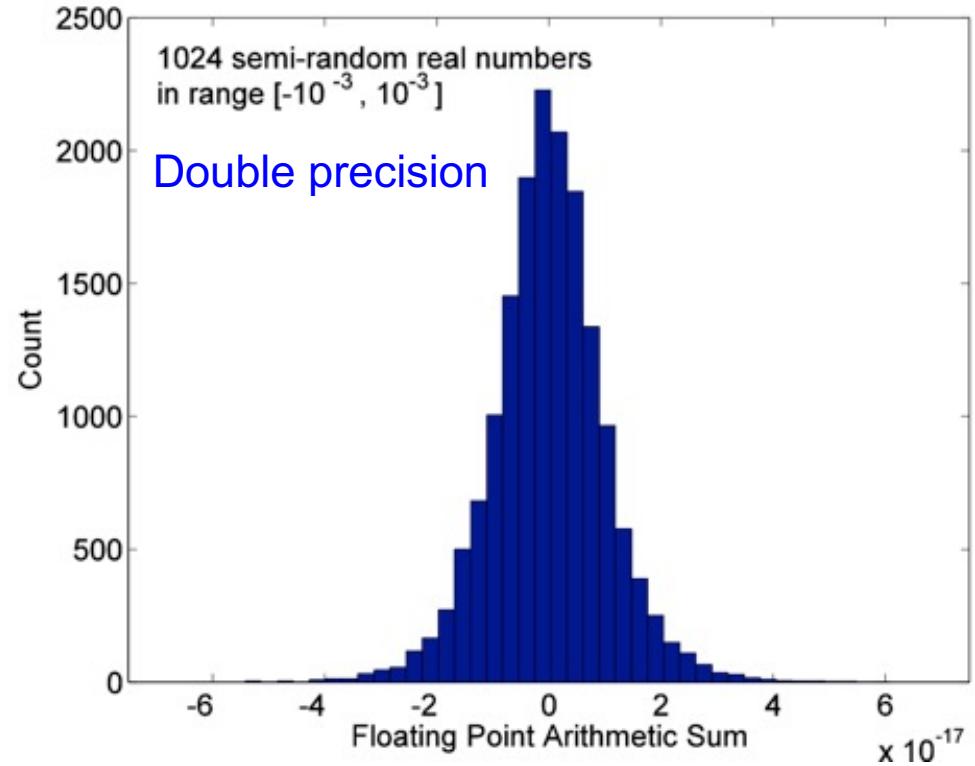


Reproducibility Challenge

- Rounding (truncation) error makes floating-point addition non-associative
 $(a + b) + c \neq a + (b + c)!$



Standard deviation of sum with random summation orders



Distribution of sum with random summation orders

- Finding: Sum becomes a random walk across the space of possible rounding error

High-Precision (HP) Method

- Propose an extension of the order-invariant, higher-precision intermediate-sum method by Hallberg & Adcroft *Par. Comput.* **40**, 140 ('14)
- The proposed variation represents a real number r using a set of N 64-bit unsigned integers, a_i ($i = 0, N-1$)

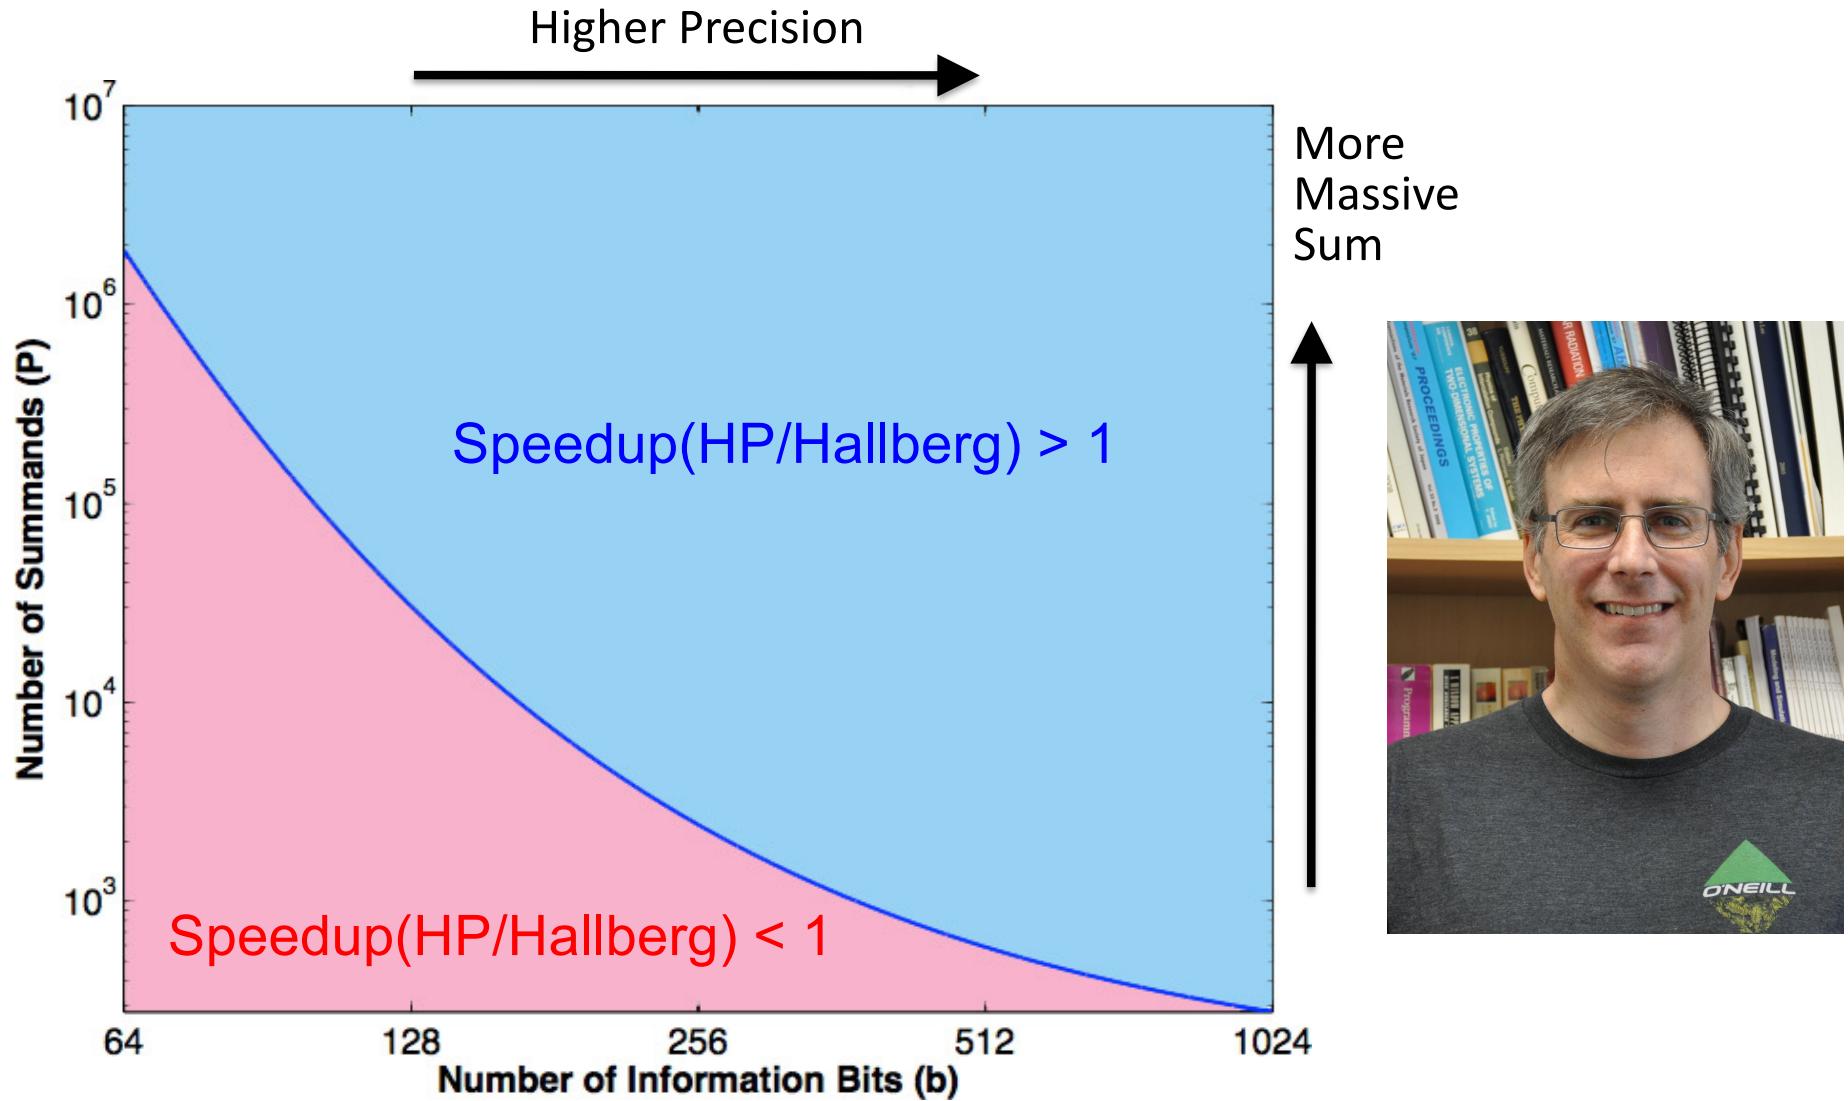
$$\begin{aligned} r &= \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)} \\ &= \overbrace{a_0 2^{64(N-k-1)} + \cdots + a_{N-k-1} 2^{64(1)}}^{\text{$N-k$}} + \cdots + \overbrace{a_{N-k} 2^{-64} + \cdots + a_{N-1} 2^{-64k}}^{\text{k}} \end{aligned}$$

- k is the number of 64-bit unsigned integers assigned to represent the fractional portion of r ($0 \leq k \leq N$), whereas $N-k$ integers represent the whole-number component
- Negative number is represented by two's complement in integer representation, using only 1 bit

If you are the first to find the problem, the simplest solution suffices to prove the concept

Performance Projection

- HP sum is faster than Hallberg sum for higher precision & larger numbers of summands



Detailed IPDPS 2016 Presentation



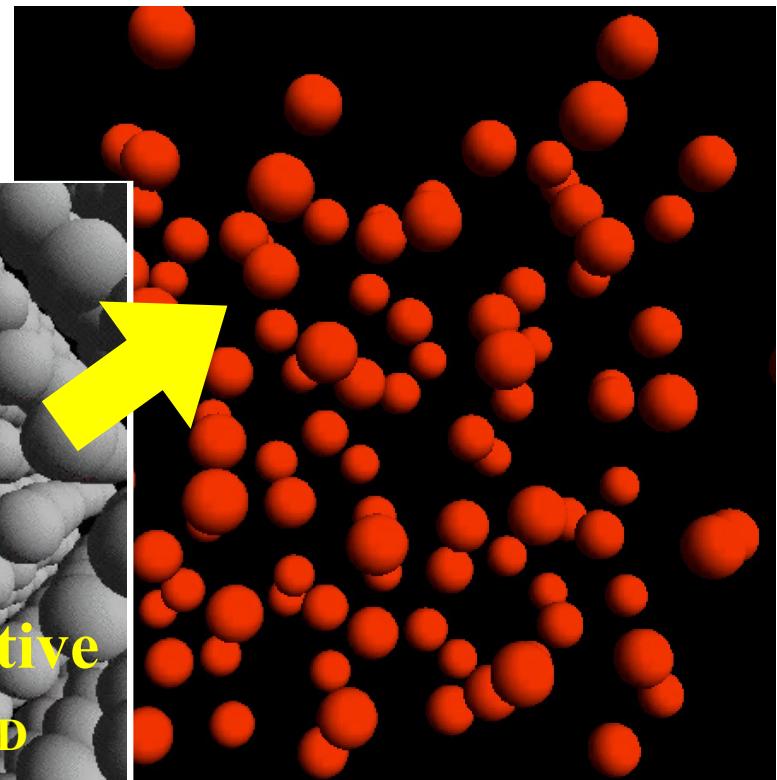
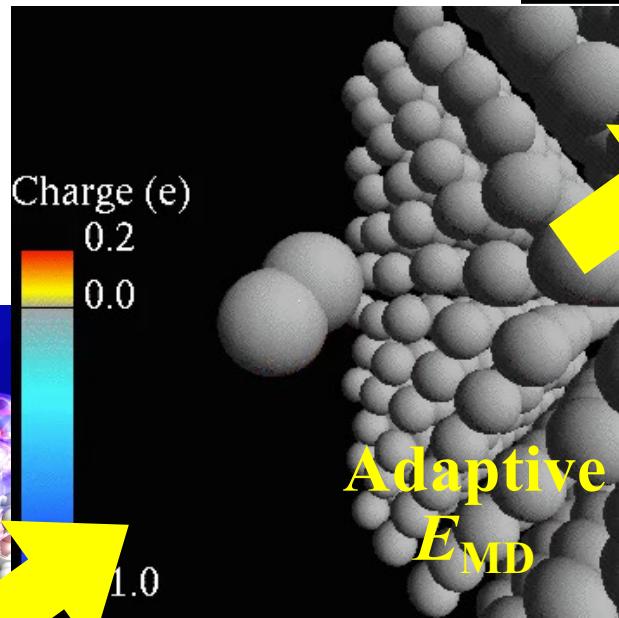
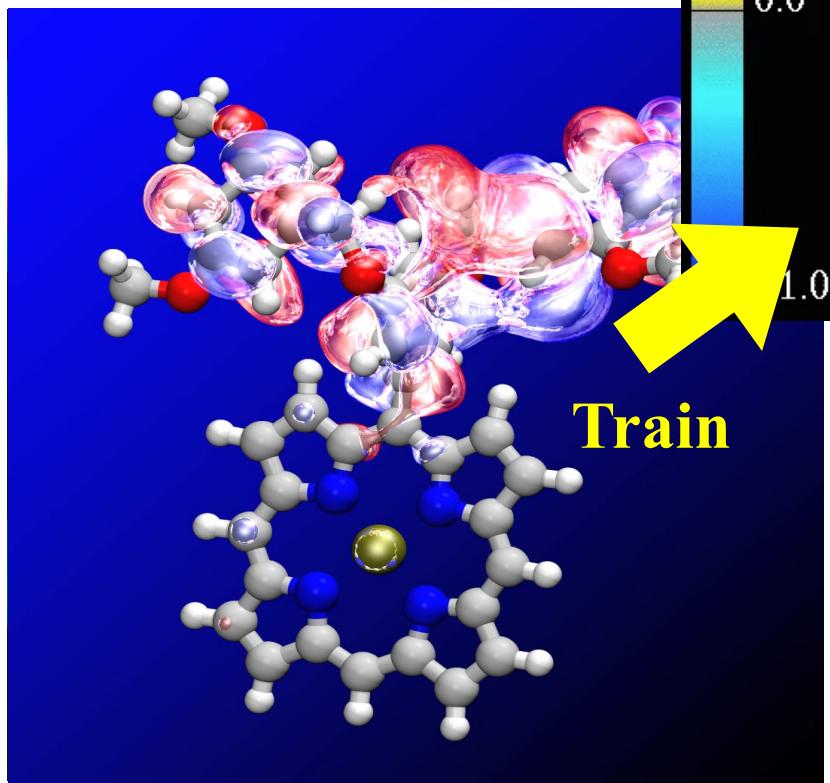
RIP, Patrick

Hierarchy of Atomistic Simulation Methods

Molecular Dynamics (MD)

Reactive MD (RMD)

Nonadiabatic quantum MD (NAQMD)

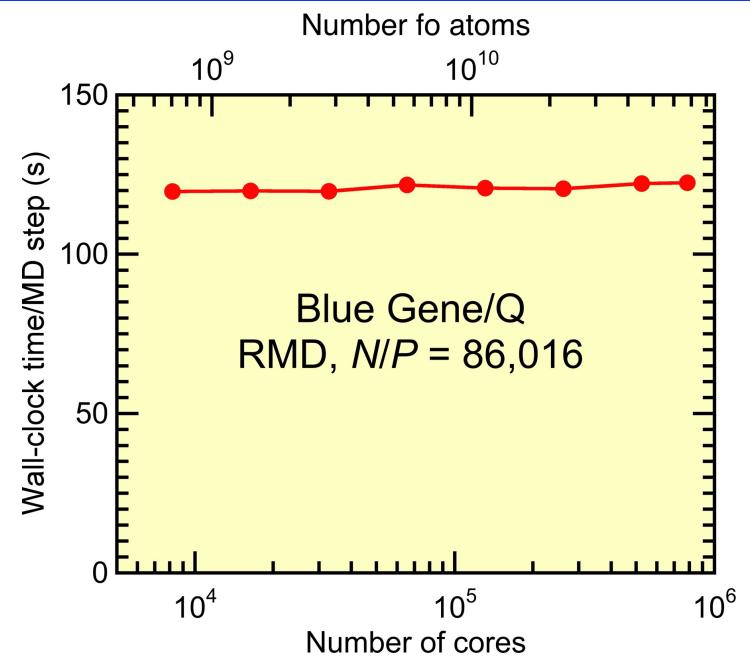
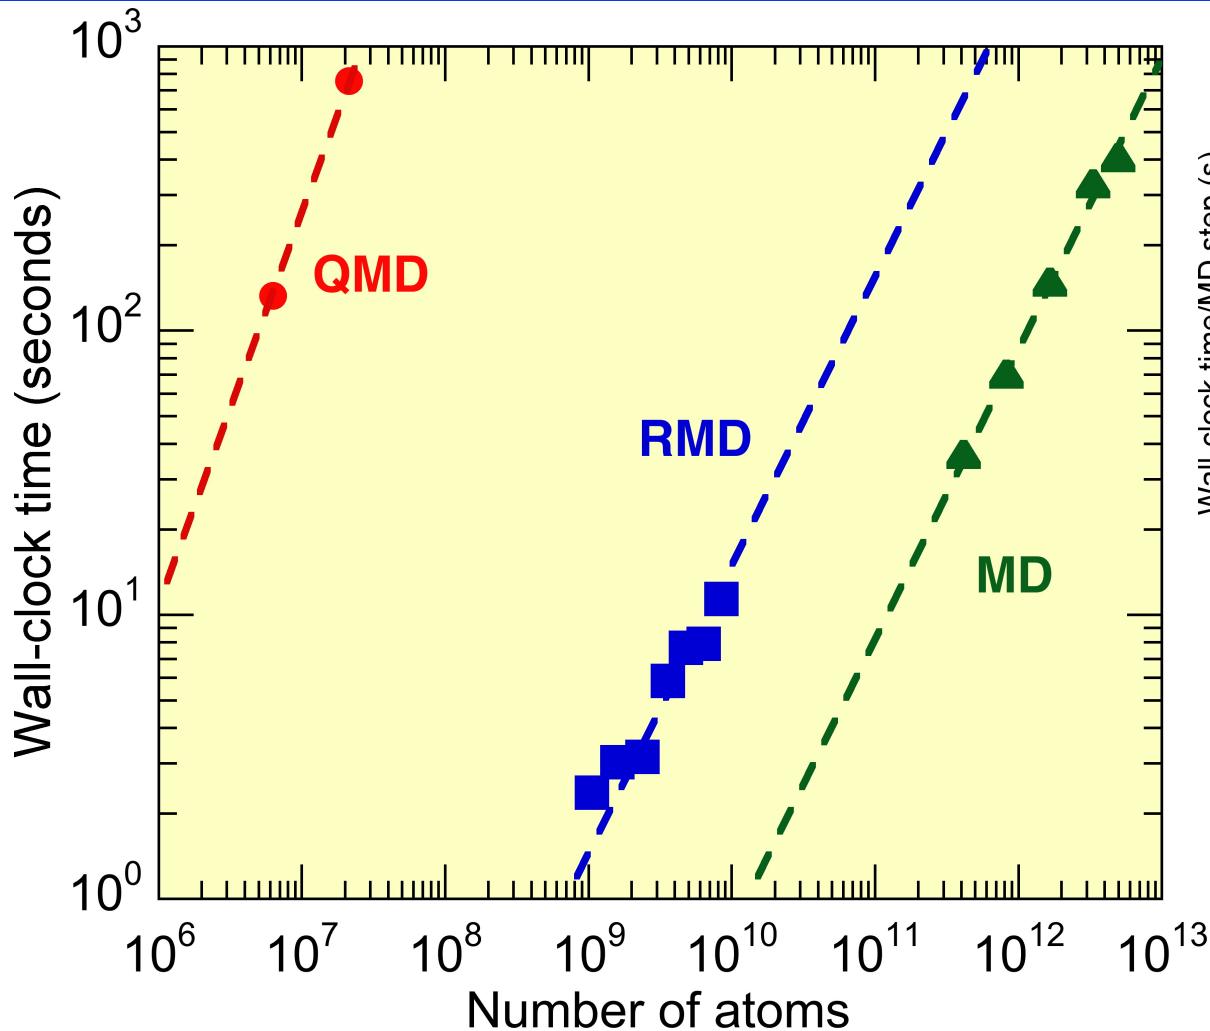


First principles-based reactive force-fields

- **Reactive bond order $\{BO_{ij}\}$**
→ Bond breakage & formation
- **Charge equilibration (QEeq) $\{q_i\}$**
→ Charge transfer

Tersoff, Brenner, Sinnott *et al.*; Streitz & Mintmire *et al.*;
van Duin & Goddard (ReaxFF)

Scalable Simulation Algorithm Suite



QMD (quantum molecular dynamics): DC-DFT
RMD (reactive molecular dynamics): F-ReaxFF
MD (molecular dynamics): MRMD

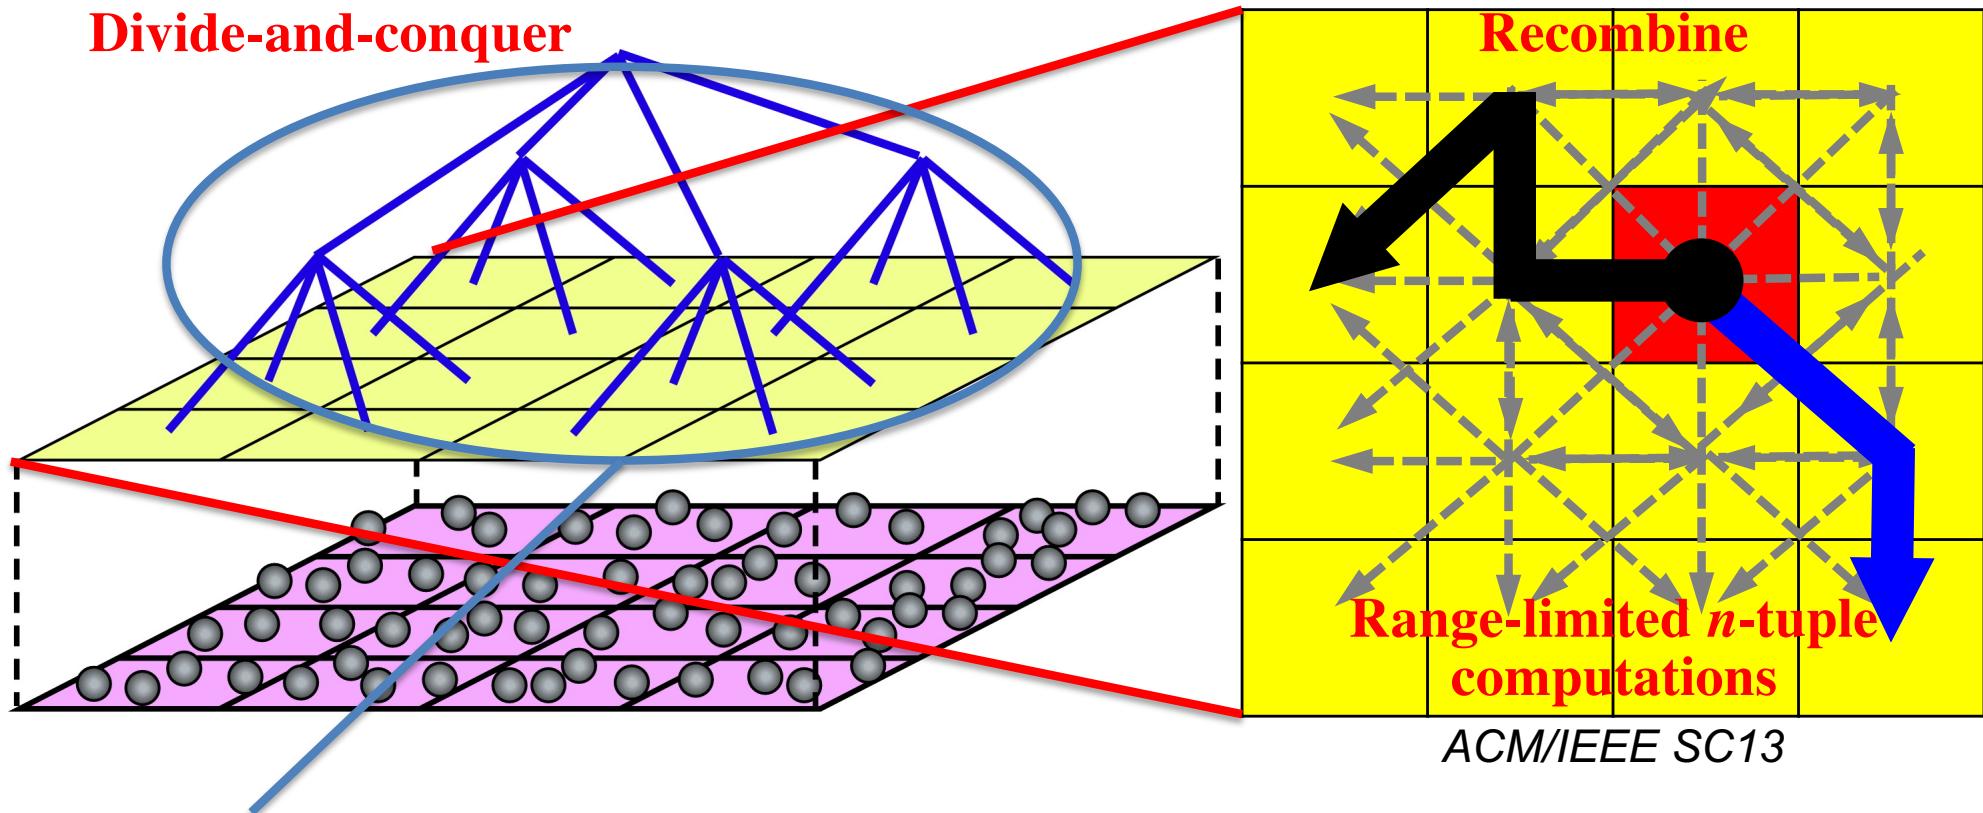
- 4.9 trillion-atom space-time multiresolution MD (MRMD) of SiO₂
- 67.6 billion-atom fast reactive force-field (F-ReaxFF) RMD of RDX
- 39.8 trillion grid points (50.3 million-atom) DC-DFT QMD of SiC
parallel efficiency 0.984 on 786,432 Blue Gene/Q cores

Exascale Computing Challenge

1. Scalability beyond million-way parallelism

J. Chem. Phys. **140**, 18A529 ('14)
IEEE/ACM SC14
IEEE Computer **48**(11), 33 ('15)

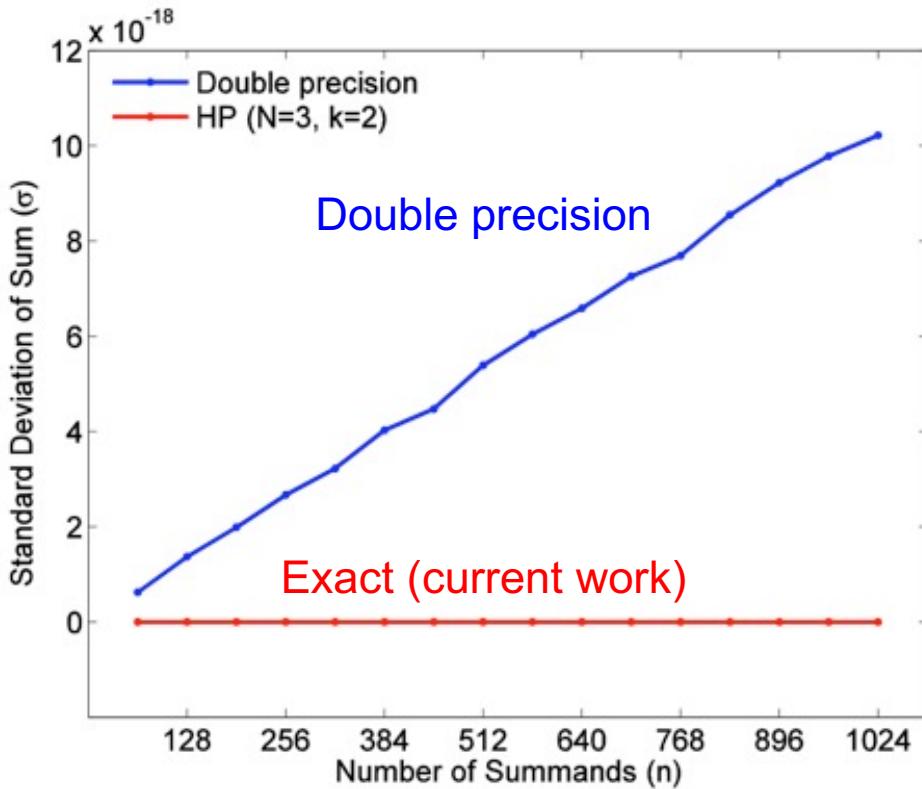
Divide-conquer-recombine (DCR) algorithmic framework
Metascalable (“design once, scale on future architectures”)



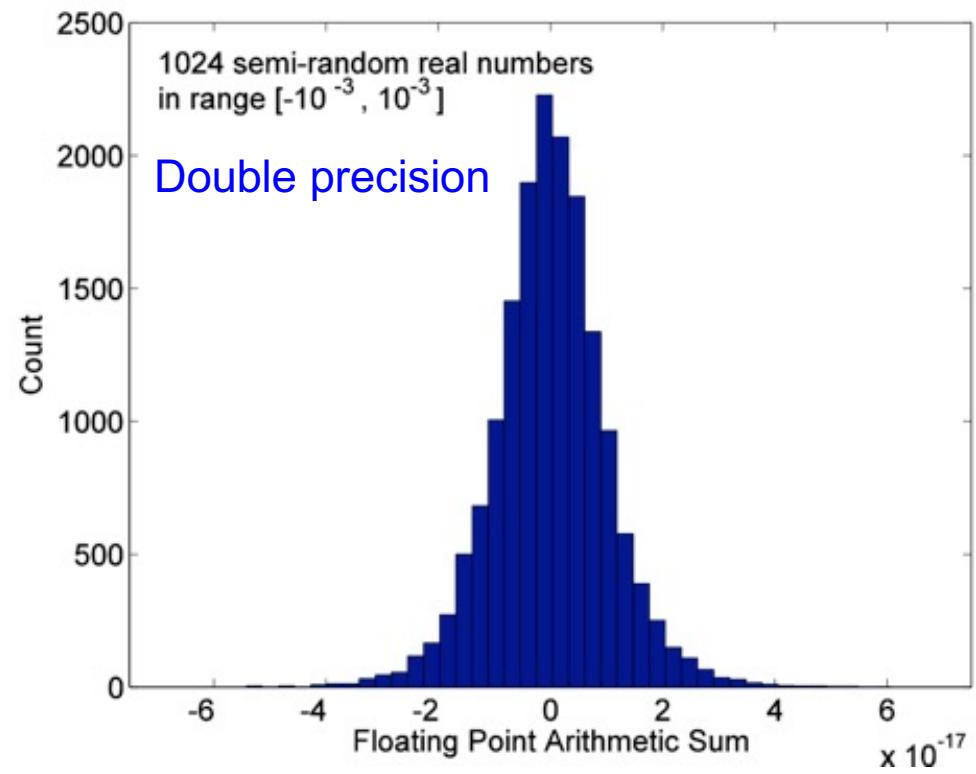
2. Reproducibility of real-number summation for multimillion summands & beyond in the global sum; double-precision arithmetic began to produce different results on different high-end architectures

Reproducibility Challenge

- Rounding (truncation) error makes floating-point addition non-associative



Standard deviation of sum with
random summation orders



Distribution of sum with random
summation orders

- Sum becomes a random walk across the space of possible rounding error

Related Works

- General-purpose arbitrary precision arithmetic
[GNU-MPL '12]
 - Extensive computation & memory usage
- Error-compensation methods
 - > Error-free transformation for tracking residuals
[Priest, '91, Higham '93, Rump '09, Demmel '13]
 - Complex implementation
 - > Summation reordering for minimizing error
[Hel '01]
 - Prohibitive at large scales
- Hardware solutions
[Gustafson '15]
 - Not available yet
- Higher-precision intermediate sums
[He '01, Hallberg '14]
 - Simple implementation, low overhead

Contributions

- Propose an extension of the order-invariant, higher-precision intermediate-sum method by Hallberg & Adcroft [*Par. Comput.* **40**, 140 ('14)]:
 - (1) Improves performance* for large ($> 10^6$) number of summands
 - (2) Eliminates the aliasing problem of the original method
- The new method outperforms the previous state-of-the-art for large problems involving million+ summands on broad systems (MPI, OpenMP, CUDA/GPU, Xeon Phi)

*Performance is defined as the computational speed

Hallberg Order-Invariant Sum

- **Integer representation with higher accuracy:** Represent a real number r using a set of N 64-bit signed integers, a_i ($i = 0, N-1$); $M (< 63)$ is a positive integer

$$r = \sum_{i=0}^{N-1} a_i 2^{\left(i - \frac{N}{2}\right)M} = 2^{-NM/2} (a_0 + a_1 2^M + a_2 2^{2M} + \dots)$$

- **Order-invariant parallel sum:** Two real numbers are added by summing N pairs of corresponding integers concurrently
- **Carry out (potential sequential dependence):** When any of the integer additions exceeds 2^M , carry out must be added to the next integer in the set
- **Carry-overhead reduction:** Carry operations are avoided up to $P = 2^{63-M}-1$ summands to expose high parallelism

Drawback of Hallberg Sum

- **Overhead:** Not all integer bits serve to provide real-number precision; $63 \cdot M$ bits per integer are dedicated to book-keeping
- **Aliasing:** Multiple integer representations could represent the same real number
- **Normalization & sum overheads** to convert the integer representation back to real

High-Precision (HP) Method

- The proposed variation of Hallberg method represents a real number r using a set of N 64-bit unsigned integers, a_i ($i = 0, N-1$)

$$r = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)} \\ = \underbrace{a_0 2^{64(N-k-1)} + \cdots + a_{N-k-1} 2^{-64}}_{N-k} + \underbrace{a_{N-k} 2^{-64} + \cdots + a_{N-1} 2^{-64k}}_k$$

- k is the number of 64-bit unsigned integers assigned to represent the fractional portion of r ($0 \leq k \leq N$), whereas $N-k$ integers represent the whole-number component
- Negative number is represented by two's complement in integer representation, using only 1 bit

HP Algorithm (1): Conversion

- **Simple procedure:** A single pass converts a double-precision number r to HP integers a_i & translates them to two's complement

$$r = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)}$$

```
dtmp = fabs(r)*264*(N-k-1);
isneg = (r < 0.0);
for (i=0; i<N-1; i++) {
    itmp = (uint64_t)dtmp;
    dtmp = (dtmp - (double)itmp)*264;
    a[i] = (isneg) ? ~itmp + (dtmp<=0.0) : itmp;
}
a[N-1] = (isneg) ? ~(uint64_t)dtmp + 1 : (uint64_t)dtmp;
```

- Inverse of this algorithm converts HP number back to double-precision

HP Algorithm (2): Addition

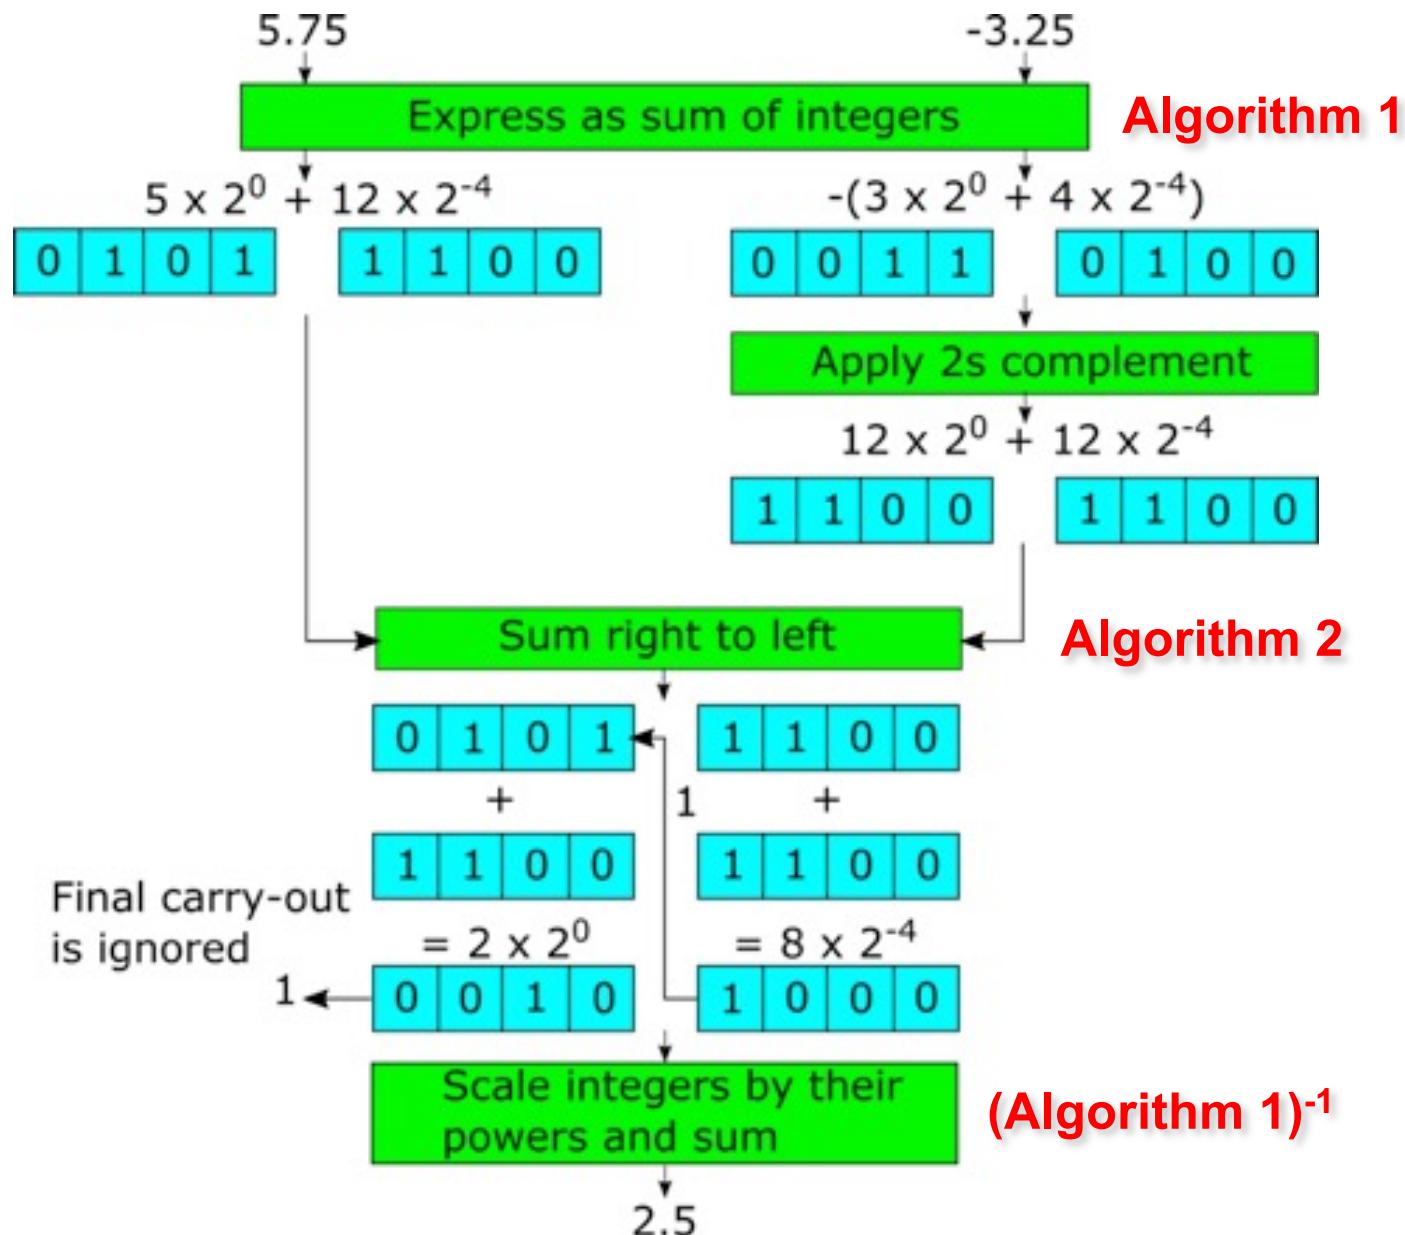
- **Addition of two HP numbers, $a \leftarrow a + b$**

$$\begin{cases} r_1 = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)} \\ r_2 = \sum_{i=0}^{N-1} b_i 2^{64(N-k-i-1)} \end{cases}$$

```
a[N-1] = a[N-1]+b[N-1];
co = (a[N-1]<b[N-1]);
for (i=N-2; i>=1; i--) {
    a[i] = a[i]+b[i]+co;
    co = (a[i]==b[i]) ? co : (a[i]<b[i]);
}
a[0] = a[0]+b[0]+co;
```

- **Overflow of the sum is detected by comparing the signs of the summands with that of the sum**

HP Sum: Example



Representation Power

- Maximum range & smallest representable HP number

$$r_{\text{HP}} = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)}$$

N	k	Bits	Maximum range	Smallest number
2	1	128	$\pm 9.223372 \times 10^{18}$	5.421011×10^{-20}
3	2	192	$\pm 9.223372 \times 10^{18}$	2.938736×10^{-39}
6	3	256	$\pm 3.138551 \times 10^{57}$	1.593092×10^{-58}
8	4	512	$\pm 5.789604 \times 10^{76}$	8.636169×10^{-78}

- Equivalency with Hallberg representation power

$$r_{\text{Hallberg}} = \sum_{i=0}^{N-1} a_i 2^{(i-N/2)M}$$

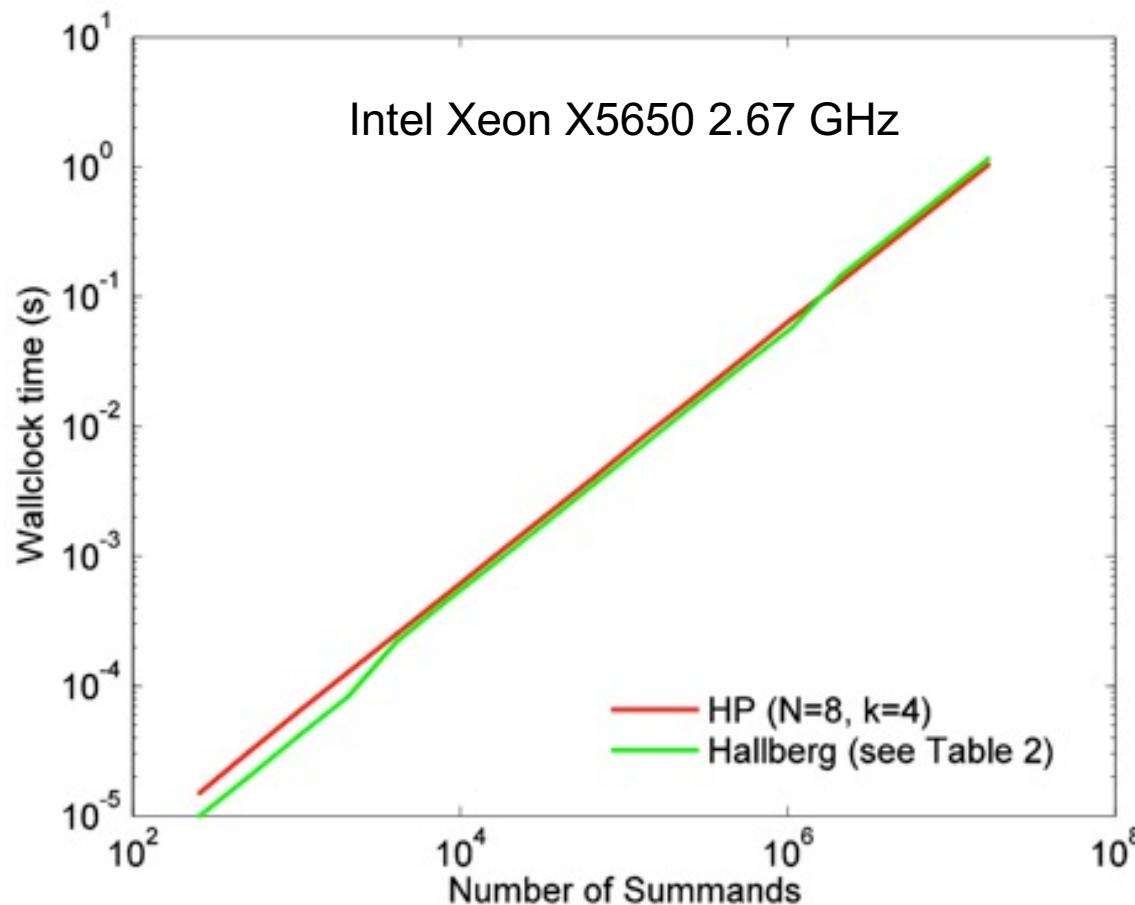
N	M	Precision bits	Max summands
10	52	520	2048
12	43	516	1 M
14	37	518	64 M

HP Method: Properties

- Invariance of sum with respect to both summation order & architecture is guaranteed with appropriate setting of N & k to provide sufficient accuracy
- Overflow & underflow can be readily detected at runtime at double-precision (DP)-to-HP conversion, HP-sum & HP-to-DP conversion steps
- Atomicity of addition (which is essential for multithreading) is guaranteed using only the widely available compare-&-swap (CAS) synchronization primitive

Performance

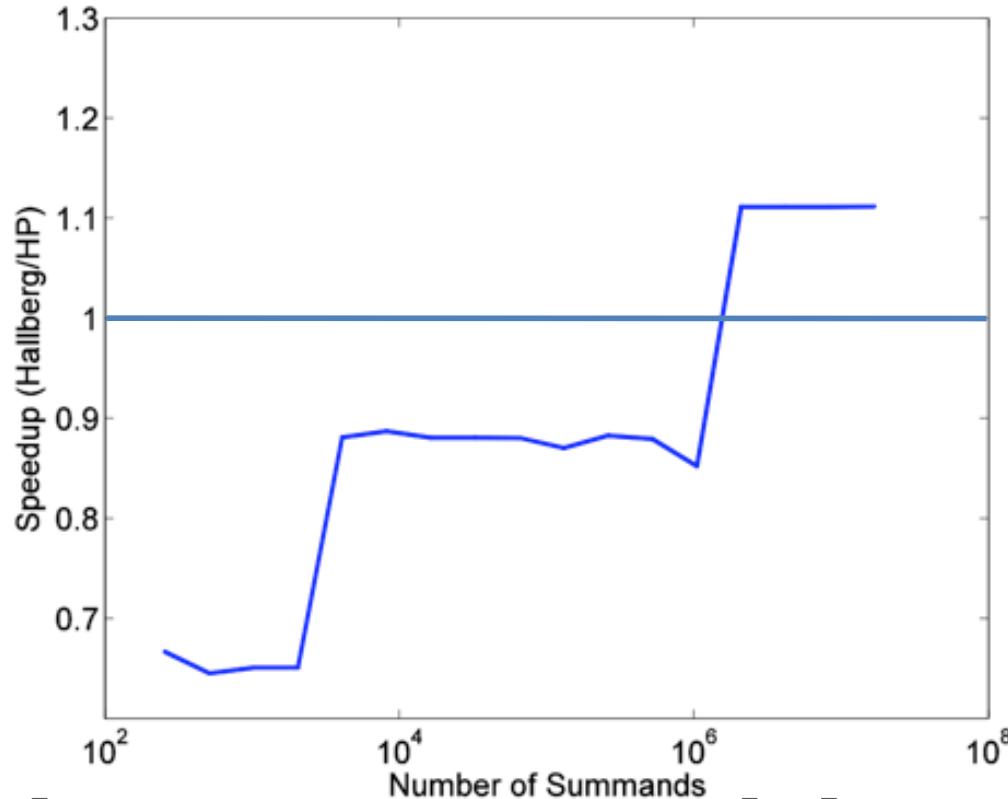
- Computing time of real-number sum using the current (HP) & Hallberg methods as a function of # of summands



- HP sum is faster than Hallberg sum over million summands

Performance Analysis

- Speedup of HP sum over Hallberg sum as a function of the number of summands

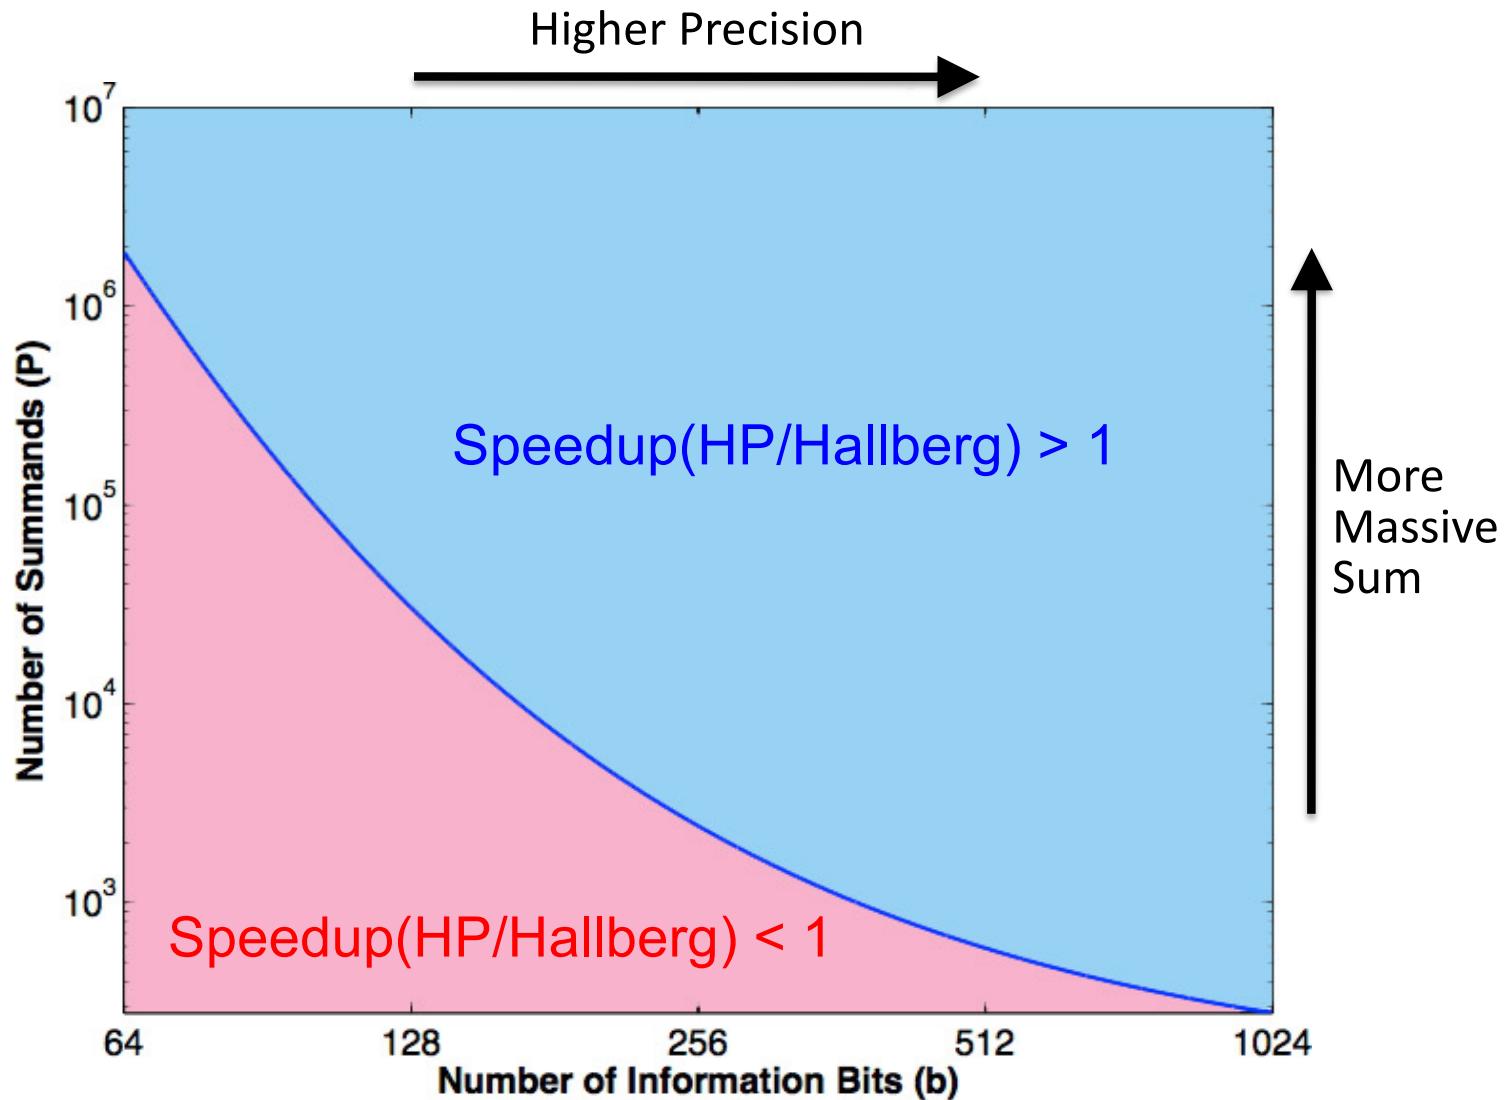


$$T_{\text{HP}} = c_{\text{HP}} \left\lceil \frac{b+1}{64} \right\rceil \quad T_{\text{Hallberg}} = c_{\text{Hallberg}} \left\lceil \frac{b}{M} \right\rceil \quad b: \text{Precision bit count}$$

$$\text{Speedup} \sim \frac{c_{\text{Hallberg}}}{c_{\text{HP}}} \frac{64b}{M(b+65)} \geq \frac{32c_{\text{Hallberg}}}{c_{\text{HP}}} \frac{1}{M} \quad 2^{63-M} \propto \#\text{summands}$$

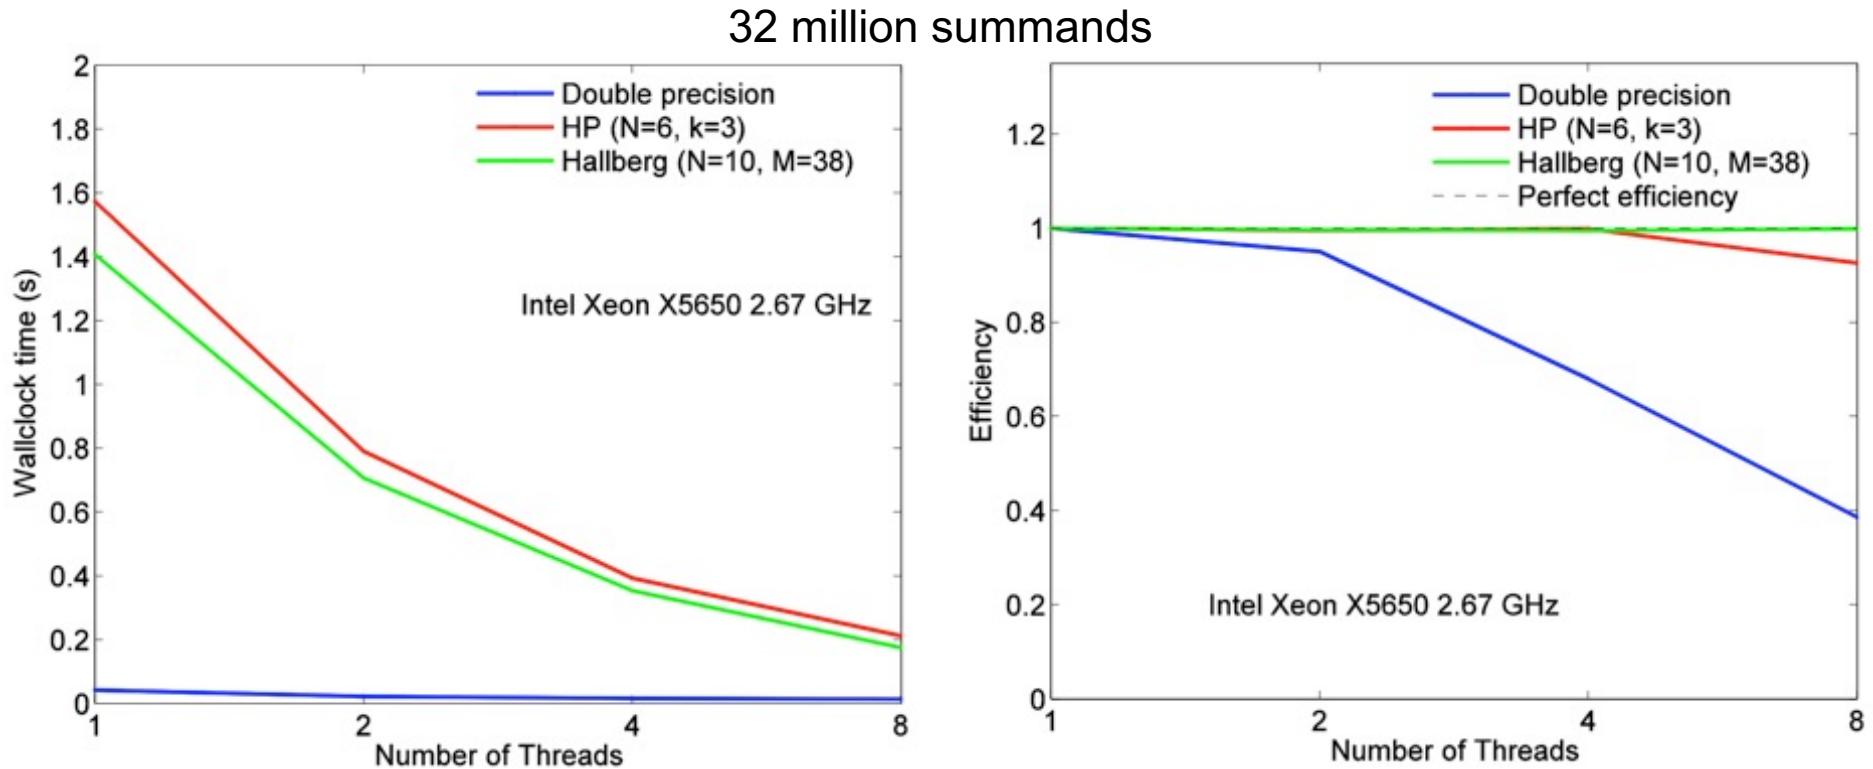
Performance Projection

- HP sum is faster than Hallberg sum for larger numbers of summands & higher precision



Parallel Efficiency with OpenMP

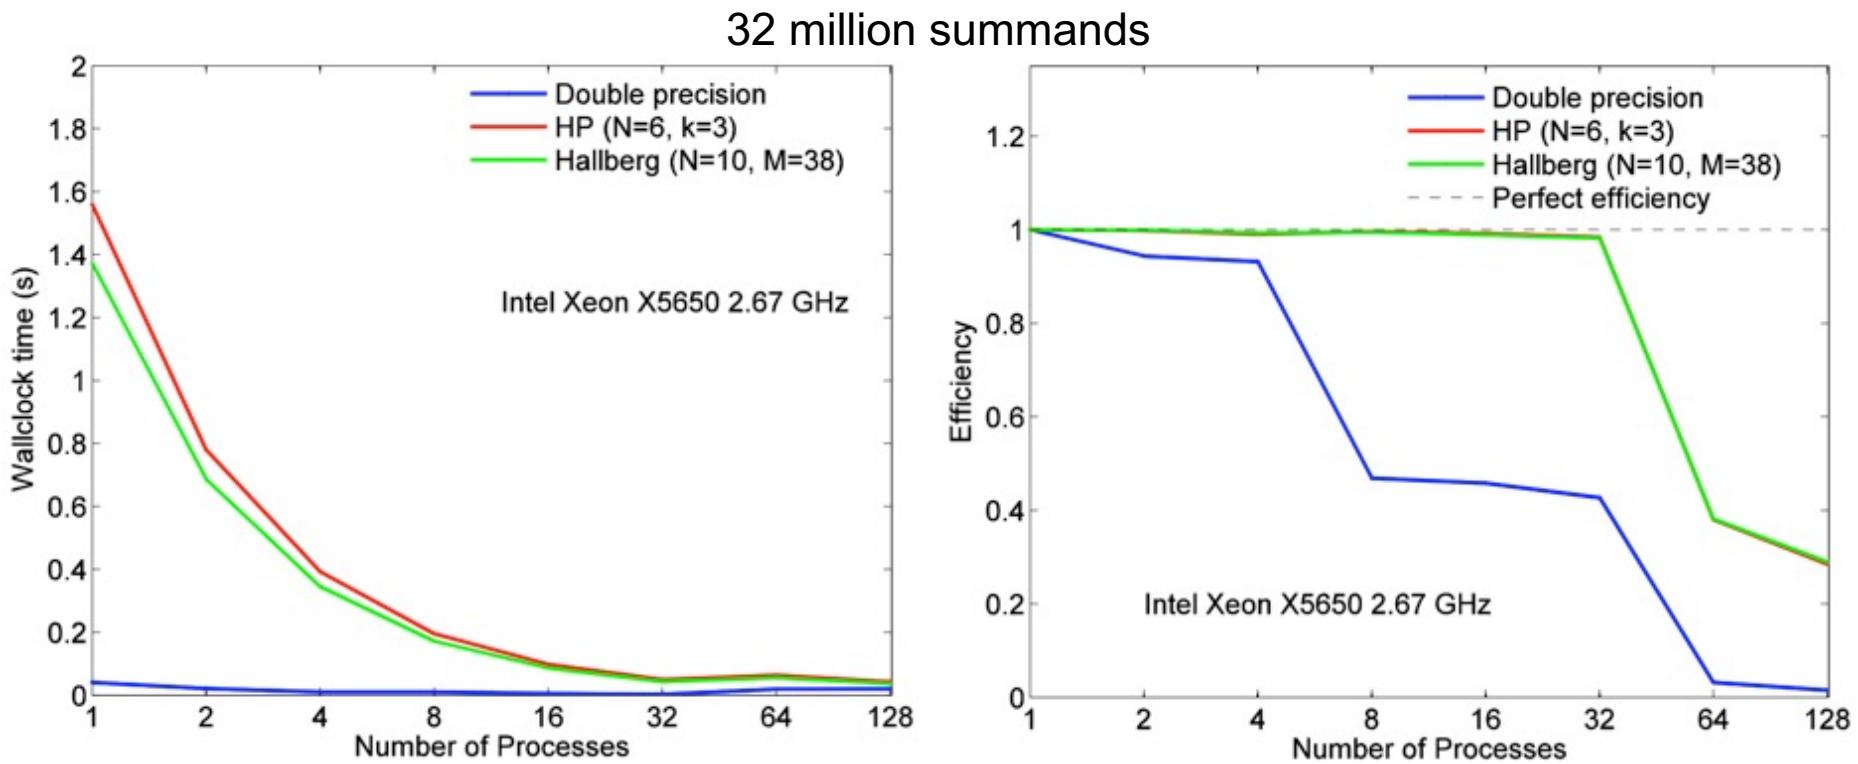
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of OpenMP threads on Xeon



- Higher parallel efficiency of HP & Hallberg sums over double-precision sum

Parallel Efficiency with MPI

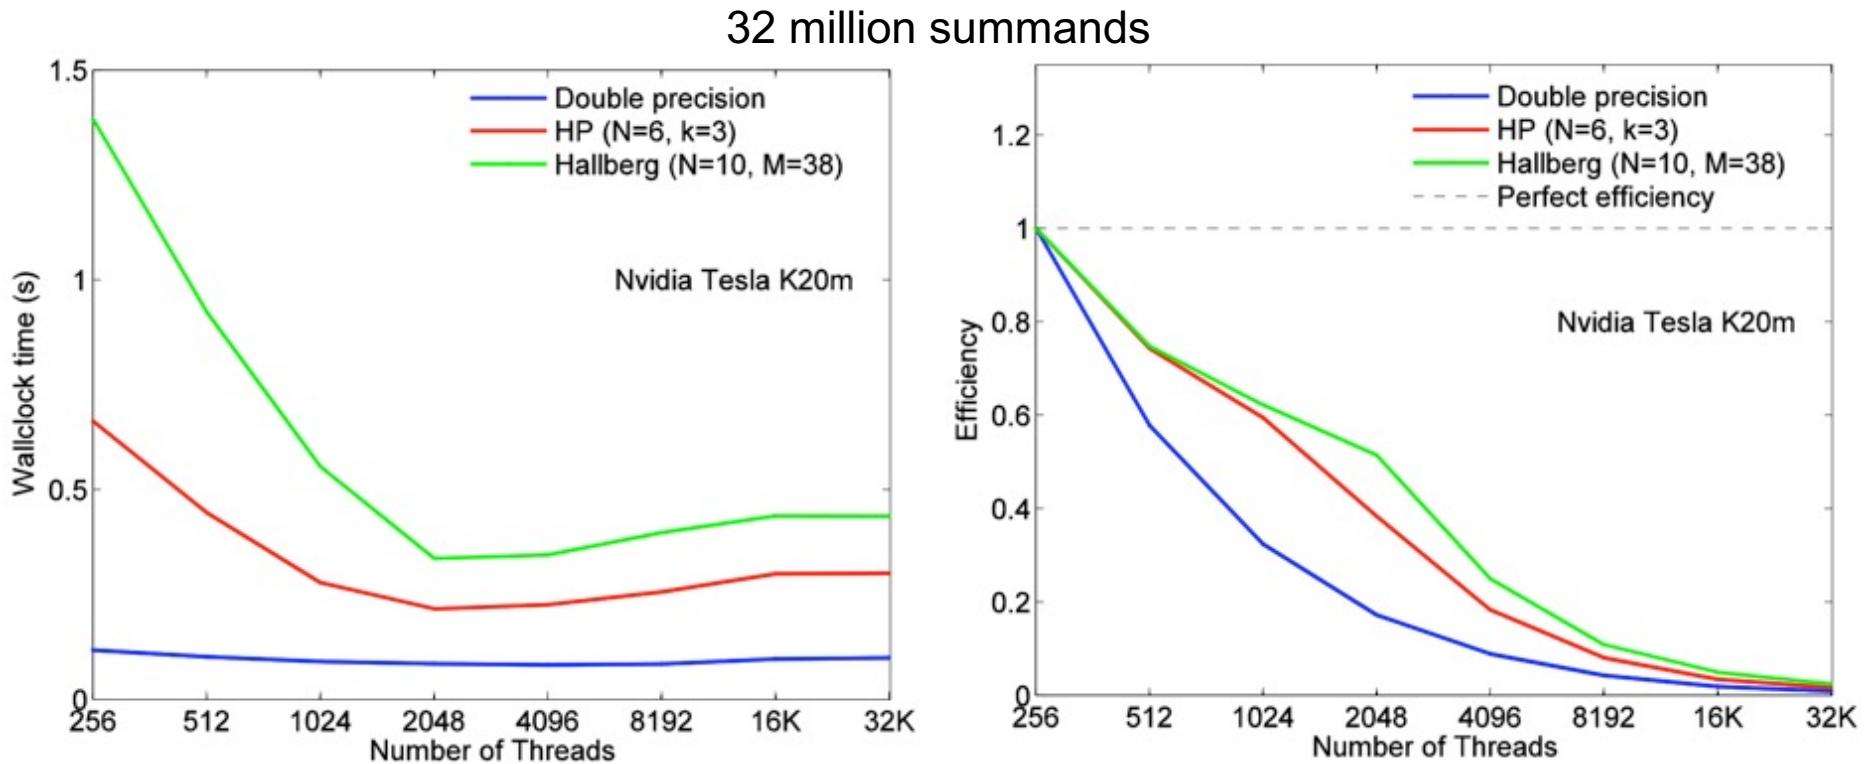
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of MPI processes on Xeon



- Higher parallel efficiency of HP & Hallberg sums over double-precision sum

Parallel Efficiency on GPGPU

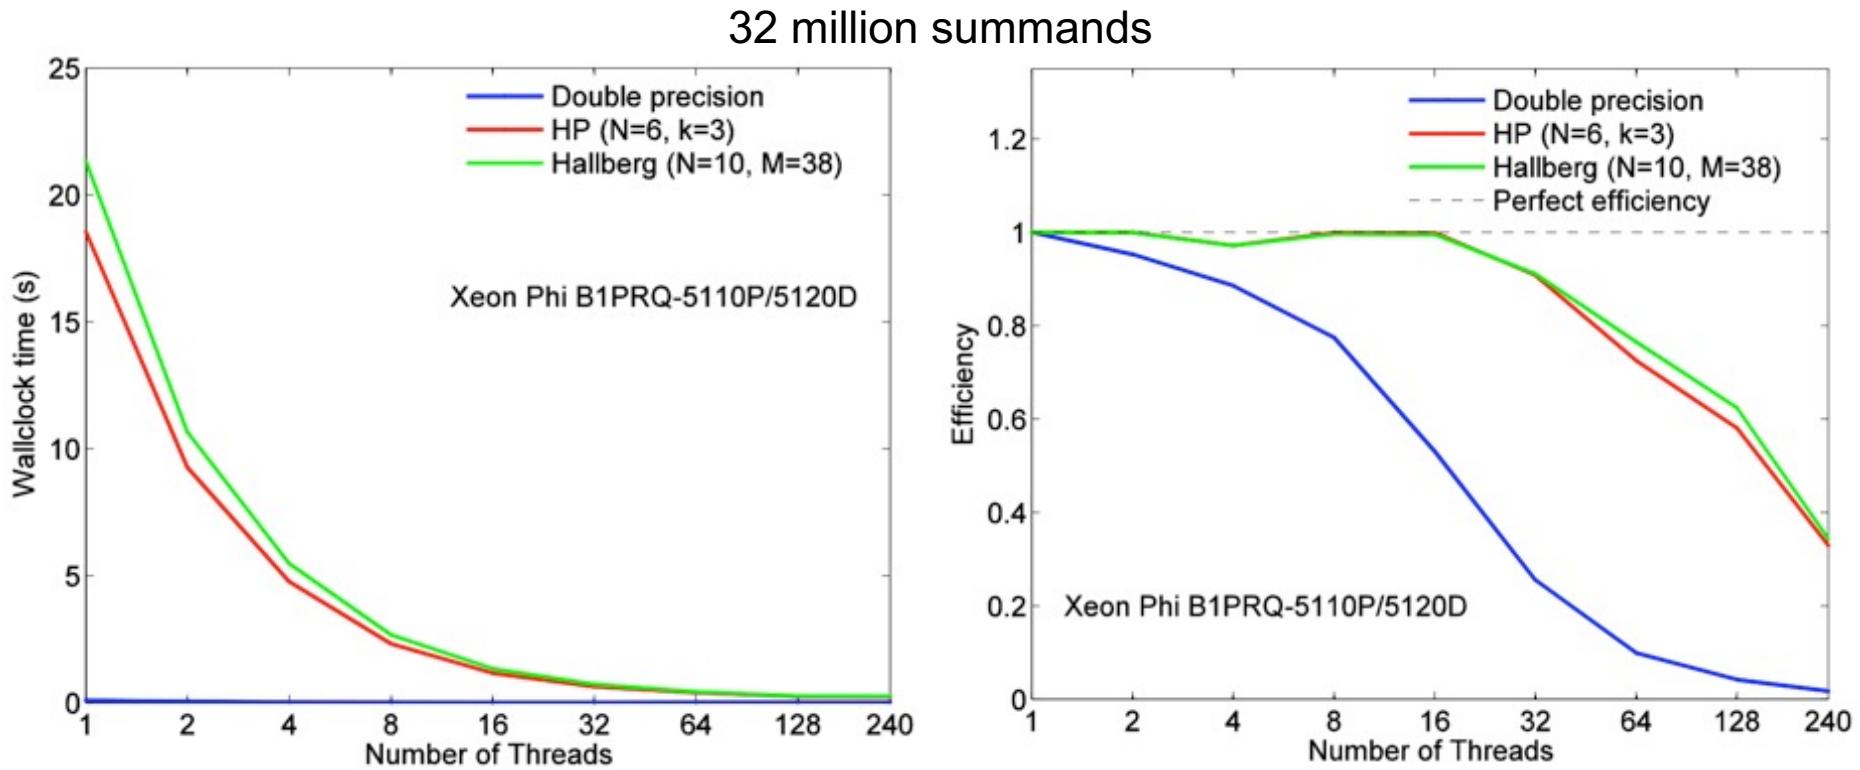
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of CUDA threads on general-purpose graphics processing unit (GPGPU)



- Faster speed of HP sum (7 reads & 6 writes on global memory) over Hallberg sum (11 reads & 10 writes)

Parallel Efficiency on Xeon Phi

- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of threads on Intel Xeon Phi co-processor



- Faster speed of HP sum over Hallberg sum

Large Production Simulations

- **16,661-atom quantum molecular dynamics (QMD) simulation on 786,432 IBM Blue Gene/Q cores suggests a rapid H₂-production technology that is industrially scalable**

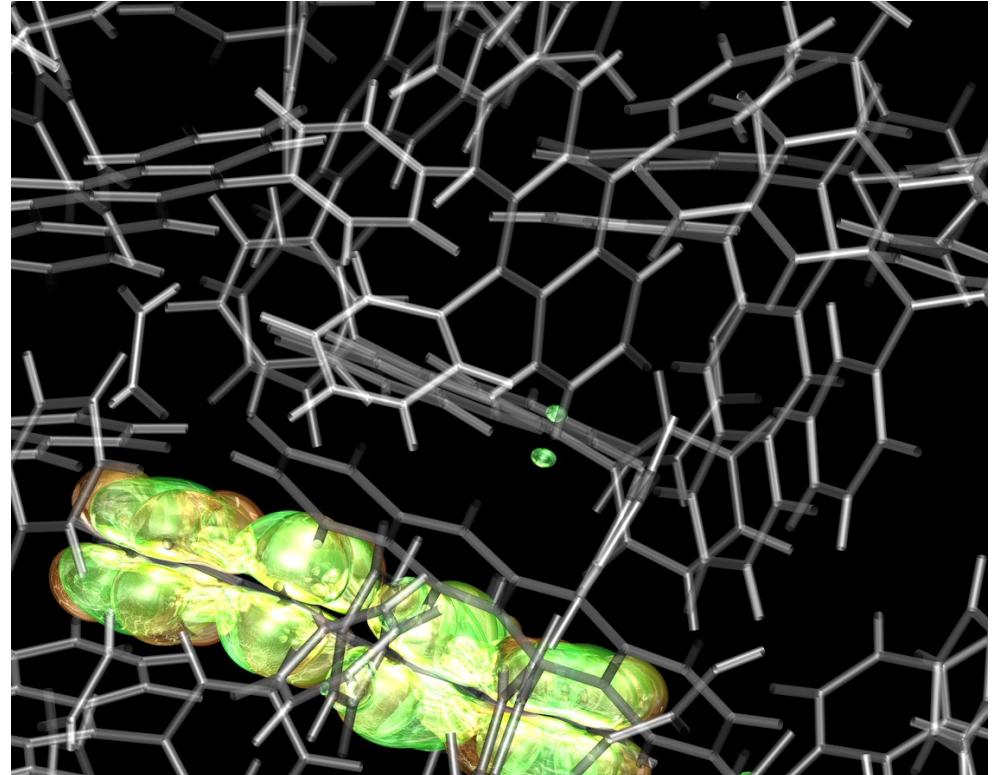
21,140 time steps (129,208 self-consistent-field iterations);
Nano Lett. **14**, 4090 ('14)

- **Up-to 6,400-atom divide-conquer-recombine nonadiabatic QMD simulation reaches experimental time scales from first principles for photoexcitation dynamics**

Appl. Phys. Lett. **102**, 173301 ('13);
Sci. Rep. **5**, 19599 ('16)

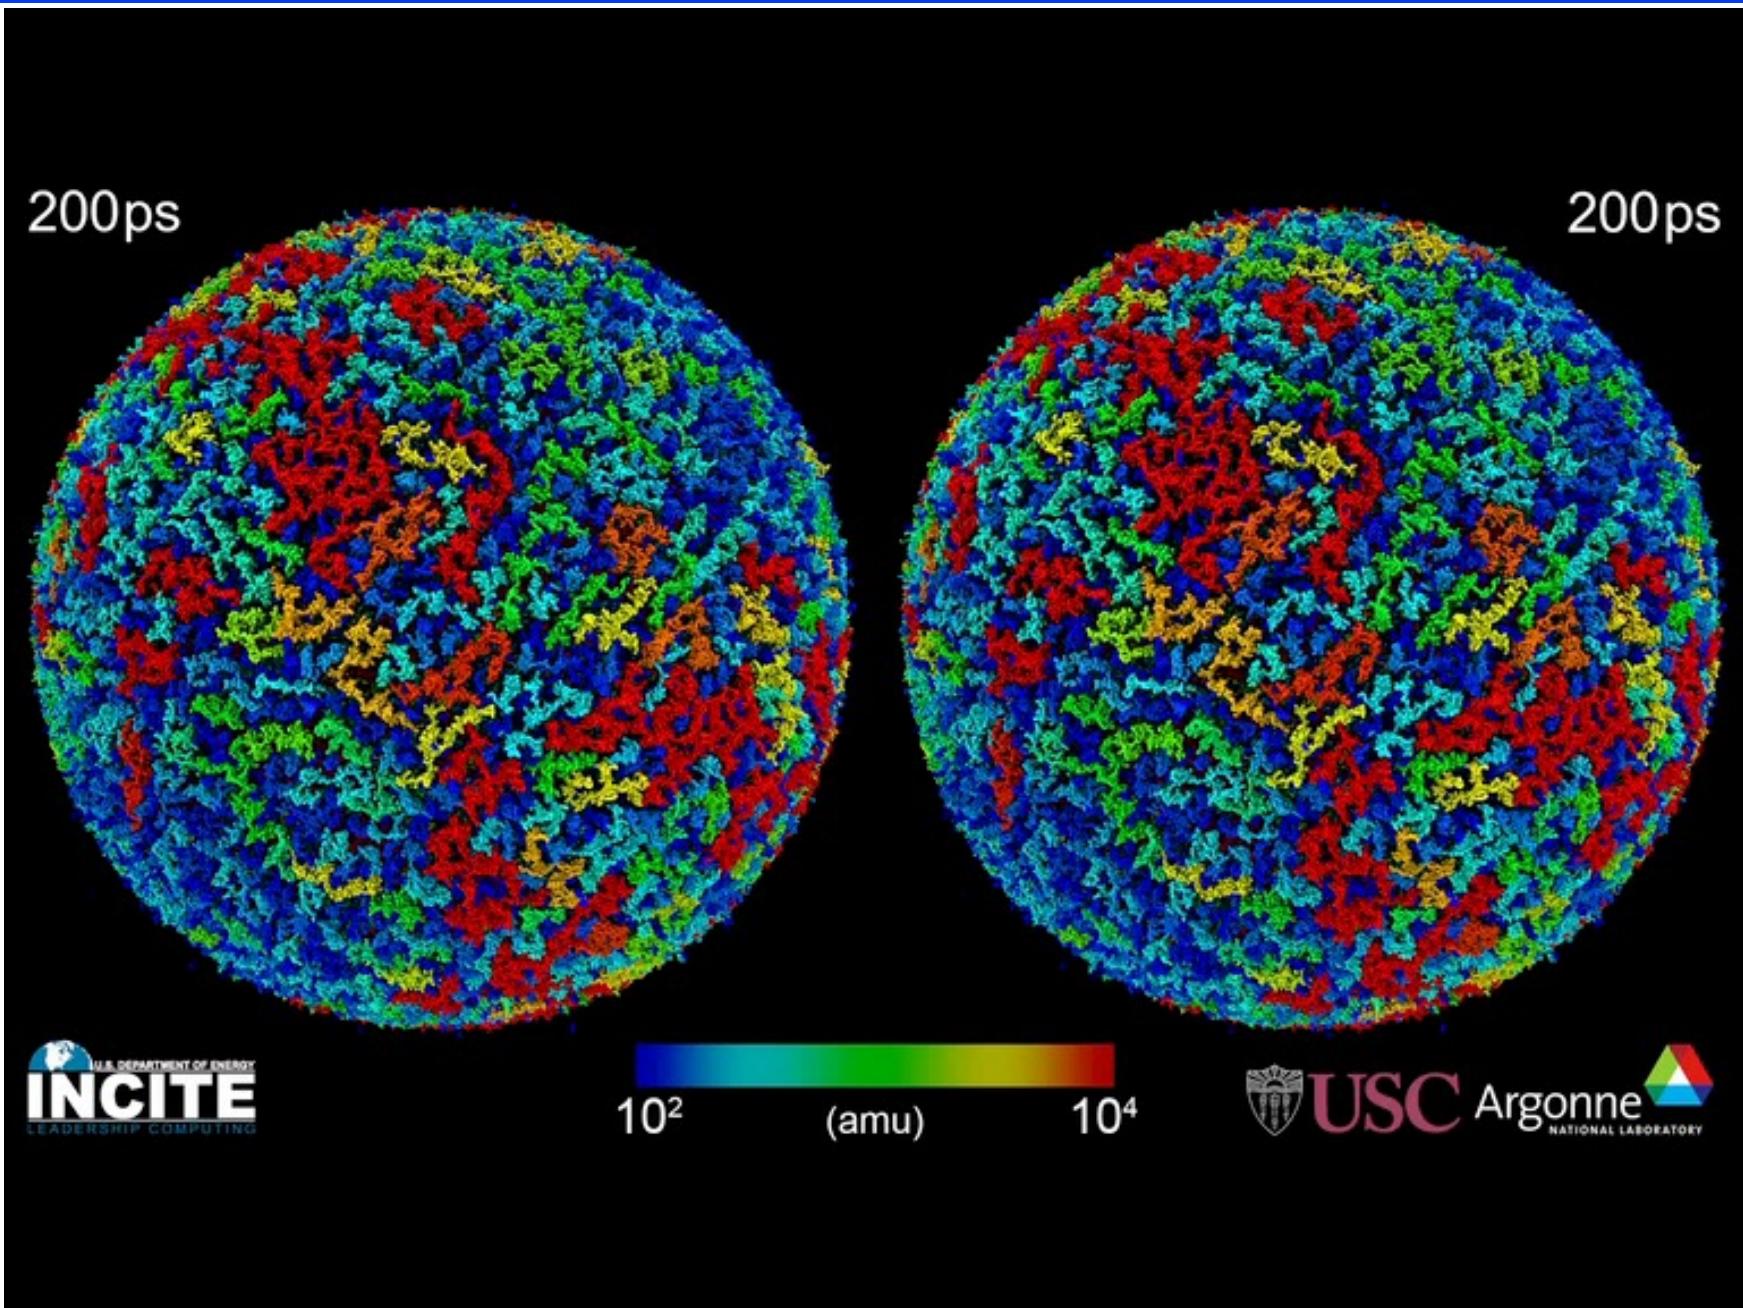
- **112 million-atom reactive molecular dynamics (RMD) simulation on 786,432 IBM Blue Gene/Q cores reveals a simple synthetic pathway to fractal graphene**

Sci. Rep. **6**, 24109 ('16)



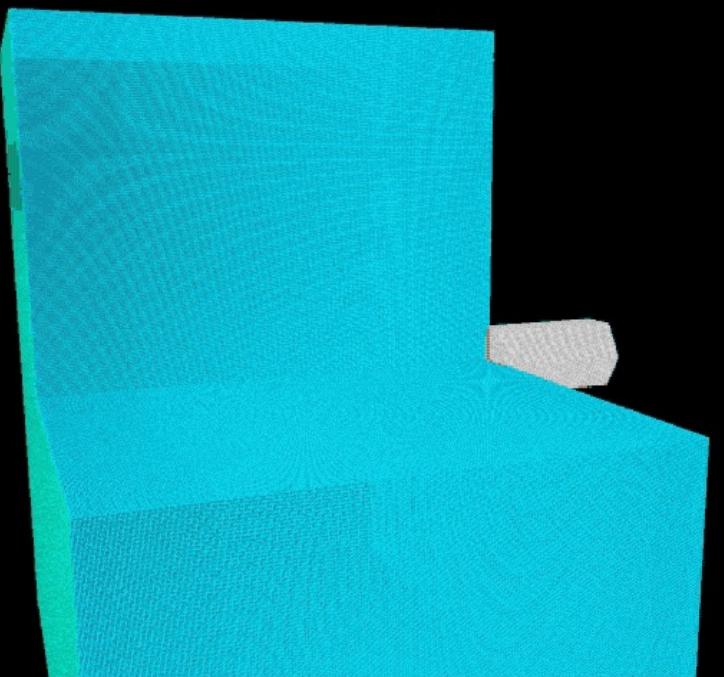
Quasi-electron Quasi-hole

Percolation Transition



Movie made by J. Insley (Argonne)

Billion-Atom Molecular Dynamics

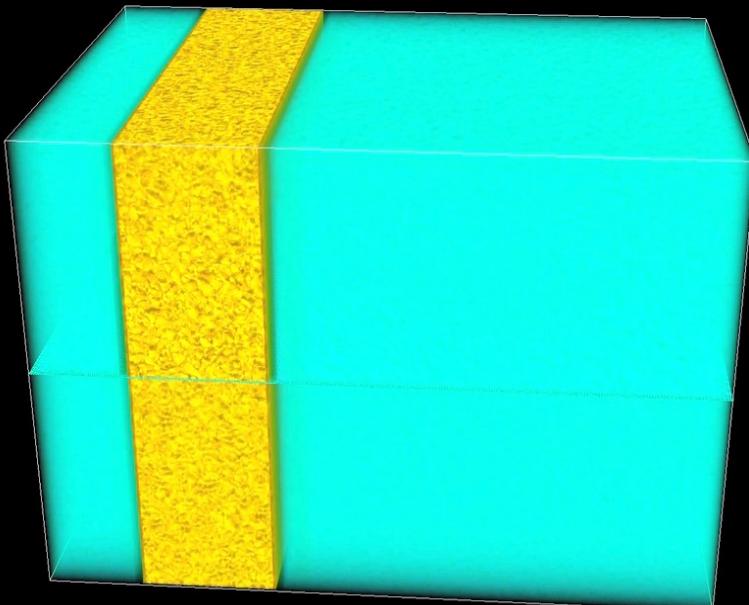


- **Hypervelocity impact on AlN**

P. S. Branicio *et al.*,
Phys. Rev. Lett. **96**, 065502 ('06)

- Shock-induced nanobubble collapse in water near silica surface (67 million core-hours of computing on 163,840 Blue Gene/P cores)

A. Shekhar *et al.*, *Phys. Rev. Lett.* **111**, 184503 ('13)



— 100 nm

Conclusion

1. An order-invariant real-number summation method has been proposed for reproducible parallel computing
2. The proposed method achieves higher computing speed than the previous state-of-the-art for million+ summands on various parallel systems (MPI, OpenMP, CUDA, Xeon Phi)

Thank You

Research supported by
DOE Grant DE-SC0014607

