

# Order-Invariant Real Number Summation: Circumventing Accuracy Loss for Multimillion Summands on Multiple Parallel Architectures

---

---

Patrick E. Small, Rajiv K. Kalia, Aiichiro Nakano, Priya Vasishta

*Collaboratory for Advanced Computing & Simulations*

*Department of Computer Science*

*Department of Physics & Astronomy*

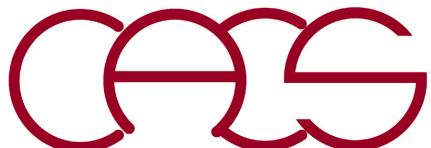
*Department of Chemical Engineering & Materials Science*

*Department of Biological Sciences*

*University of Southern California*

Email: (patrices, rkalia, anakano, priyav)@usc.edu

Proc. IEEE International Parallel & Distributed Processing Symposium, IPDPS, p. 152 ('16)

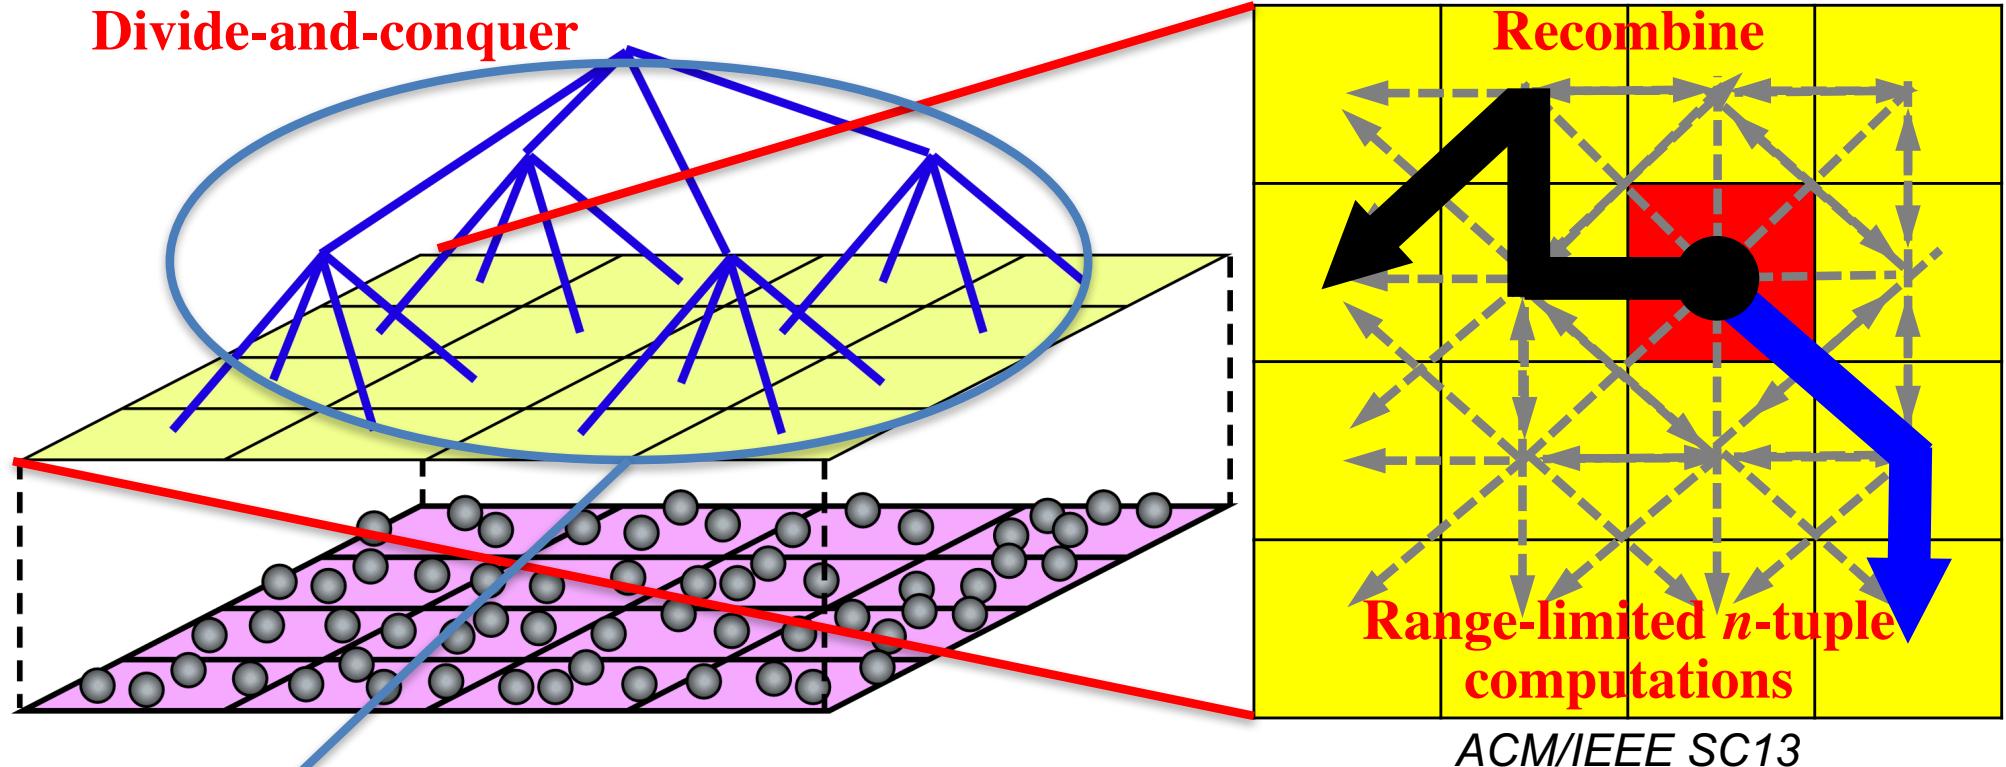


# Exascale Computing Challenge

## 1. Scalability for billion-way parallelism

*J. Chem. Phys.* **140**, 18A529 ('14)  
*IEEE/ACM SC14*  
*IEEE Computer* **48**(11), 33 ('15)

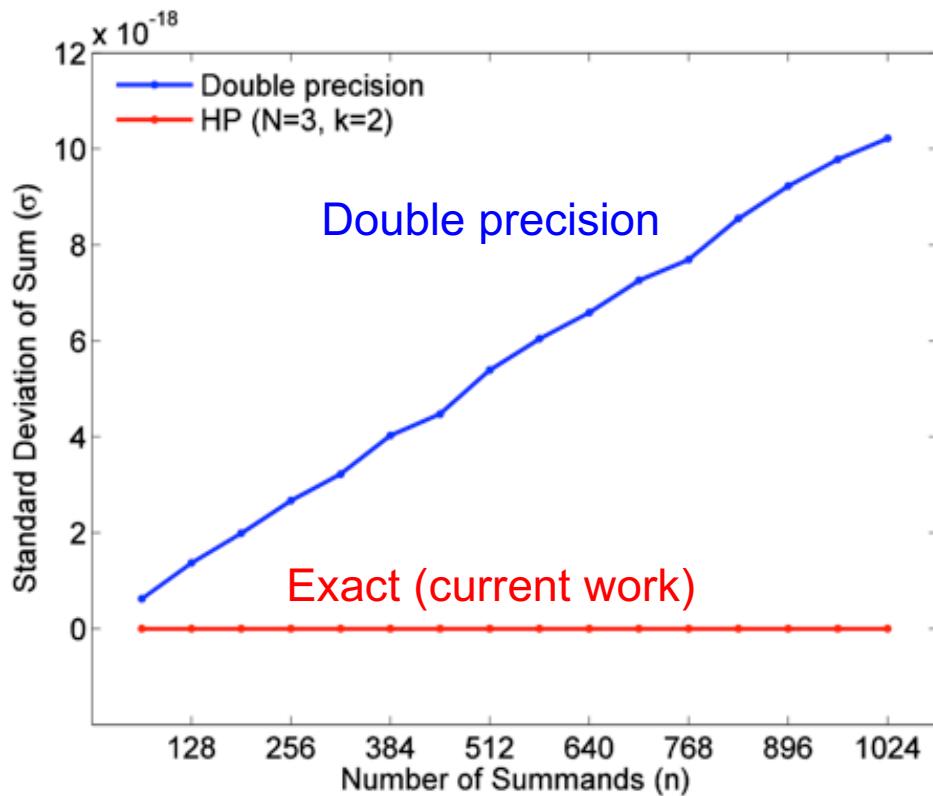
Divide-conquer-recombine (DCR) algorithmic framework  
Metascalable (“design once, scale on future architectures”)



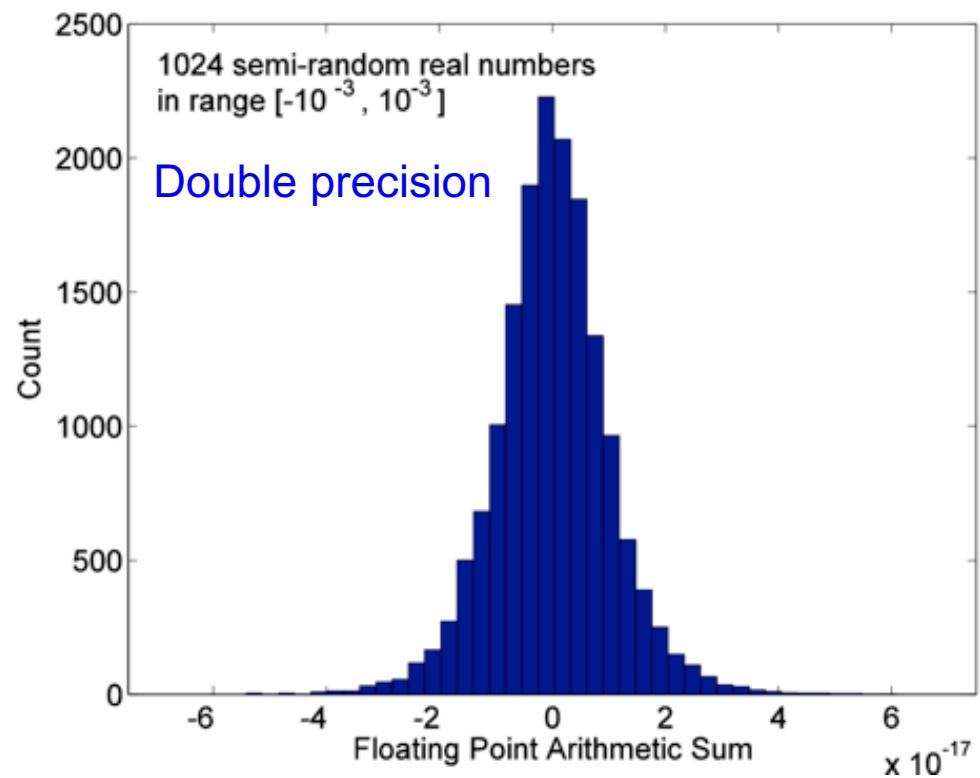
## 2. Reproducibility of real-number summation for multibillion summands in the global sum; double-precision arithmetic began to produce different results on different high-end architectures

# Reproducibility Challenge

- Rounding (truncation) error makes floating-point addition non-associative



Standard deviation of sum with  
random summation orders



Distribution of sum with random  
summation orders

- Sum becomes a random walk across the space of possible rounding error

# High-Precision (HP) Method

---

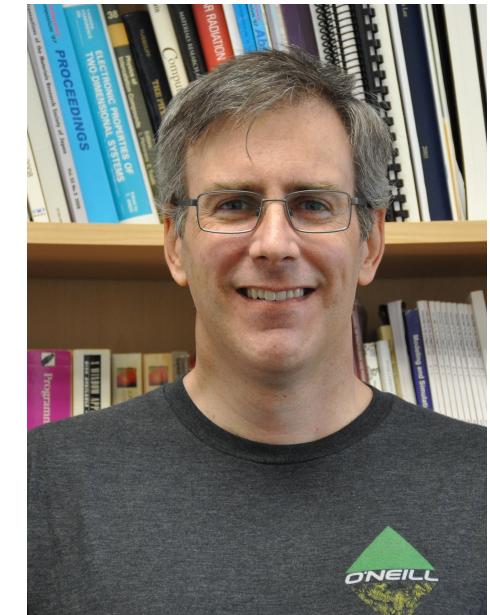
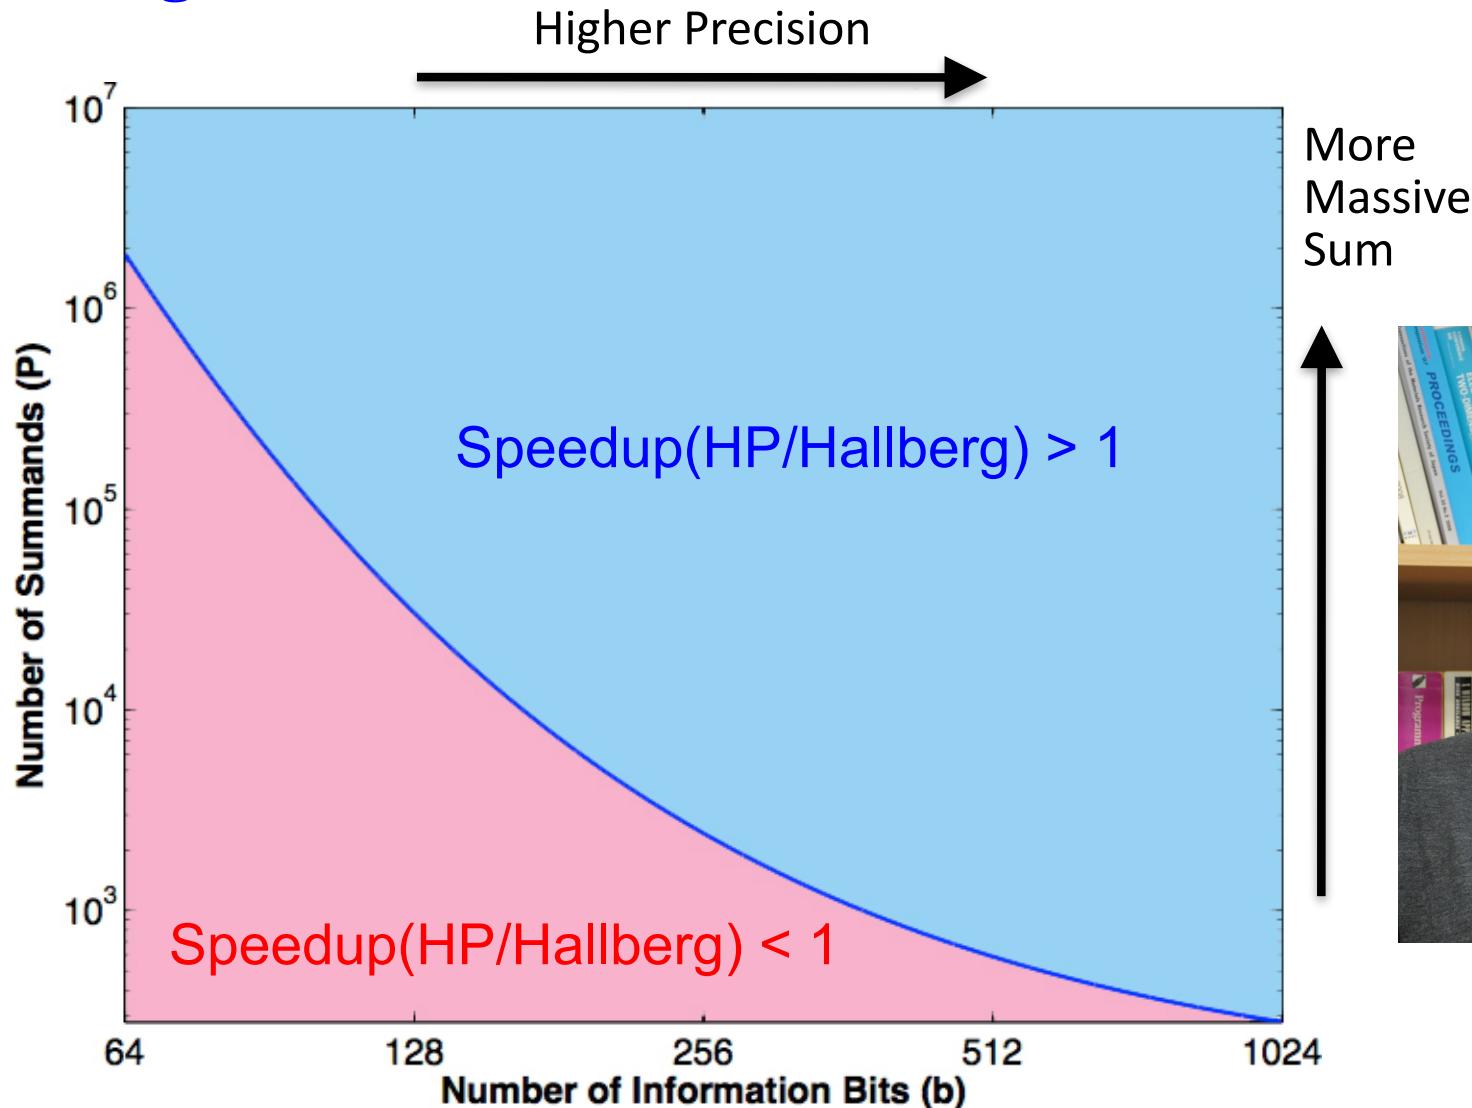
- Propose an extension of the order-invariant, higher-precision intermediate-sum method by Hallberg & Adcroft [*Par. Comput.* **40**, 140 ('14)]
- The proposed variation represents a real number  $r$  using a set of  $N$  64-bit unsigned integers,  $a_i$  ( $i = 0, N-1$ )

$$\begin{aligned} r &= \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)} \\ &= \underbrace{a_0 2^{64(N-k-1)} + \cdots + a_{N-k-1}}_{N-k} + \cdots + \underbrace{a_{N-k} 2^{-64} + \cdots + a_{N-1} 2^{-64k}}_k \end{aligned}$$

- $k$  is the number of 64-bit unsigned integers assigned to represent the fractional portion of  $r$  ( $0 \leq k \leq N$ ), whereas  $N-k$  integers represent the whole-number component
- Negative number is represented by two's complement in integer representation, using only 1 bit

# Performance Projection

- HP sum is faster than Hallberg sum for higher precision & larger numbers of summands



# Detailed IPDPS 2016 Presentation

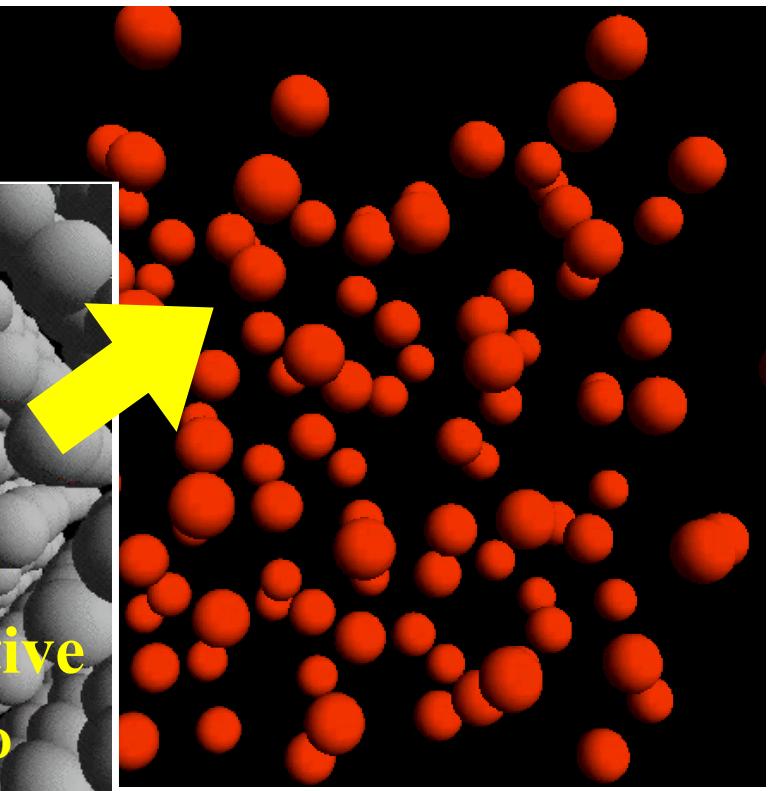
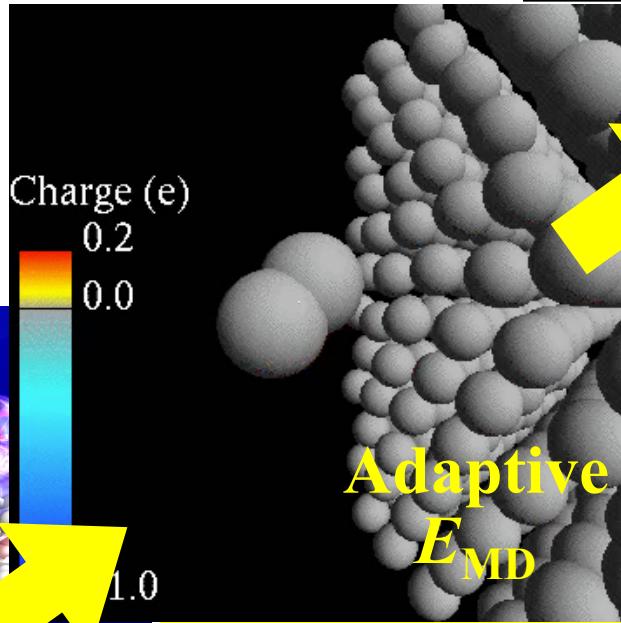
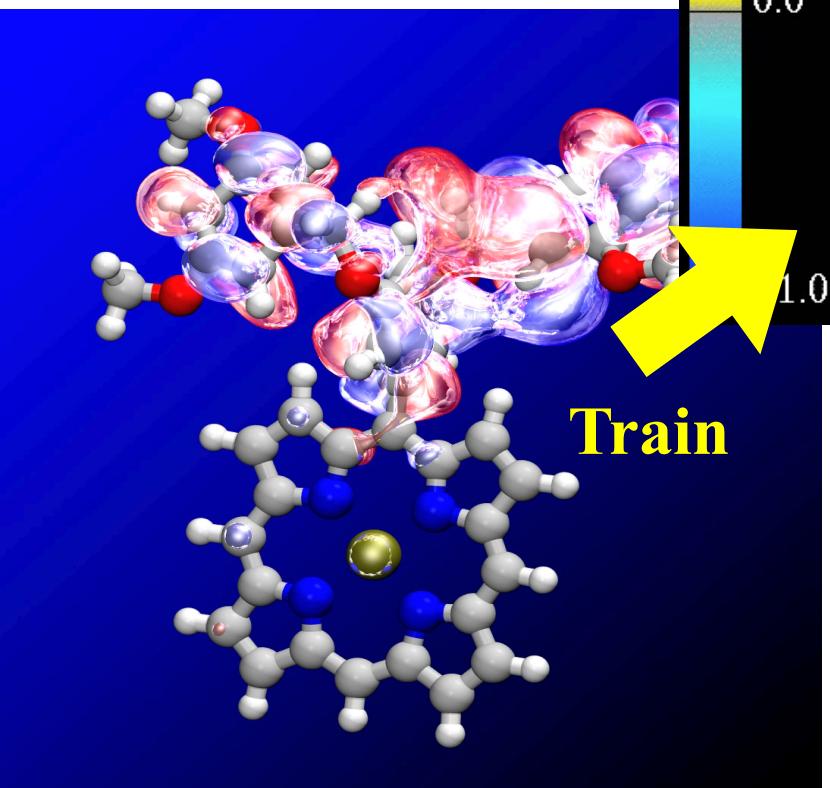
---

# Hierarchy of Atomistic Simulation Methods

## Molecular Dynamics (MD)

### Reactive MD (RMD)

### Nonadiabatic quantum MD (NAQMD)

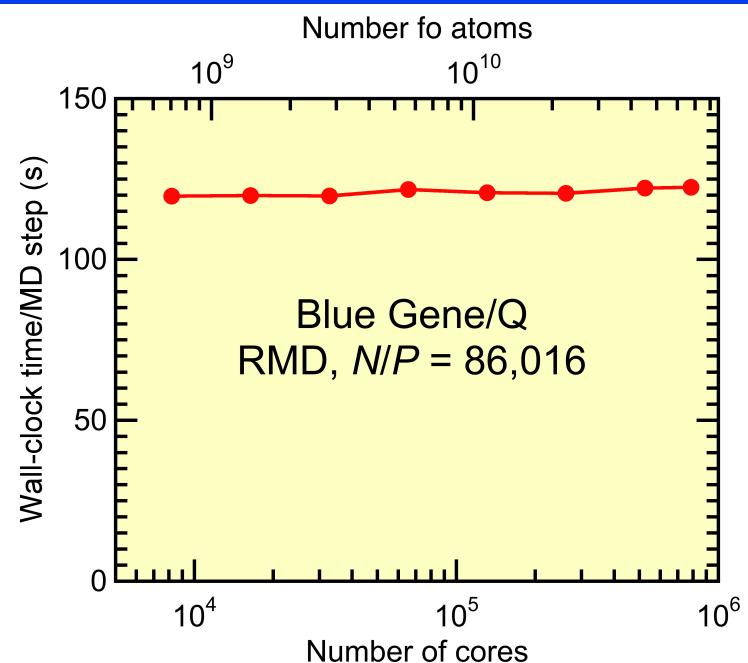
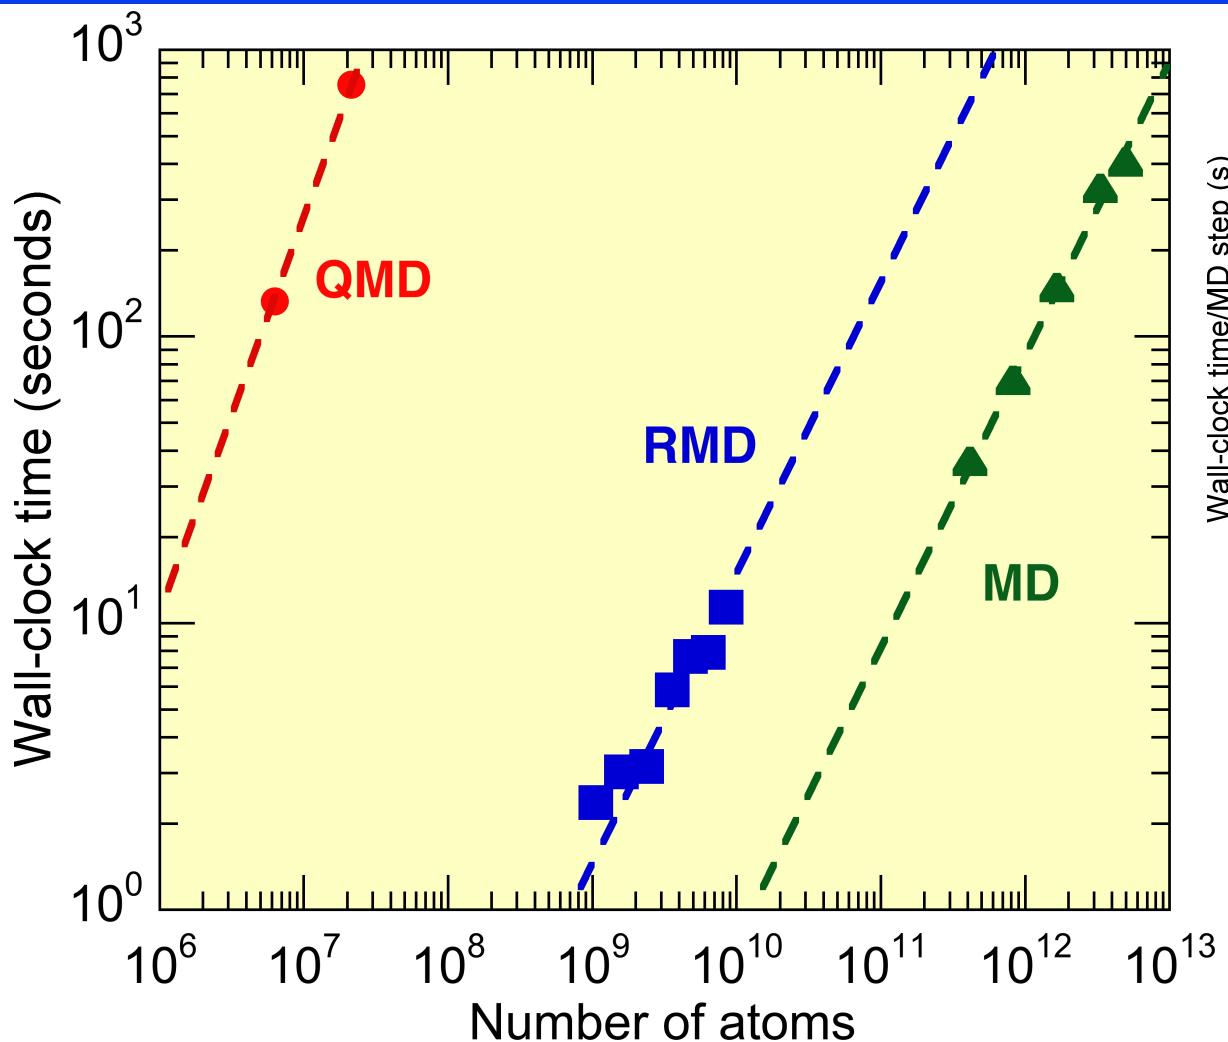


### First principles-based reactive force-fields

- **Reactive bond order  $\{BO_{ij}\}$**   
→ Bond breakage & formation
- **Charge equilibration (QE<sub>q</sub>)  $\{q_i\}$**   
→ Charge transfer

Tersoff, Brenner, Sinnott *et al.*; Streitz & Mintmire *et al.*;  
van Duin & Goddard (ReaxFF)

# Scalable Simulation Algorithm Suite



**QMD (quantum molecular dynamics): DC-DFT**  
**RMD (reactive molecular dynamics): F-ReaxFF**  
**MD (molecular dynamics): MRMD**

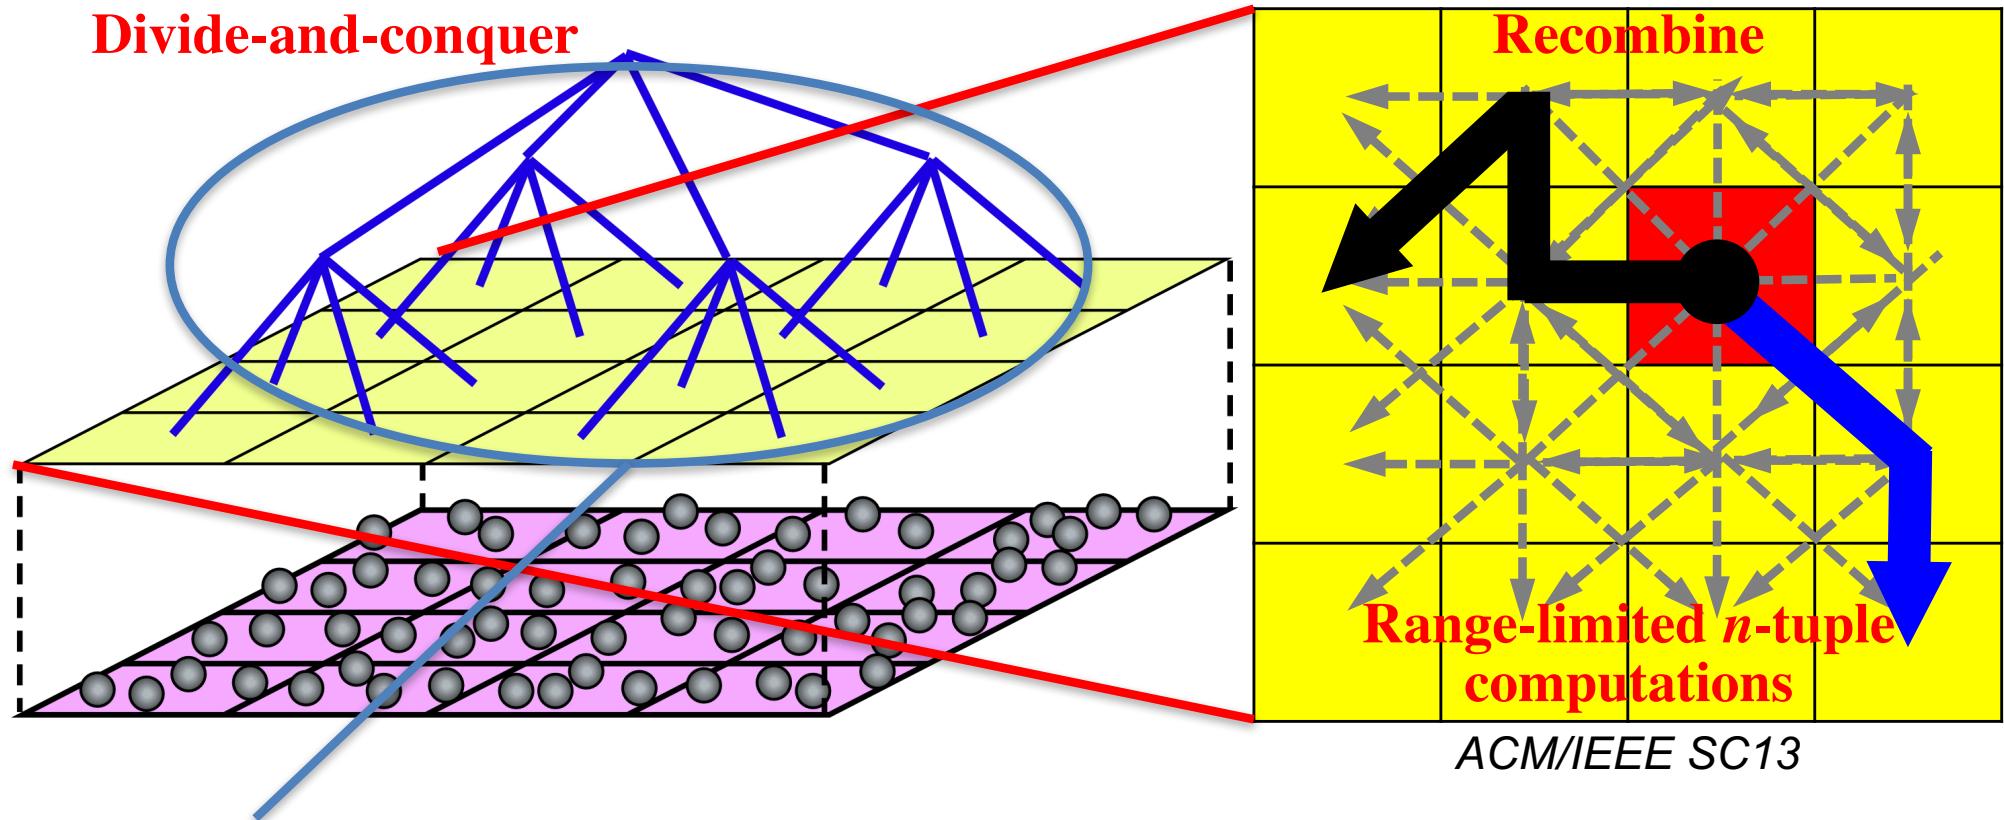
- 4.9 trillion-atom space-time multiresolution MD (MRMD) of  $\text{SiO}_2$
- 67.6 billion-atom fast reactive force-field (F-ReaxFF) RMD of RDX
- 39.8 trillion grid points (50.3 million-atom) DC-DFT QMD of SiC  
parallel efficiency 0.984 on 786,432 Blue Gene/Q cores

# Exascale Computing Challenge

## 1. Scalability beyond million-way parallelism

*J. Chem. Phys.* **140**, 18A529 ('14)  
*IEEE/ACM SC14*  
*IEEE Computer* **48**(11), 33 ('15)

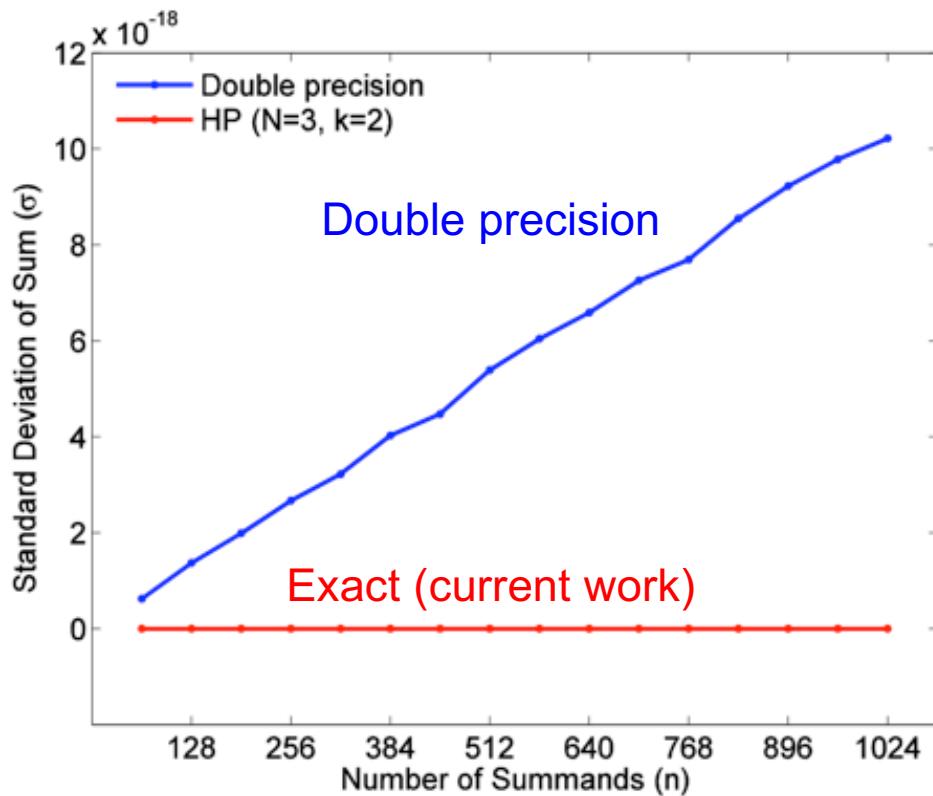
Divide-conquer-recombine (DCR) algorithmic framework  
Metascalable (“design once, scale on future architectures”)



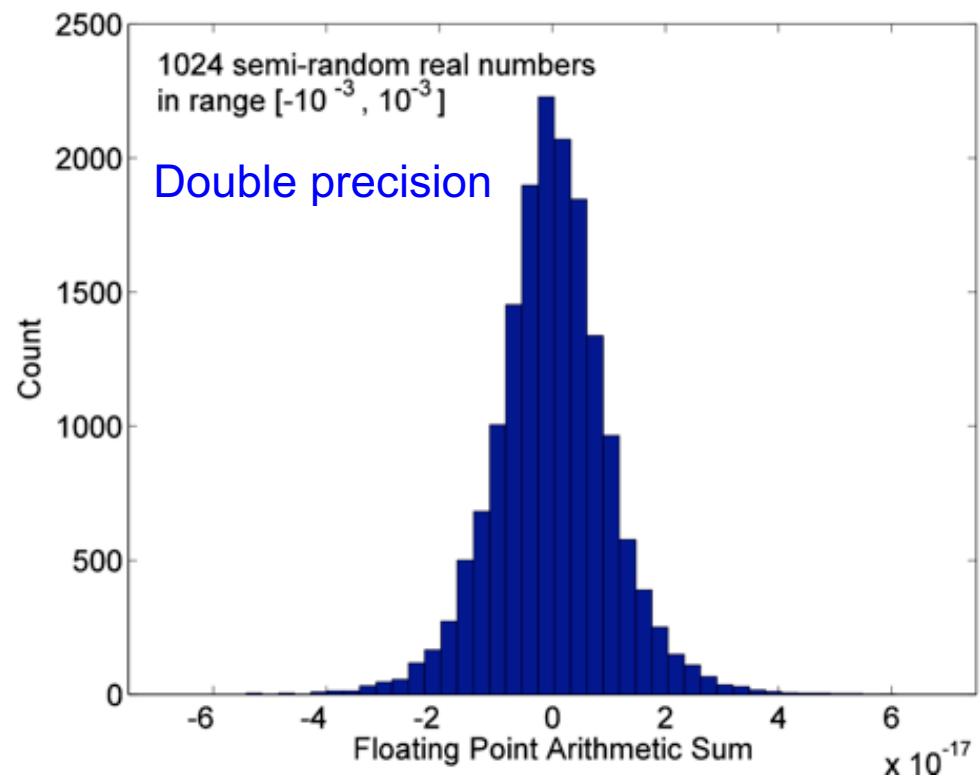
## 2. Reproducibility of real-number summation for multimillion summands & beyond in the global sum; double-precision arithmetic began to produce different results on different high-end architectures

# Reproducibility Challenge

- Rounding (truncation) error makes floating-point addition non-associative



Standard deviation of sum with  
random summation orders



Distribution of sum with random  
summation orders

- Sum becomes a random walk across the space of possible rounding error

# Related Works

---

---

- **General-purpose arbitrary precision arithmetic**  
[GNU-MPL '12]
  - Extensive computation & memory usage
- **Error-compensation methods**
  - > **Error-free transformation for tracking residuals**  
[Priest, '91, Higham '93, Rump '09, Demmel '13]
    - Complex implementation
  - > **Summation reordering for minimizing error**  
[Hel '01]
    - Prohibitive at large scales
- **Hardware solutions**  
[Gustafson '15]
  - Not available yet
- **Higher-precision intermediate sums**  
[He '01, Hallberg '14]
  - Simple implementation, low overhead

# Contributions

---

---

- Propose an extension of the order-invariant, higher-precision intermediate-sum method by Hallberg & Adcroft [*Par. Comput.* **40**, 140 ('14)]:
  - (1) Improves performance\* for large ( $> 10^6$ ) number of summands
  - (2) Eliminates the aliasing problem of the original method
- The new method outperforms the previous state-of-the-art for large problems involving million+ summands on broad systems (MPI, OpenMP, CUDA/GPU, Xeon Phi)

\*Performance is defined as the computational speed

# Hallberg Order-Invariant Sum

---

- **Integer representation with higher accuracy:** Represent a real number  $r$  using a set of  $N$  64-bit signed integers,  $a_i$  ( $i = 0, N-1$ );  $M (< 63)$  is a positive integer

$$r = \sum_{i=0}^{N-1} a_i 2^{\left(i - \frac{N}{2}\right)M} = 2^{-NM/2} (a_0 + a_1 2^M + a_2 2^{2M} + \dots)$$

- **Order-invariant parallel sum:** Two real numbers are added by summing  $N$  pairs of corresponding integers concurrently
- **Carry out (potential sequential dependence):** When any of the integer additions exceeds  $2^M$ , carry out must be added to the next integer in the set
- **Carry-overhead reduction:** Carry operations are avoided up to  $P = 2^{63-M}-1$  summands to expose high parallelism

# Drawback of Hallberg Sum

---

- **Overhead:** Not all integer bits serve to provide real-number precision;  $63 \cdot M$  bits per integer are dedicated to book-keeping
- **Aliasing:** Multiple integer representations could represent the same real number
- **Normalization & sum overheads** to convert the integer representation back to real

# High-Precision (HP) Method

---

- The proposed variation of Hallberg method represents a real number  $r$  using a set of  $N$  64-bit unsigned integers,  $a_i$  ( $i = 0, N-1$ )

$$r = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)}$$
$$= \underbrace{a_0 2^{64(N-k-1)} + \cdots + a_{N-k-1}}_{N-k} + \underbrace{a_{N-k} 2^{-64} + \cdots + a_{N-1} 2^{-64k}}_k$$

- $k$  is the number of 64-bit unsigned integers assigned to represent the fractional portion of  $r$  ( $0 \leq k \leq N$ ), whereas  $N-k$  integers represent the whole-number component
- Negative number is represented by two's complement in integer representation, using only 1 bit

# HP Algorithm (1): Conversion

---

- **Simple procedure:** A single pass converts a double-precision number  $r$  to HP integers  $a_i$  & translates them to two's complement

$$r = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)}$$

```
dtmp = fabs(r)*264*(N-k-1);
isneg = (r < 0.0);
for (i=0; i<N-1; i++) {
    itmp = (uint64_t)dtmp;
    dtmp = (dtmp - (double)itmp)*264;
    a[i] = (isneg) ? ~itmp + (dtmp<=0.0) : itmp;
}
a[N-1] = (isneg) ? ~(uint64_t)dtmp + 1 : (uint64_t)dtmp;
```

- Inverse of this algorithm converts HP number back to double-precision

# HP Algorithm (2): Addition

---

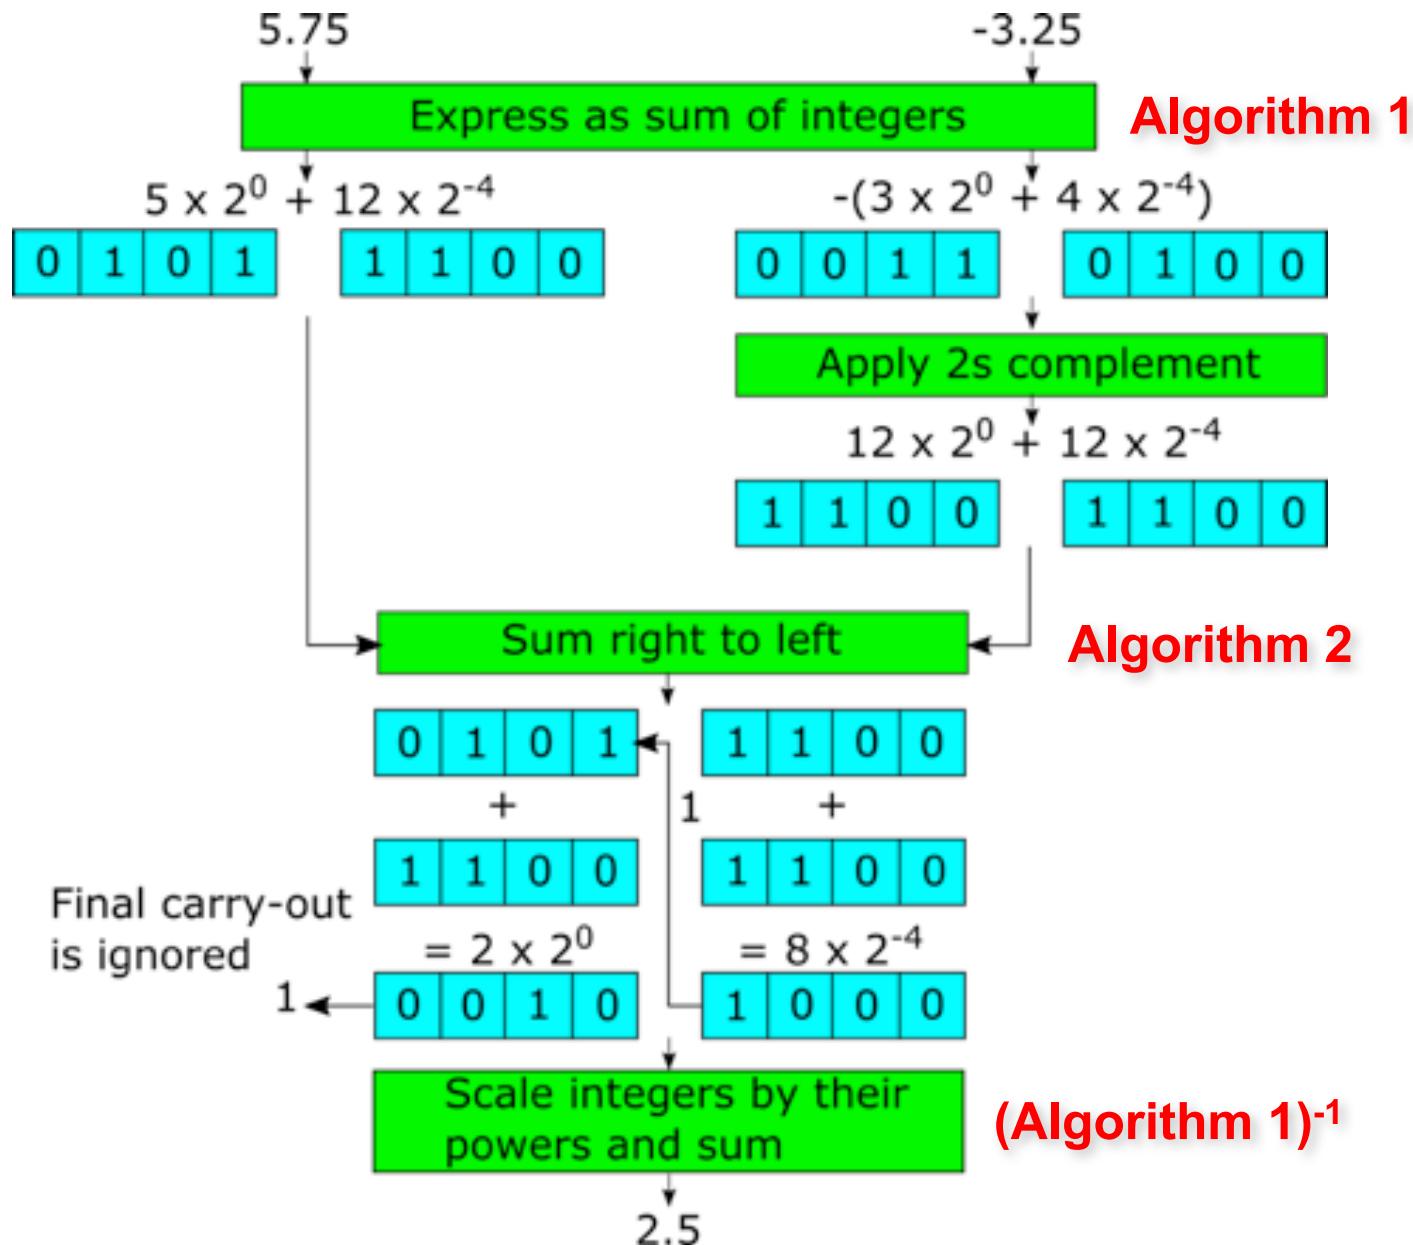
- **Addition of two HP numbers,  $a \leftarrow a + b$**

$$\begin{cases} r_1 = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)} \\ r_2 = \sum_{i=0}^{N-1} b_i 2^{64(N-k-i-1)} \end{cases}$$

```
a[N-1] = a[N-1]+b[N-1];
co = (a[N-1]<b[N-1]);
for (i=N-2; i>=1; i--) {
    a[i] = a[i]+b[i]+co;
    co = (a[i]==b[i]) ? co : (a[i]<b[i]);
}
a[0] = a[0]+b[0]+co;
```

- **Overflow of the sum is detected by comparing the signs of the summands with that of the sum**

# HP Sum: Example



# Representation Power

- Maximum range & smallest representable HP number

$$r_{\text{HP}} = \sum_{i=0}^{N-1} a_i 2^{64(N-k-i-1)}$$

<b><math>N</math></b>	<b><math>k</math></b>	<b>Bits</b>	<b>Maximum range</b>	<b>Smallest number</b>
2	1	128	$\pm 9.223372 \times 10^{18}$	$5.421011 \times 10^{-20}$
3	2	192	$\pm 9.223372 \times 10^{18}$	$2.938736 \times 10^{-39}$
6	3	256	$\pm 3.138551 \times 10^{57}$	$1.593092 \times 10^{-58}$
8	4	512	$\pm 5.789604 \times 10^{76}$	$8.636169 \times 10^{-78}$

- Equivalency with Hallberg representation power

$$r_{\text{Hallberg}} = \sum_{i=0}^{N-1} a_i 2^{(i-N/2)M}$$

<b><math>N</math></b>	<b><math>M</math></b>	<b>Precision bits</b>	<b>Max summands</b>
10	52	520	2048
12	43	516	1 M
14	37	518	64 M

# HP Method: Properties

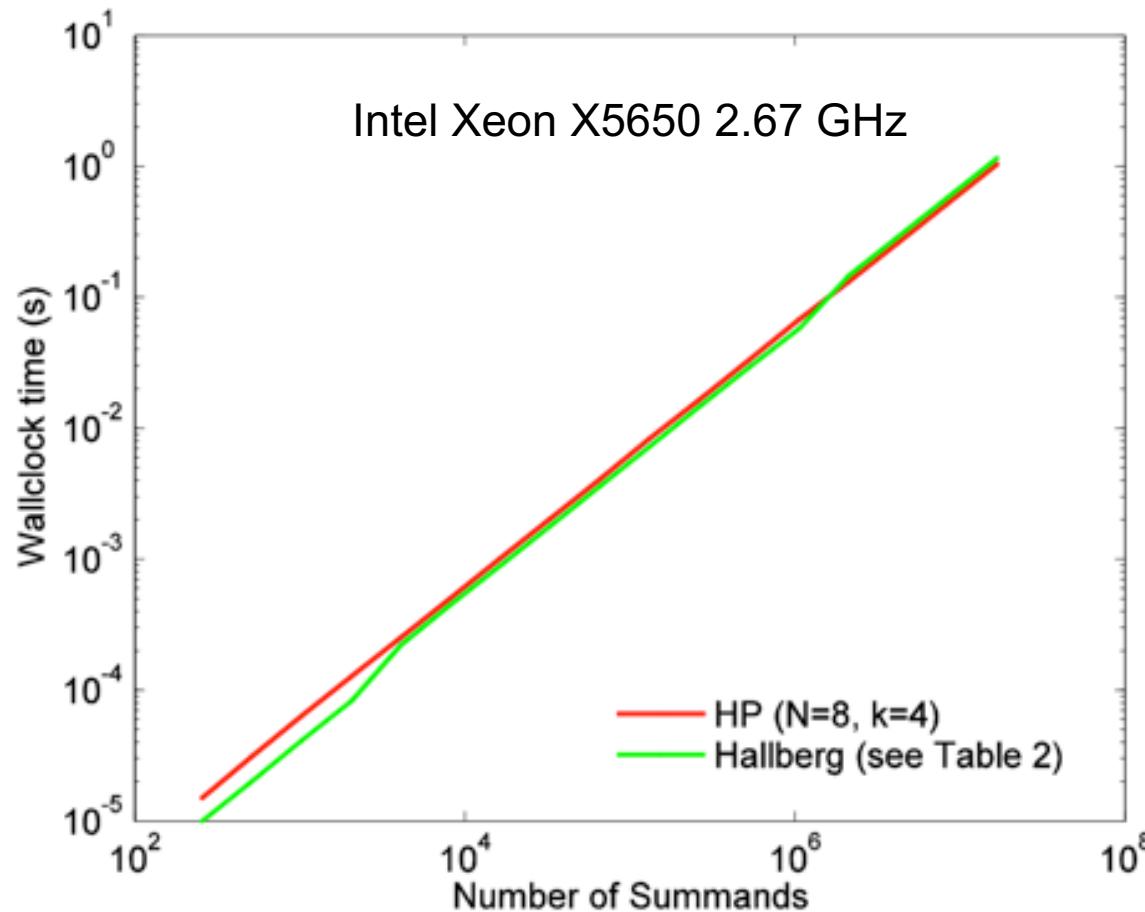
---

---

- Invariance of sum with respect to both summation order & architecture is guaranteed with appropriate setting of  $N$  &  $k$  to provide sufficient accuracy
- Overflow & underflow can be readily detected at runtime at double-precision (DP)-to-HP conversion, HP-sum & HP-to-DP conversion steps
- Atomicity of addition (which is essential for multithreading) is guaranteed using only the widely available compare-&-swap (CAS) synchronization primitive

# Performance

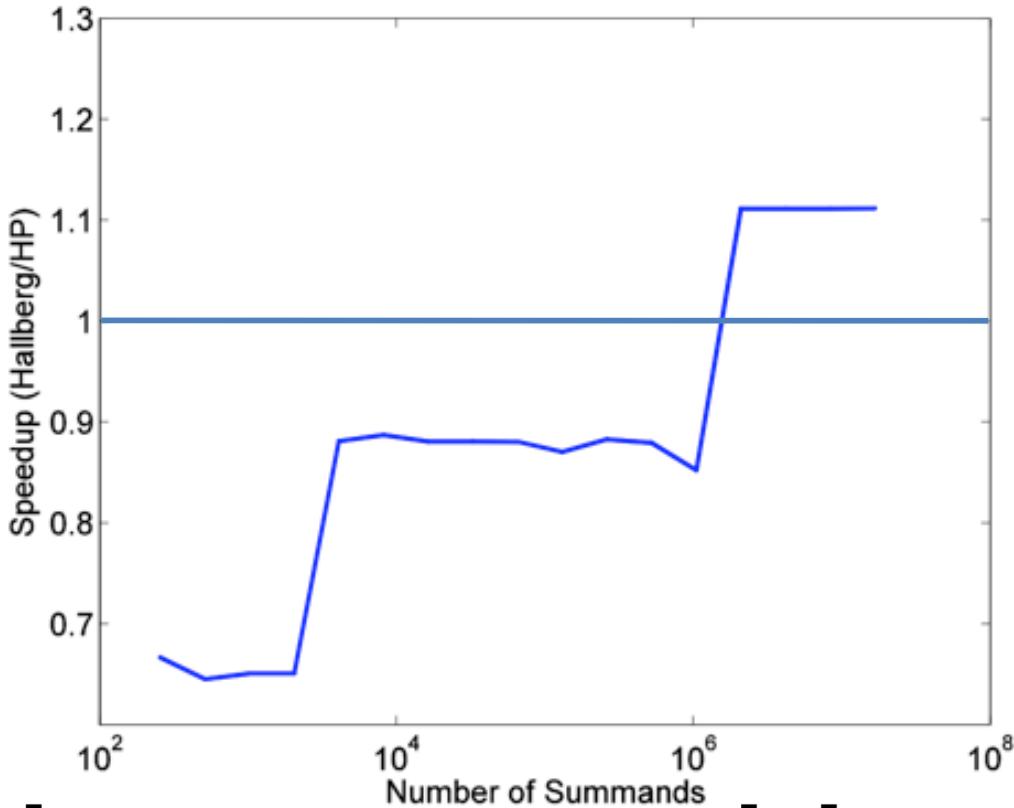
- Computing time of real-number sum using the current (HP) & Hallberg methods as a function of # of summands



- HP sum is faster than Hallberg sum over million summands

# Performance Analysis

- Speedup of HP sum over Hallberg sum as a function of the number of summands

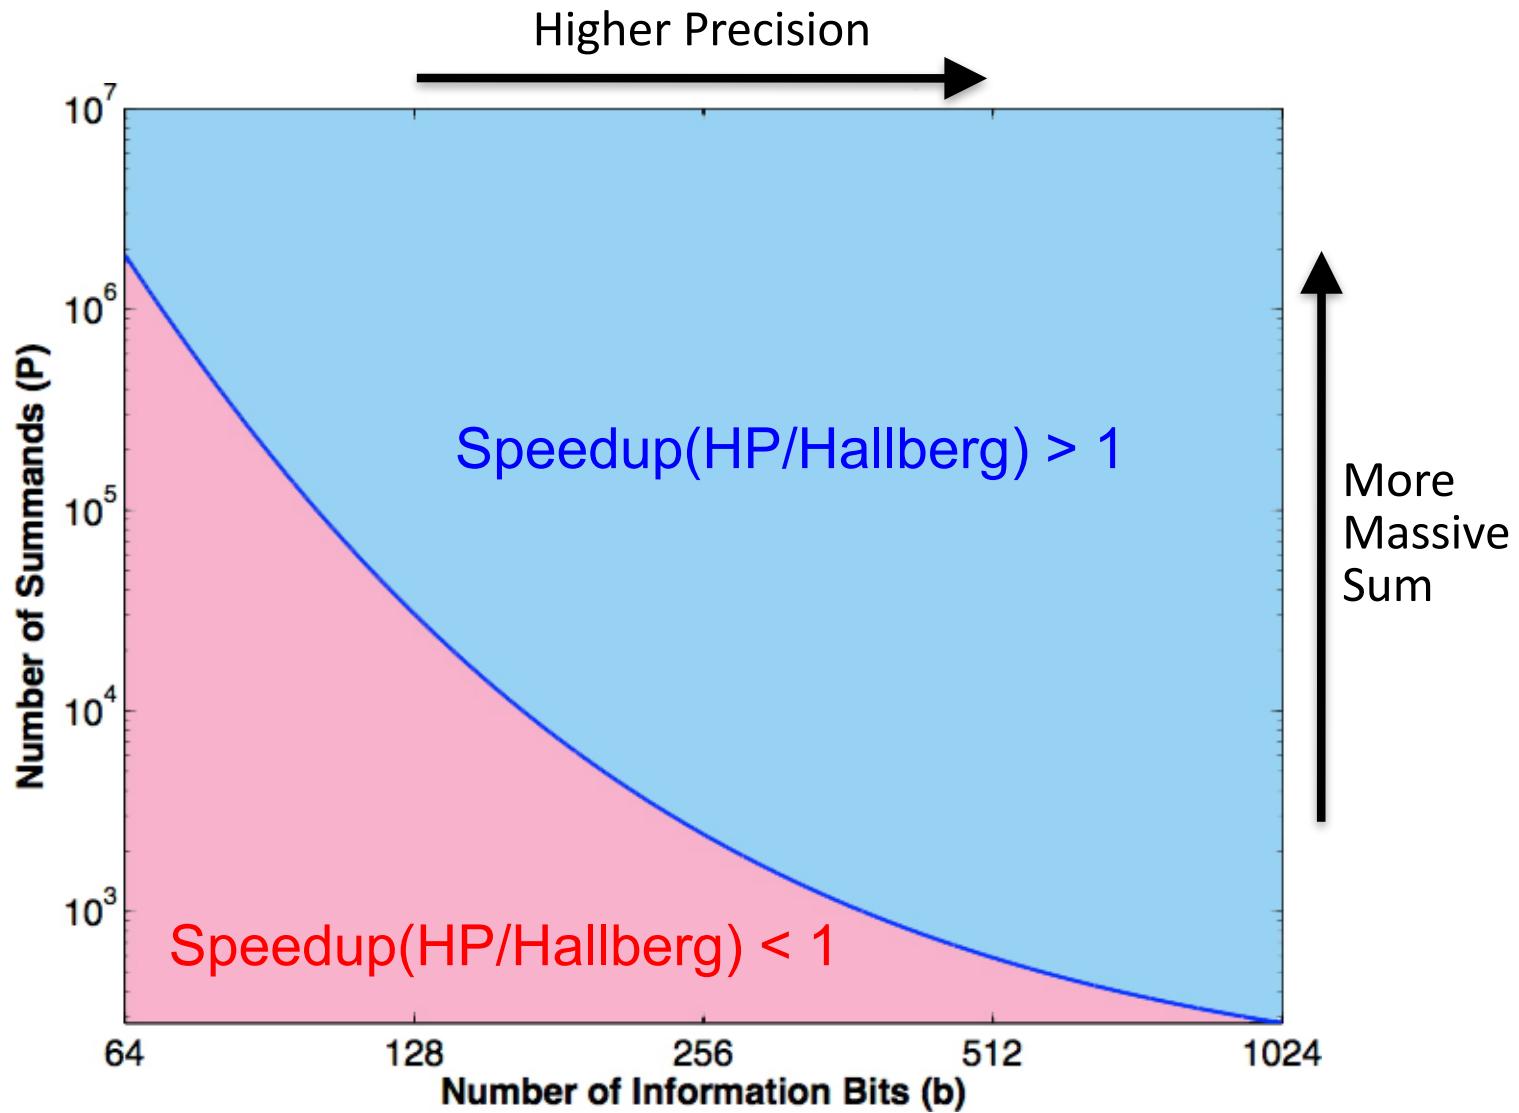


$$T_{\text{HP}} = c_{\text{HP}} \left\lceil \frac{b+1}{64} \right\rceil \quad T_{\text{Hallberg}} = c_{\text{Hallberg}} \left\lceil \frac{b}{M} \right\rceil \quad b: \text{Precision bit count}$$

$$\text{Speedup} \sim \frac{c_{\text{Hallberg}}}{c_{\text{HP}}} \frac{64b}{M(b+65)} \geq \frac{32c_{\text{Hallberg}}}{c_{\text{HP}}} \frac{1}{M} \quad 2^{63-M} \propto \#\text{summands}$$

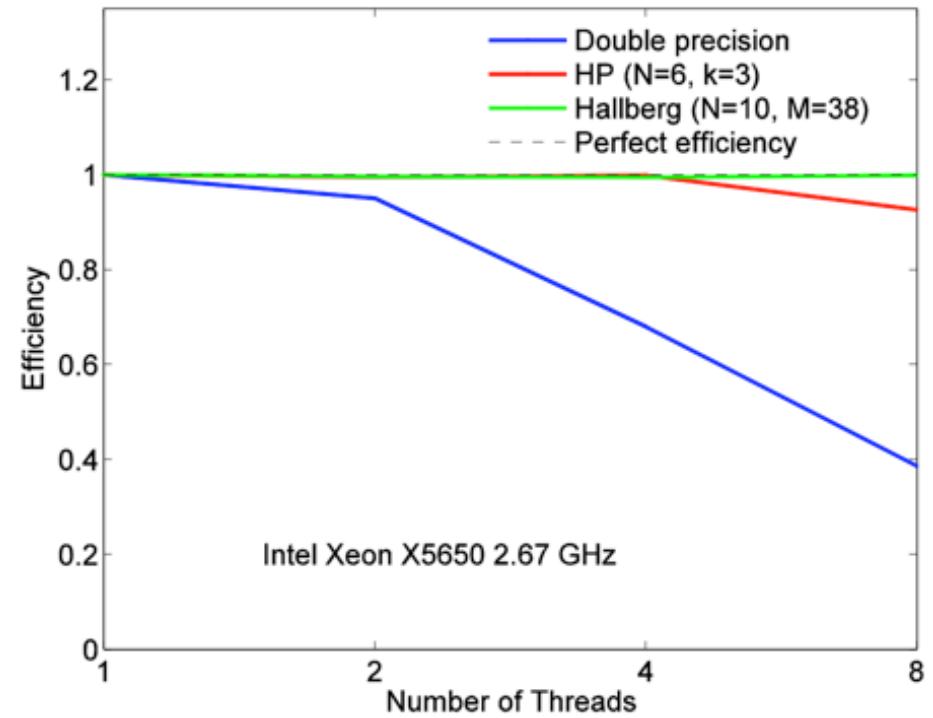
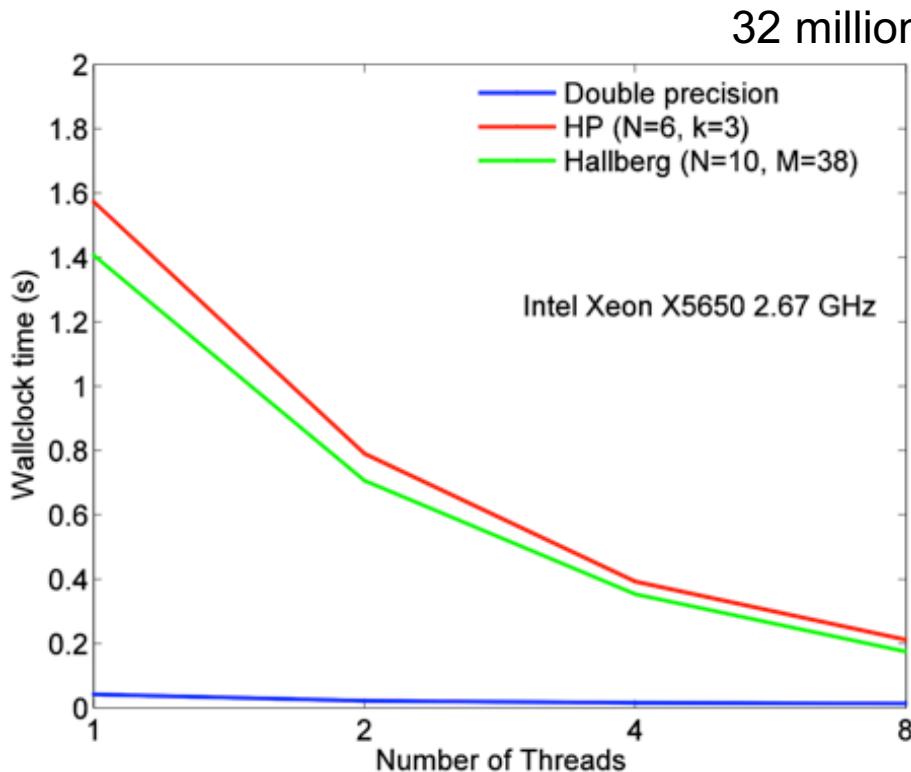
# Performance Projection

- HP sum is faster than Hallberg sum for larger numbers of summands & higher precision



# Parallel Efficiency with OpenMP

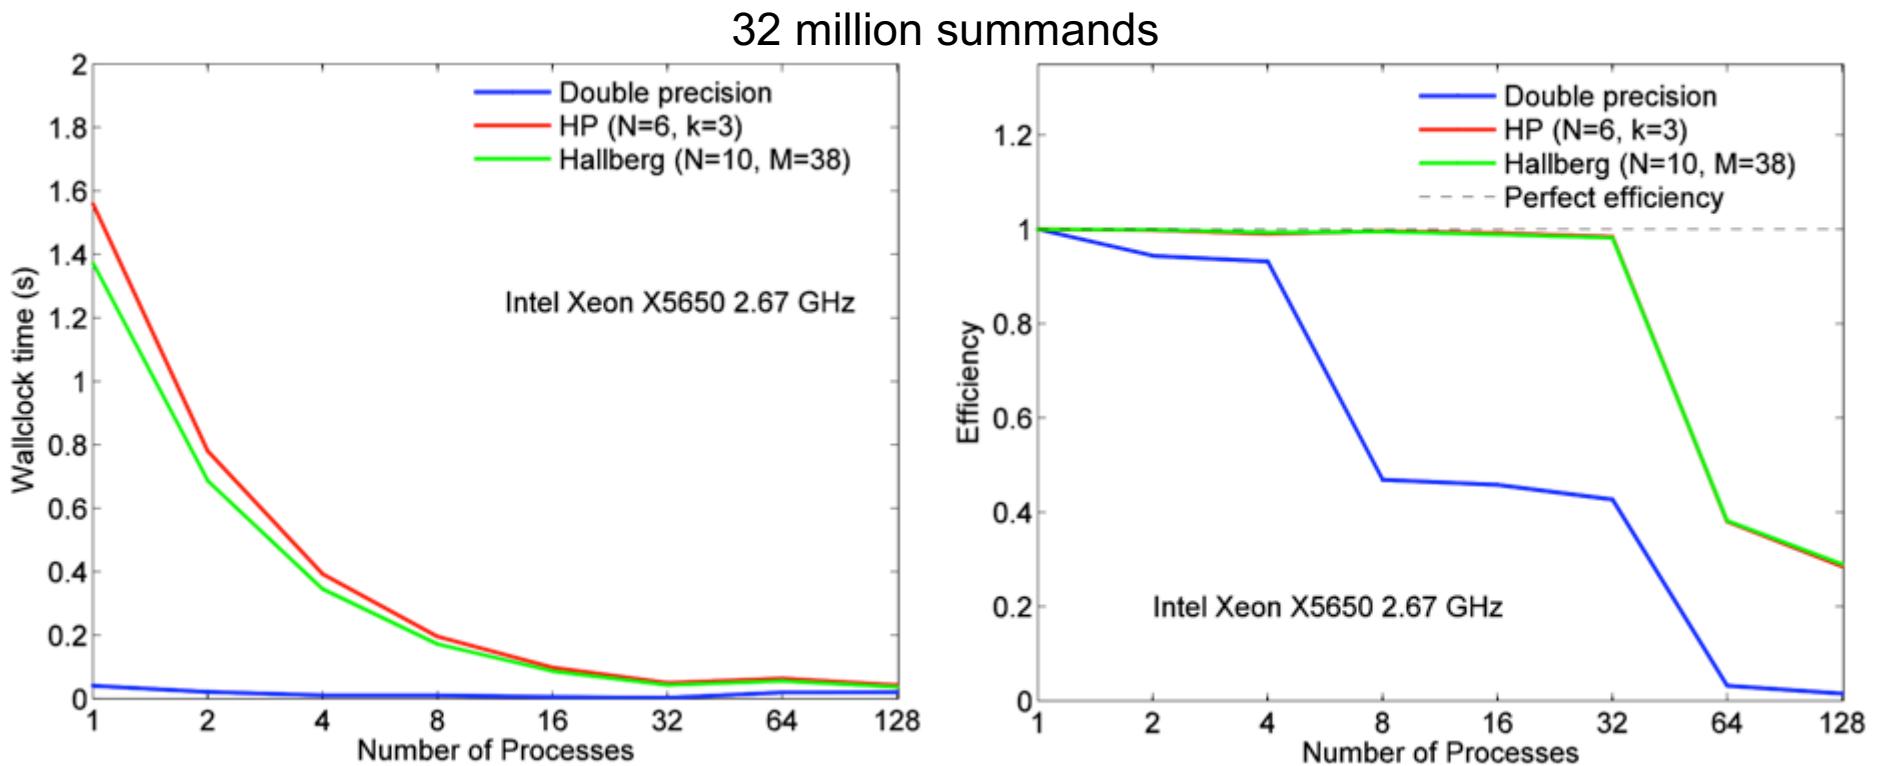
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of OpenMP threads on Xeon



- Higher parallel efficiency of HP & Hallberg sums over double-precision sum

# Parallel Efficiency with MPI

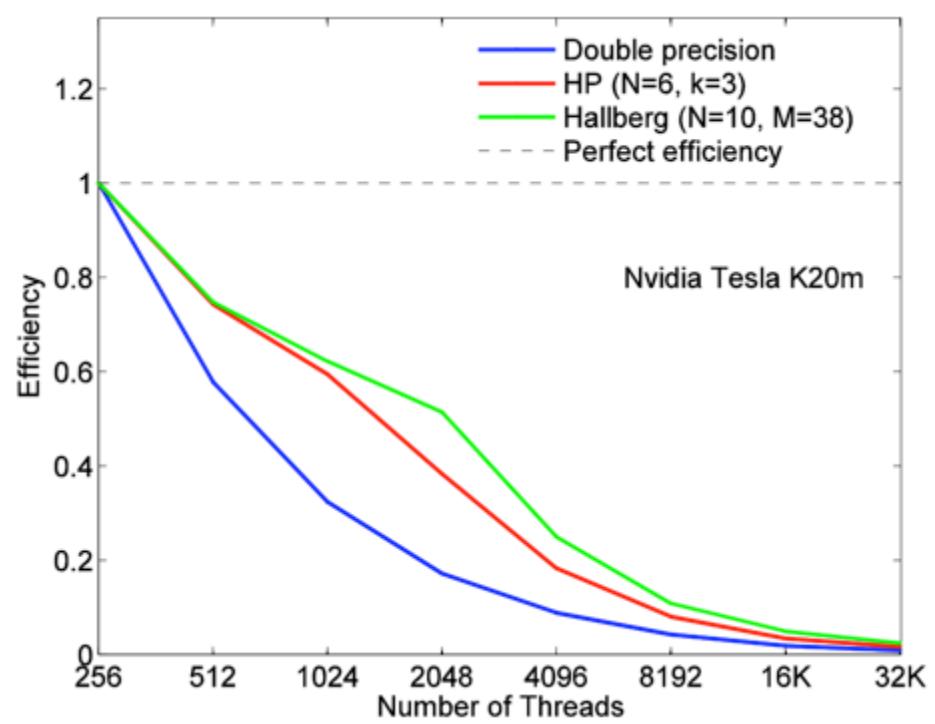
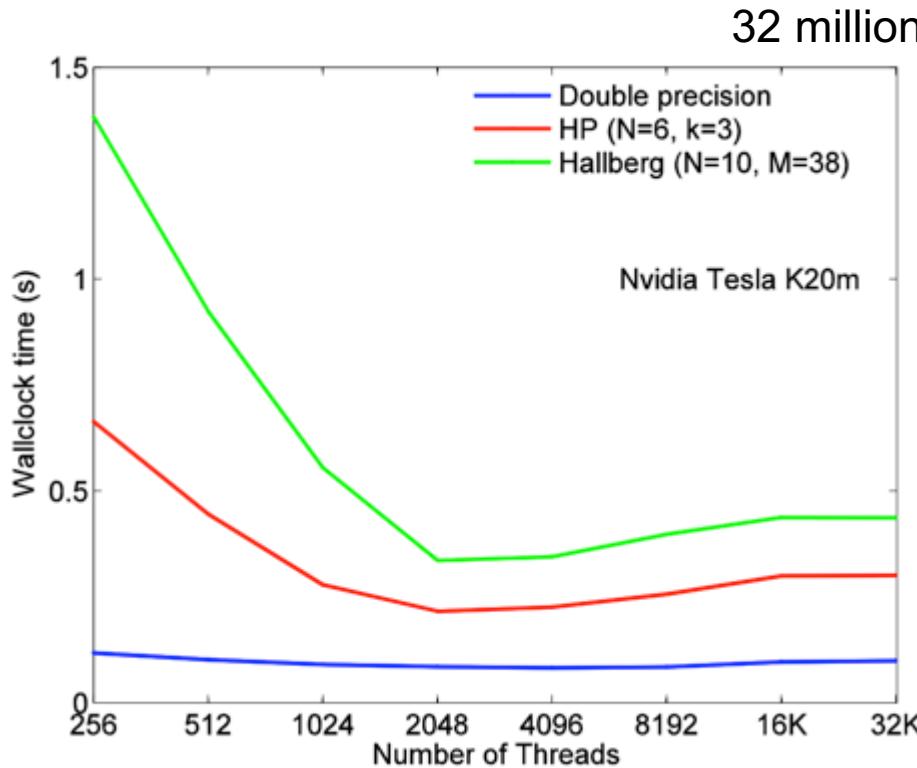
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of MPI processes on Xeon



- Higher parallel efficiency of HP & Hallberg sums over double-precision sum

# Parallel Efficiency on GPGPU

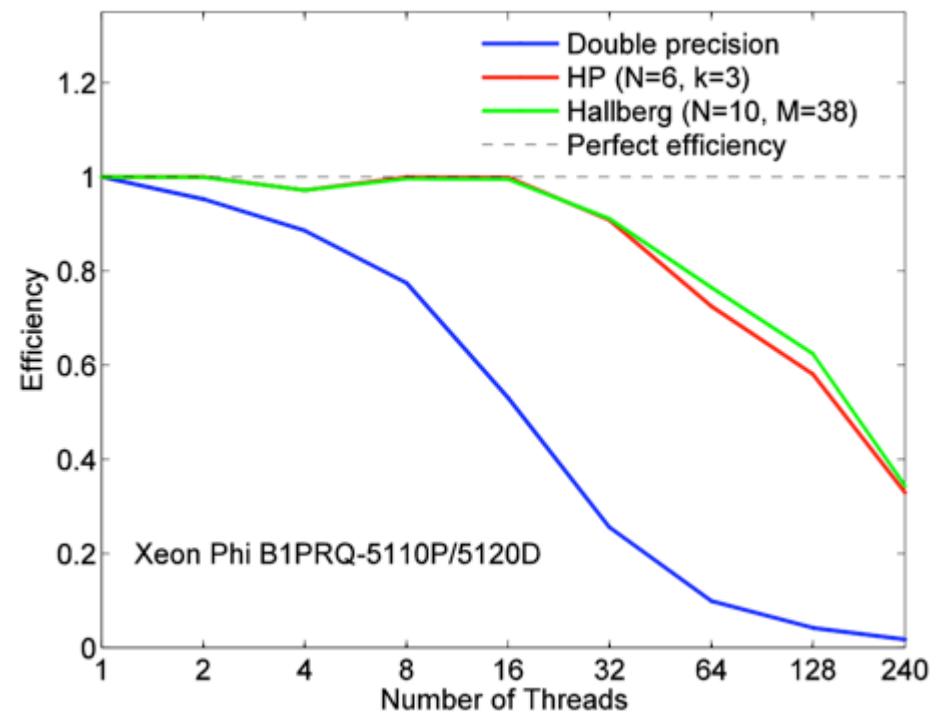
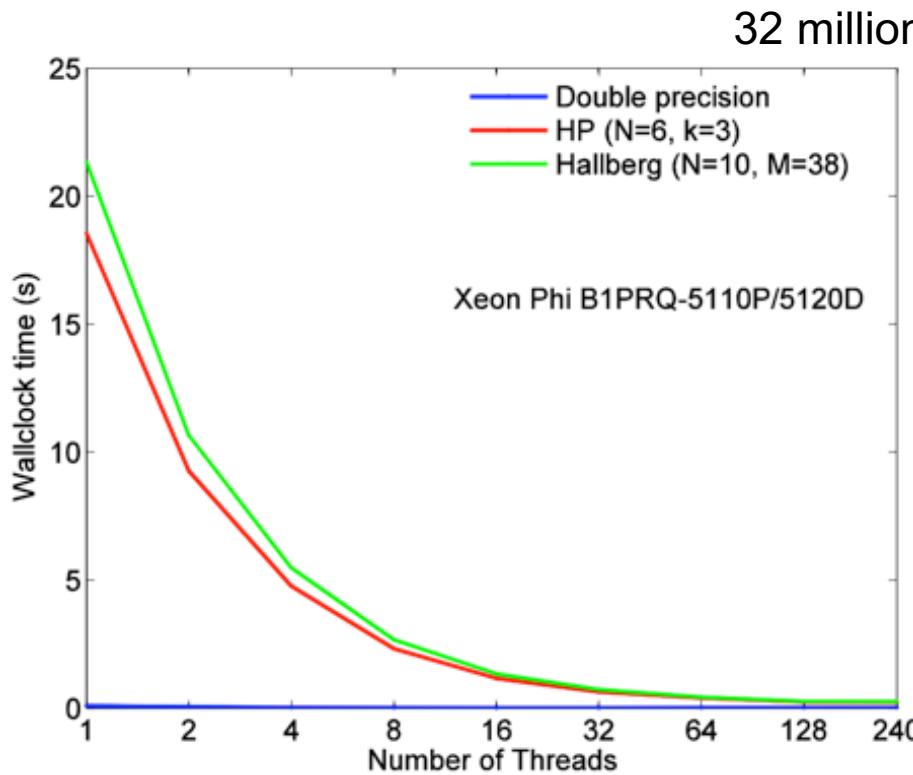
- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of CUDA threads on general-purpose graphics processing unit (GPGPU)



- Faster speed of HP sum (7 reads & 6 writes on global memory) over Hallberg sum (11 reads & 10 writes)

# Parallel Efficiency on Xeon Phi

- Runtime & strong-scaling parallel efficiency of HP, Hallberg & (order-sensitive) double-precision sums as a function of the number of threads on Intel Xeon Phi co-processor



- Faster speed of HP sum over Hallberg sum

# Large Production Simulations

- **16,661-atom quantum molecular dynamics (QMD) simulation on 786,432 IBM Blue Gene/Q cores suggests a rapid H<sub>2</sub>-production technology that is industrially scalable**

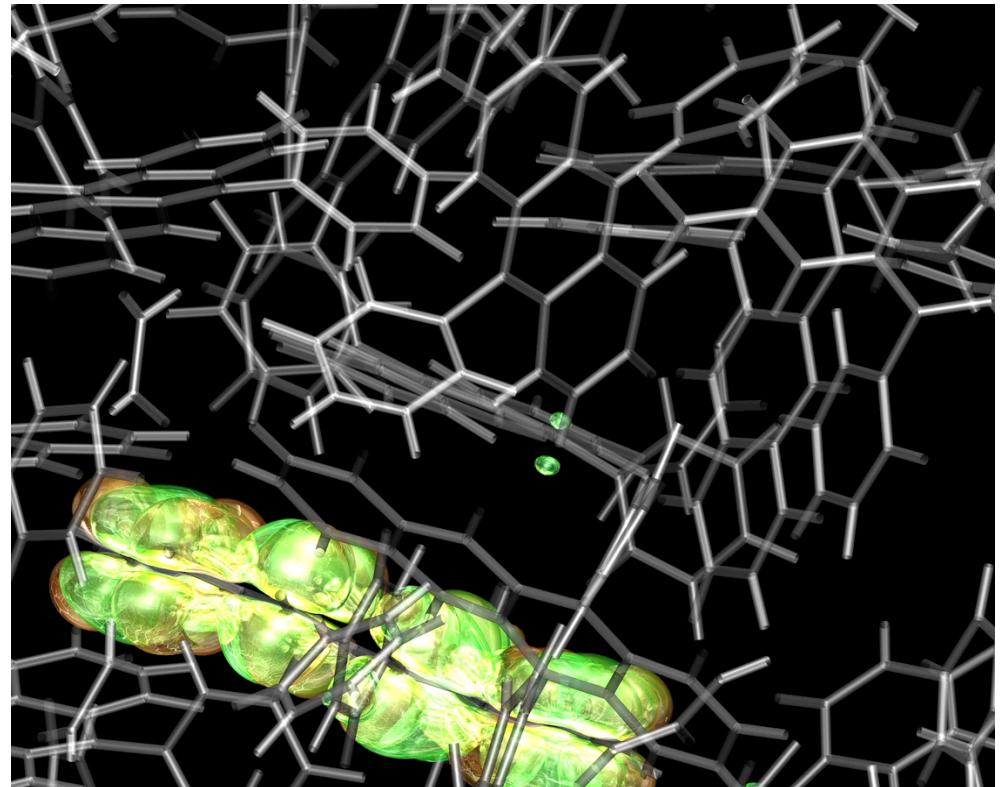
21,140 time steps (129,208 self-consistent-field iterations);  
*Nano Lett.* **14**, 4090 ('14)

- **Up-to 6,400-atom divide-conquer-recombine nonadiabatic QMD simulation reaches experimental time scales from first principles for photoexcitation dynamics**

*Appl. Phys. Lett.* **102**, 173301 ('13);  
*Sci. Rep.* **5**, 19599 ('16)

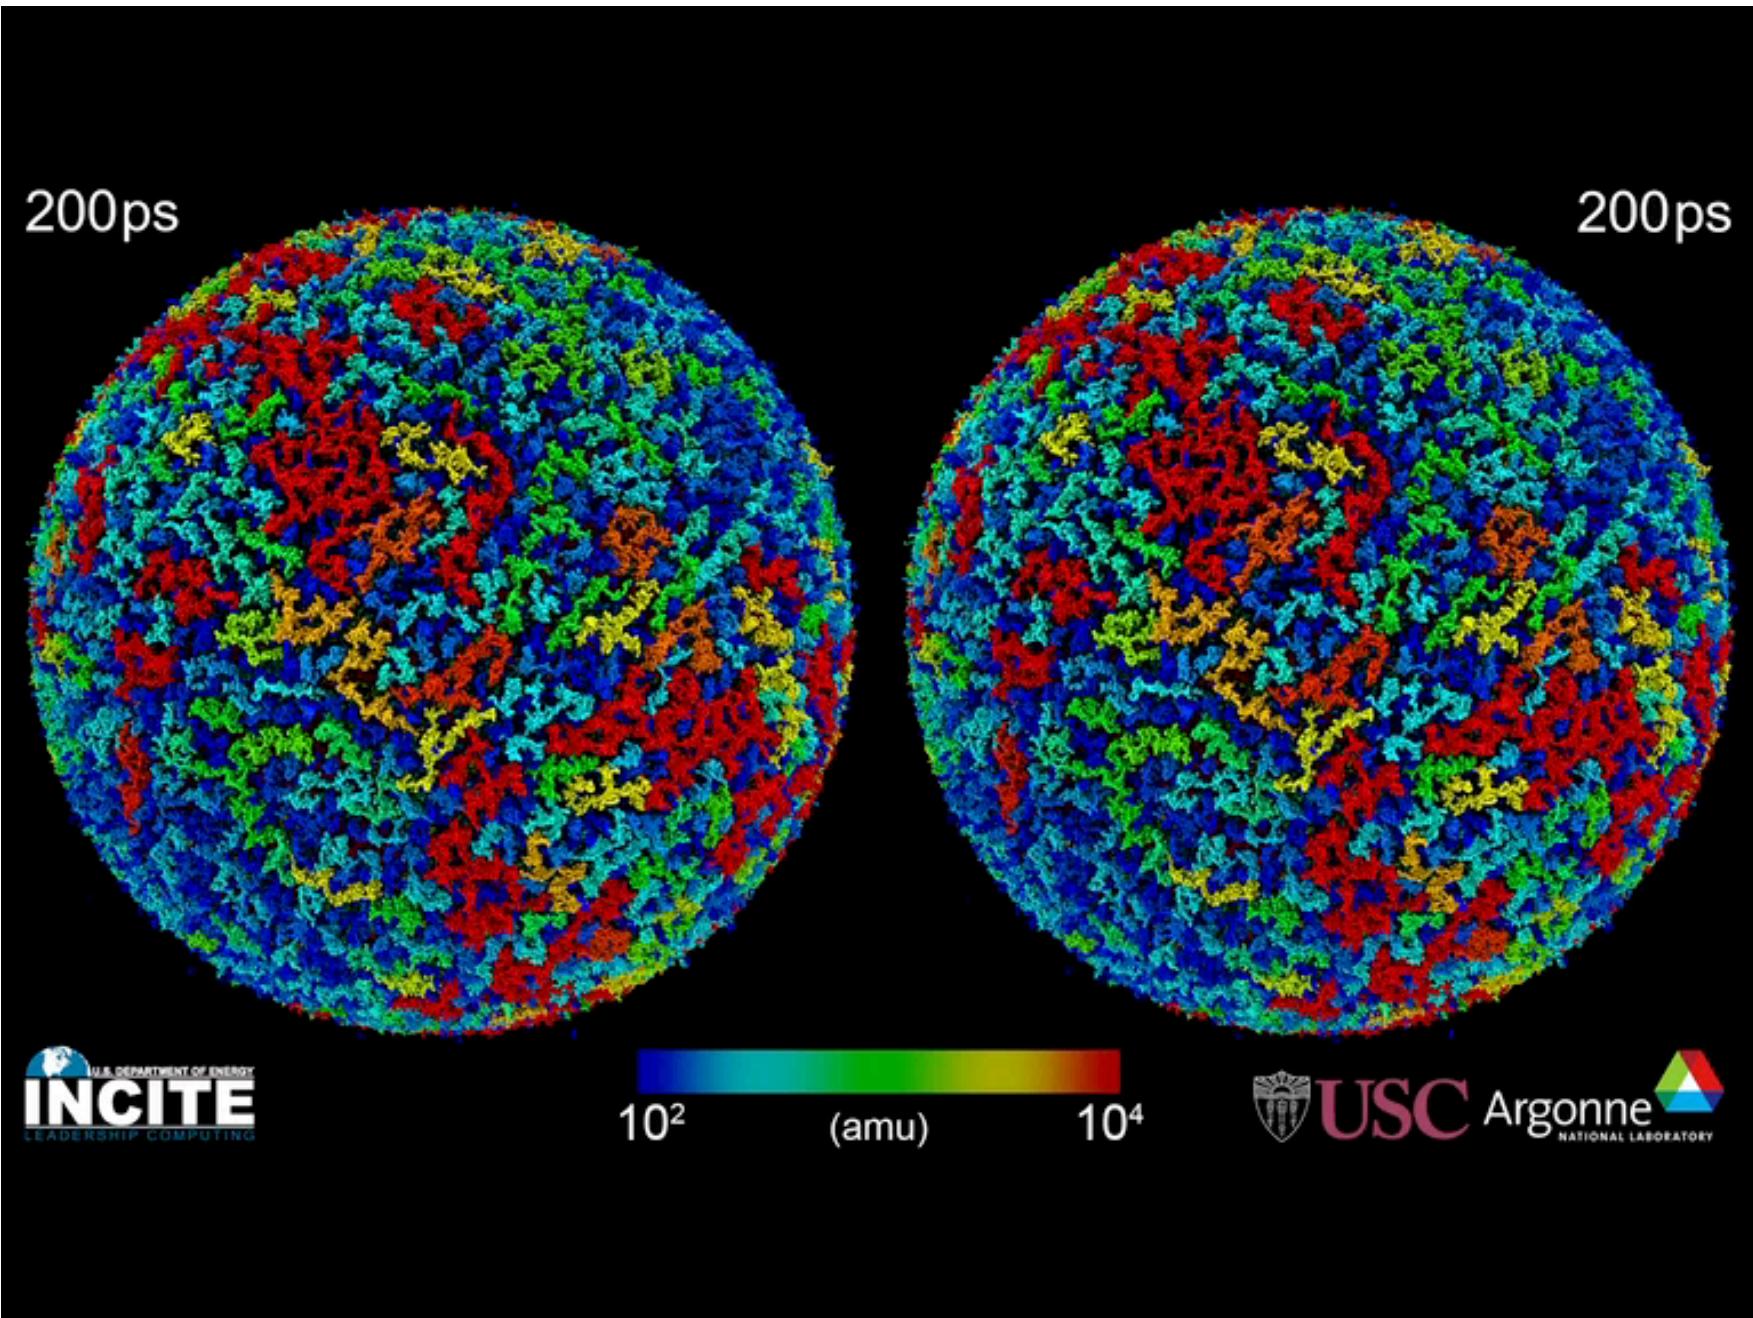
- **112 million-atom reactive molecular dynamics (RMD) simulation on 786,432 IBM Blue Gene/Q cores reveals a simple synthetic pathway to fractal graphene**

*Sci. Rep.* **6**, 24109 ('16)



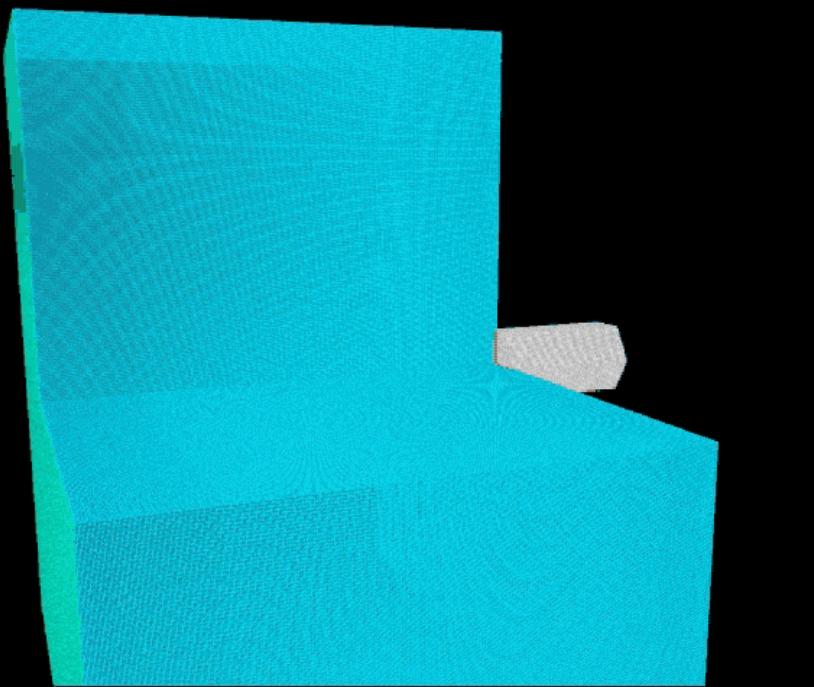
Quasi-electron Quasi-hole

# Percolation Transition



Movie made by J. Insley (Argonne)

# Billion-Atom Molecular Dynamics

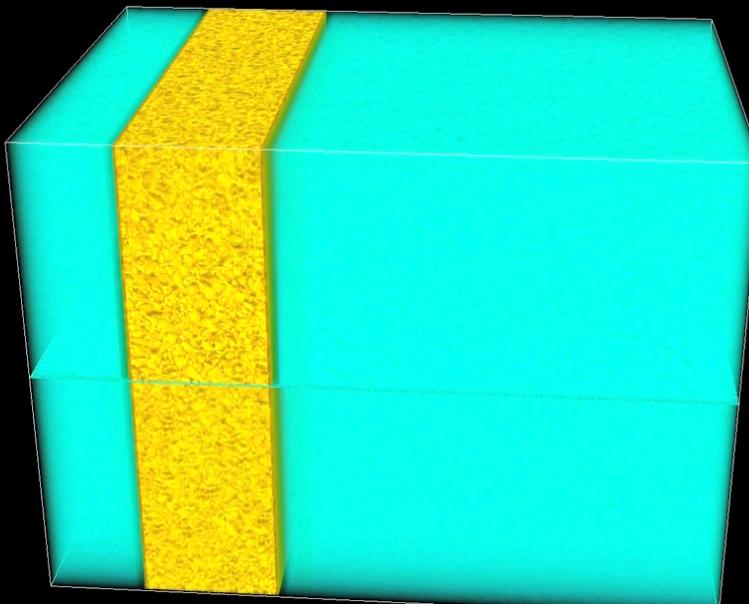


- **Hypervelocity impact on AlN**

P. S. Branicio *et al.*,  
*Phys. Rev. Lett.* **96**, 065502 ('06)

- Shock-induced nanobubble collapse in water near silica surface (67 million core-hours of computing on 163,840 Blue Gene/P cores)

A. Shekhar *et al.*, *Phys. Rev. Lett.* **111**, 184503 ('13)



— 100 nm

# Conclusion

---

- 1. An order-invariant real-number summation method has been proposed for reproducible parallel computing**
- 2. The proposed method achieves higher computing speed than the previous state-of-the-art for million+ summands on various parallel systems (MPI, OpenMP, CUDA, Xeon Phi)**

**Thank You**

**Research supported by  
DOE Grant DE-SC0014607**

